

Finite difference methods for wave motion

Hans Petter Langtangen^{1,2}

¹Center for Biomedical Computing, Simula Research Laboratory

²Department of Informatics, University of Oslo

Nov 14, 2012

Note: **VERY PRELIMINARY VERSION!** (Expect lots of typos!)

Contents

1	Finite difference methods for waves on a string	2
1.1	A finite difference method	2
1.2	Generalization: including a source term	7
2	Implementation	9
2.1	Making a solver function	9
2.2	Verification: exact quadratic solution	10
2.3	Visualization: animating $u(x, t)$	11
3	Vectorization	14
3.1	Operations on slices of arrays	15
3.2	Finite difference schemes expressed as slices	17
4	Exercises	18
5	Generalization: reflecting boundaries	19
5.1	Neumann boundary condition	19
5.2	Discretization of derivatives at the boundary	19
5.3	Implementation of Neumann conditions	20
5.4	Alternative implementation via ghost cells	21
6	Generalization: variable wave velocity	22
6.1	The model PDE with a variable coefficient	22
6.2	Discretizing the variable coefficient	23
6.3	Computing the coefficient between mesh points	23
6.4	How a variable coefficient affects the stability	24
6.5	Implementation of variable coefficients	24
6.6	A more general model PDE with variable coefficients	25
6.7	Generalization: including damping	25

7	Building a general 1D wave equation solver	26
7.1	User action function as a class	26
7.2	Collection of initial conditions	28
7.3	Calling functions from the command line	29
8	Exercises	30
9	Finite difference methods for 2D and 3D wave equations	31
9.1	Multi-dimensional wave equations	31
9.2	Mesh	33
9.3	Discretization	33
10	Implementation	35
10.1	Scalar computations	36
10.2	Vectorized computations	37
10.3	Verification	39
10.4	Migrating loops to Cython	39
10.5	Migrating loops to Fortran	43
10.6	Migrating loops to C via Cython	47
10.7	Migrating loops to C via f2py	50
10.8	Migrating loops to C via Instant	51
10.9	Migrating loops to C++ via f2py	51
10.10	Using classes to implement a simulator	52
10.11	Callbacks to Python from Fortran or C	52
11	Analysis of the continuous and discrete solutions	52
11.1	Properties of the solution of the wave equation	52
11.2	Analysis of the finite difference scheme	53
11.3	Extending the analysis to 2D and 3D	57
12	Applications of wave equations	59
12.1	Waves on a string	59
12.2	Waves on a membrane	62
12.3	Elastic waves in a rod	62
12.4	The acoustic model for seismic waves	62
12.5	Sound waves in liquids and gases	64
12.6	Spherical waves	65
12.7	The linear shallow water equations	66
12.8	Waves in blood vessels	68
12.9	Electromagnetic waves	70
13	Exercises	71

List of exercises

Exercise	1	Add storage of solution in a user action function ...	p. 18
Exercise	2	Use a class for the user action function	p. 18
Exercise	3	Compare several Courant numbers in one movie	p. 18
Exercise	4	Use ghost cells to implement Neumann conditions ...	p. 30
Exercise	5	Find a symmetry boundary condition	p. 30
Exercise	6	Prove symmetry of a 1D wave problem computationally ...	p. 30
Exercise	7	Prove symmetry of a 1D wave problem analytically ...	p. 30
Exercise	8	Send pulse waves through a layered medium	p. 31
Exercise	9	Explain why numerical noise occurs	p. 31
Exercise	10	Investigate harmonic averaging in a 1D model	p. 31
Exercise	11	Simulate elastic waves in a rod	p. 71
Exercise	12	Test the efficiency of compiled loops in 3D	p. 71
Exercise	13	Earthquake-generated tsunami in a 1D model	p. 71
Exercise	14	Implement an open boundary condition	p. 72
Exercise	15	Earthquake-generated tsunami over a subsea ...	p. 74
Exercise	16	Implement Neumann conditions in 2D	p. 75
Exercise	18	Earthquake-generated tsunami over a 3D hill	p. 76
Exercise	19	Implement loops in compiled languages	p. 77
Exercise	20	Write a complete program in Fortran or C	p. 77
Exercise	21	Investigate Matplotlib for visualization	p. 77
Exercise	22	Investigate Mayavi for visualization	p. 77
Exercise	23	Investigate Paraview for visualization	p. 77
Exercise	24	Investigate OpenDX for visualization	p. 77
Exercise	25	Investigate harmonic vs arithmetic mean	p. 77
Exercise	26	Simulate seismic waves in 2D	p. 78

A very wide range of physical processes lead to wave motion, where signals are propagated through a medium in space and time, normally with little or no permanent movement of the medium itself. The shape of the signals may undergo changes as they travel through matter, but usually not so much that the signals cannot be recognized at some later point in space and time. Many types of wave motion can be described by the *wave equation* $u_{tt} = \nabla \cdot (c^2 \nabla u) + f$, which we will solve in the forthcoming text by finite difference methods.

1 Finite difference methods for waves on a string

We begin our study of wave equations by simulating one-dimensional waves on a string, say on a guitar or violin string. Let the string in the deformed state coincide with the interval $[0, L]$ on the x axis, and let $u(x, t)$ be the displacement at time t in the y direction of a point initially at x . The displacement function u is governed by the mathematical model

$$\frac{\partial^2 u}{\partial t^2} = c^2 \frac{\partial^2 u}{\partial x^2}, \quad x \in (0, L), \quad t \in (0, T] \quad (1)$$

$$u(x, 0) = I(x), \quad x \in [0, L] \quad (2)$$

$$\frac{\partial}{\partial t} u(x, 0) = 0, \quad x \in [0, L] \quad (3)$$

$$u(0, t) = 0, \quad t \in (0, T], \quad (4)$$

$$u(L, t) = 0, \quad t \in (0, T]. \quad (5)$$

The constant c and the function $I(x)$ must be prescribed.

Equation (1) is known as the one-dimensional *wave equation*. Since this PDE contains a second-order derivative in time, we need *two initial conditions*, here (2) specifying the initial shape of the string, $I(x)$, and (3) reflecting that the initial velocity of the string is zero. In addition, PDEs need *boundary conditions*, here (4) and (5), specifying that the string is fixed at the ends, i.e., that the displacement u is zero at the ends.

Sometimes we will use a more compact notation for the partial derivatives to save space:

$$u_t = \frac{\partial u}{\partial t}, \quad u_{tt} = \frac{\partial^2 u}{\partial t^2}, \quad (6)$$

and similar for derivatives with respect to other variables. Then the wave equation can be written compactly as $u_{tt} = c^2 u_{xx}$.

1.1 A finite difference method

The PDE problem (1)-(5) will now be discretized in space and time by a finite difference method.

Step 1: Discretizing the Domain. The temporal domain $[0, T]$ is represented by a finite number of mesh points

$$t_0 = 0 < t_1 < t_2 < \cdots < t_{N-1} < t_N = T. \quad (7)$$

Similarly, the spatial domain $[0, L]$ is replaced by a set of mesh points

$$0 = x_0 < x_1 < x_2 < \cdots < x_{N_x-1} < x_{N_x} = L. \quad (8)$$

One may view the mesh as two-dimensional in the x, t plane, consisting of points (x_i, t_n) , $i = 0, \dots, N_x$, $n = 0, \dots, N$.

For uniformly distributed mesh points we can introduce the constant mesh spacings Δt and Δx . We have that

$$x_i = i\Delta x, \quad i = 0, \dots, N_x, \quad (9)$$

and

$$t_i = n\Delta t, \quad n = 0, \dots, N. \quad (10)$$

We also have that $\Delta x = x_i - x_{i-1}$, $i = 1, \dots, N_x$, and $\Delta t = t_n - t_{n-1}$, $n = 1, \dots, N$. Figure 1 displays a mesh in the x - t plane with $N = 5$, $N_x = 5$, and constant mesh spacings.




Figure 1: Mesh in space and time for a 1D wave equation.

The solution $u(x, t)$ is sought at the mesh points: $u(x_i, t_n)$ for $i = 0, \dots, N_x$ and $n = 0, \dots, N$. The numerical approximation at mesh point (x_i, t_n) is denoted by u_i^n . The circles in Figure 1 illustrates neighboring mesh points where values of u_i^n are connected through a discrete equation. In this particular case, $u_2^1, u_1^2, u_2^2, u_3^2$, and u_2^3 are connected in a discrete equation associated with the center point $(2, 2)$. The term *stencil* is often used about the discrete equation at a mesh point, and the geometry of a typical stencil is illustrated in Figure 1.

Step 2: Fulfilling the equation at the mesh points. For a numerical solution by the finite difference method, we relax the condition that (1) holds at all points in the domain $(0, L) \times (0, T]$ to the requirement that the PDE is fulfilled at the mesh points:

$$\frac{\partial^2}{\partial t^2} u(x_i, t_n) = c^2 \frac{\partial^2}{\partial x^2} u(x_i, t_n), \quad (11)$$

for $i = 1, \dots, N_x - 1$ and $n = 1, \dots, N$. For $n = 0$ we have the initial conditions $u = I(x)$ and $u_t = 0$, and at the boundaries $i = 0, N_x$ we have the boundary condition $u = 0$.

Step 3: Replacing derivatives by finite differences. The second-order derivatives can be replaced by central differences. The most widely used difference approximation of the second-order derivative is

$$\frac{\partial^2}{\partial t^2} u(x_i, t_n) \approx \frac{u_i^{n+1} - 2u_i^n + u_i^{n-1}}{\Delta t^2} = [D_t D_t u]_i^n,$$

and a similar definition of the second-order derivative in x direction:

$$\frac{\partial^2}{\partial x^2} u(x_i, t_n) \approx \frac{u_{i+1}^n - 2u_i^n + u_{i-1}^n}{\Delta x^2} = [D_x D_x u]_i^n.$$

We can now replace the derivatives in (11) and get

$$\frac{u_i^{n+1} - 2u_i^n + u_i^{n-1}}{\Delta t^2} = c^2 \frac{u_{i+1}^n - 2u_i^n + u_{i-1}^n}{\Delta x^2}, \quad (12)$$

or written more compactly using the operator notation:

$$[D_t D_t u = c^2 D_x D_x]_i^n. \quad (13)$$

We also need to replace the derivative in the initial condition (3) by a finite difference approximation. A centered difference of the type

$$\frac{\partial}{\partial t} u(x_i, t_n) \approx \frac{u_i^1 - u_i^{-1}}{2\Delta t} = [D_{2t} u]_i^0,$$

seems appropriate. In operator notation the initial condition is written as

$$[D_{2t} u]_i^0 = 0.$$

Writing out this equation and ordering the terms give

$$u_i^{-1} = u_i^1, \quad i = 0, \dots, N_x. \quad (14)$$

Step 4: Formulating a Recursive Algorithm. As for ordinary differential equations, we assume that u_i^n and u_i^{n-1} are already computed for $i = 0, \dots, N_x$. The only unknown quantity in (12) is therefore u_i^{n+1} , which we can solve for:

$$u_i^{n+1} = -u_i^{n-1} + 2u_i^n + C^2 (u_{i+1}^n - 2u_i^n + u_{i-1}^n), \quad (15)$$

where we have introduced the parameter

$$C = c \frac{\Delta t}{\Delta x}, \quad (16)$$

known as the (dimensionless) *Courant number*. We see that the discrete version of the PDE features only one parameter, C , which is therefore the key parameter that governs the quality of the numerical solution. Both the primary physical parameter c and the numerical parameters Δx and Δt are lumped together in C .

Given that u_i^{n-1} and u_i^n are computed for $i = 0, \dots, N_x$, we find new values at the next time level by applying the formula (15) for $i = 1, \dots, N_x - 1$. Figure 1 illustrates the points that are used to compute u_2^3 . For the boundary points, $i = 0$ and $i = N_x$, we apply the boundary conditions: $u_0^{n+1} = 0$ and $u_{N_x}^{n+1} = 0$.

A problem with (15) arises when $n = 0$ since the formula for u_i^1 involves u_i^{-1} , which is an undefined quantity outside the time mesh and the time domain. However, we can use the boundary condition (14) in combination with (15) when $n = 0$ to arrive at a special formula for u_i^1 :

$$u_i^1 = u_i^0 - \frac{1}{2}C^2 (u_{i+1}^n - 2u_i^n + u_{i-1}^n) . \quad (17)$$

Figure 2 illustrates how (17) connects four instead of five points: u_2^1 , u_1^0 , u_2^0 , and u_3^0 .




Figure 2: Modified stencil for the first time step.

We can now summarize the computational algorithm:

- compute $u_i^0 = I(x_i)$ for $i = 0, \dots, N_x$
- compute u_i^1 by (17) and set $u_i^1 = 0$ for the boundary points $i = 0$ and $i = N_x$, for $n = 1, 2, \dots, N - 1$,
- apply (15) for $i = 1, \dots, N_x - 1$ and set $u_i^{n+1} = 0$ for the boundary points $i = 0$ and $i = N_x$.

The algorithm essentially consists of moving a finite difference stencil through all the mesh points, which is illustrated by an animation¹

¹http://hplgit.github.com/INF5620/doc/notes/mov-wave/wave1D_PDE_Dirichlet_stencil_gpl/index.html

Sketch of an implementation. In a Python implementation of this algorithm, we use the array elements $u[i]$ to store u_i^{n+1} , $u_1[i]$ to store u_i^n , and $u_2[i]$ to store u_i^{n-1} . The algorithm only needs to access the three most recent time levels, so we need only three arrays for u_i^{n+1} , u_i^n , and u_i^{n-1} , $i = 0, \dots, N_x$. Storing all the solutions in a two-dimensional array of size $(N_x + 1) \times (N + 1)$ would be possible in this simple one-dimensional PDE problem, but is normally out of the question in three-dimensional (3D) and large two-dimensional (2D) problems. We shall therefore in all our programs for solving PDEs have the unknown in memory at as few time levels as possible.

The following Python snippet realizes the steps in the computational algorithm.

```
# Given mesh points as arrays x and t (x[i], t[n])
dx = x[1] - x[0]; dt = t[1] - t[0]
C = c*dt/dx
N = len(t)-1
C2 = C**2                # help variable in the scheme

# Set initial condition u(x,0) = I(x)
for i in range(0, Nx+1):
    u_1[i] = I(x[i])

# Apply special formula for first step, incorporating du/dt=0
for i in range(1, Nx):
    u[i] = u_1[i] - 0.5*C**2*(u_1[i+1] - 2*u_1[i] + u_1[i-1])
u[0] = 0; u[Nx] = 0

u_2[:,], u_1[:,] = u_1, u

for n in range(1, N):
    # Update all inner mesh points at time t[n+1]
    for i in range(1, Nx):
        u[i] = 2*u_1[i] - u_2[i] - \
            C**2*(u_1[i+1] - 2*u_1[i] + u_1[i-1])

    # Insert boundary conditions
    u[0] = 0; u[Nx] = 0

    # Switch variables before next step
    u_2[:,], u_1[:,] = u_1, u
```

1.2 Generalization: including a source term

Before implementing the algorithm, it is convenient to add a source term to the PDE (1) since it gives us more freedom in finding test problems for verification. In particular, the source term allows us to use *manufactured solutions* for software testing, where we simply choose some function as solution, fit the corresponding source term, and define consistent boundary and initial conditions. Such solutions will seldom fulfill the initial condition (3) so we need to generalize this condition to $u_t = V(x)$.

The extended problem. We now address the following extended initial-boundary value problem for one-dimensional wave phenomena:

$$u_{tt} = c^2 u_{xx} + f(x, t), \quad x \in (0, L), \quad t \in (0, T] \quad (18)$$

$$u(x, 0) = I(x), \quad x \in [0, L] \quad (19)$$

$$u_t(x, 0) = V(x), \quad x \in [0, L] \quad (20)$$

$$u(0, t) = 0, \quad t > 0, \quad (21)$$

$$u(L, t) = 0, \quad t > 0. \quad (22)$$

Discretization. Sampling the PDE at (x_i, t_n) and using the same finite difference approximations as above, yields

$$[D_t D_t u = c^2 D_x D_x + f]_i^n. \quad (23)$$

Writing this out and solving for the unknown u_i^{n+1} results in

$$u_i^{n+1} = -u_i^{n-1} + 2u_i^n + C^2(u_{i+1}^n - 2u_i^n + u_{i-1}^n) + \Delta t^2 f_i^n. \quad (24)$$

The equation for the first time step must be rederived. The discretization of the initial condition $u_t = V(x)$ at $t = 0$ becomes

$$[D_{2t} u = V]_i^0 \Rightarrow u_i^{-1} = u_i^1 - 2\Delta t V_i,$$

which, when inserted in (24) for $n = 0$, gives

$$u_i^1 = u_i^0 - \Delta t V_i + \frac{1}{2} C^2 (u_{i+1}^0 - 2u_i^0 + u_{i-1}^0) + \frac{1}{2} \Delta t^2 f_i^0. \quad (25)$$

Constructing an exact solution of the discrete equations. For verification purposes we shall use a solution that is quadratic in space and linear in time. More specifically, we choose

$$u_e(x, t) = x(L - x)(1 + \frac{1}{2}t),$$

which by insertion in the PDE leads to $f(x, t) = 2(1 + t)c^2$. Moreover, this u fulfills the boundary conditions and is compatible with $I(x) = x(L - x)$ and $V(x) = \frac{1}{2}x(L - x)$.

To see if the discrete values $u_e(x_i, t_n)$ fulfills the discrete equation as well, we first establish the results

$$[D_t D_t t^2]^n = \frac{t_{n+1}^2 - 2t_n^2 + t_{n-1}^2}{\Delta t^2} = (n+1)^2 - n^2 + (n-1)^2 = 2, \quad (26)$$

$$[D_t D_t t]^n = \frac{t_{n+1} - 2t_n + t_{n-1}}{\Delta t^2} = \frac{((n+1) - n + (n-1))\Delta t}{\Delta t^2} = 0. \quad (27)$$

Hence,

$$[D_t D_t u_e]_i^n = x_i(L - x_i)[D_t D_t(1 + \frac{1}{2}t)]^n = x_i(L - x_i)\frac{1}{2}[D_t D_t t]^n = 0,$$

and

$$\begin{aligned} [D_x D_x u_e]_i^n &= (1 + \frac{1}{2}t_n)[D_x D_x(xL - x^2)]_i = (1 + \frac{1}{2}t_n)[LD_x D_x x - D_x D_x x^2]_i \\ &= -2(1 + \frac{1}{2}t_n). \end{aligned}$$

Since $f_i^n = 2(1 + \frac{1}{2}t_n)c^2$, we see that the scheme accepts u_e as a solution. Moreover, $u_e(x_i, 0) = I(x_i)$, $\partial u_e / \partial t = V(x_i)$ at $t = 0$, and $u_e(x_0, t) = u_e(x_{N_x}, 0) = 0$. Also the modified scheme for the first time step is fulfilled by $u_e(x_i, t_n)$. Therefore, the exact solution $u_e(x, t)$ of the PDE problem is also an exact solution of the discrete problem. We can use this result to check that the computed u_i^n vales from an implementation equals $u_e(x_i, t_n)$ within machine precision, *regardless of the mesh spacings Δx and Δt* ! Nevertheless, there might be stability restrictions on Δx and Δt , so the test can only be run for a mesh that is compatible with the stability criterion (which is $C \leq 1$, to be derived later).

A product of quadratic or linear expressions in the various independent variables, as shown above, will often fulfill both the continuous and discrete PDE problem and can therefore be very useful solutions for verifying implementations. However, for 1D wave equations of the type $u_t = c^2 u_{xx}$ we shall see that there is always another much more powerful way of generating exact solutions (just set $C = 1$).

2 Implementation

A real implementation of the basic computational algorithm can be encapsulated in a function, taking all the input data for the problem as arguments. The physical input data consists of c , $I(x)$, $V(x)$, $f(x, t)$, L , and T . The numerical input is the mesh parameters Δt and Δx . One possibility is to specify N_x and the Courant number $C = c\Delta t / \Delta x$. The latter is convenient to prescribe instead of Δt when performing numerical investigations, because the numerical accuracy depends directly on C .

The solution at all spatial points at a new time level is stored in an array \mathbf{u} (of length $N_x + 1$). We need to decide what do to with this solution, e.g., visualize the curve, analyze the values, or write the array to file for later use. The decision what to do is left to the user in a supplied function

```
def user_action(u, x, t, n)
```

where \mathbf{u} is the solution at the spatial points \mathbf{x} at time $\mathbf{t}[\mathbf{n}]$.

2.1 Making a solver function

A first attempt at a solver function is listed below.

```
from numpy import *

def solver(I, V, f, c, L, Nx, C, T, user_action=None):
    """Solve  $u_{tt}=c^2u_{xx} + f$  on  $(0,L) \times (0,T]$ ."""
    x = linspace(0, L, Nx+1) # mesh points in space
    dx = x[1] - x[0]
    dt = C*dx/c
    N = int(round(T/dt))
    t = linspace(0, N*dt, N+1) # mesh points in time
    C2 = C**2 # help variable in the scheme
    if f is None or f == 0 :
        f = lambda x, t: 0
    if V is None or V == 0:
        V = lambda x: 0

    u = zeros(Nx+1) # solution array at new time level
    u_1 = zeros(Nx+1) # solution at 1 time level back
    u_2 = zeros(Nx+1) # solution at 2 time levels back

    import time; t0 = time.clock() # for measuring CPU time

    # Load initial condition into u_1
    for i in range(0, Nx+1):
        u_1[i] = I(x[i])

    if user_action is not None:
        user_action(u_1, x, t, 0)

    # Special formula for first time step
    n = 0
    for i in range(1, Nx):
        u[i] = u_1[i] + dt*V(x[i]) + \
            0.5*C2*(u_1[i-1] - 2*u_1[i] + u_1[i+1]) + \
            0.5*dt**2*f(x[i], t[n])
    u[0] = 0; u[Nx] = 0

    if user_action is not None:
        user_action(u, x, t, 1)

    u_2[:,], u_1[:,] = u_1, u

    for n in range(1, N):
        # Update all inner points at time t[n+1]
        for i in range(1, Nx):
            u[i] = -u_2[i] + 2*u_1[i] + \
                C2*(u_1[i-1] - 2*u_1[i] + u_1[i+1]) + \
                dt**2*f(x[i], t[n])

        # Insert boundary conditions
        u[0] = 0; u[Nx] = 0
        if user_action is not None:
            if user_action(u, x, t, n+1):
                break

        # Switch variables before next step
        u_2[:,], u_1[:,] = u_1, u
```

```

cpu_time = t0 - time.clock()
return u, x, t, cpu_time

```

2.2 Verification: exact quadratic solution

We use the test problem derived in Section 1.2 for verification. Here is a function realizing this verification as a nose test:

```

import nose.tools as nt

def test_quadratic():
    """Check that  $u(x,t)=x(L-x)(1+t)$  is exactly reproduced."""
    def exact_solution(x, t):
        return x*(L-x)*(1 + 0.5*t)

    def I(x):
        return exact_solution(x, 0)

    def V(x):
        return 0.5*exact_solution(x, 0)

    def f(x, t):
        return 2*(1 + 0.5*t)*c**2

    L = 2.5
    c = 1.5
    Nx = 3 # very coarse mesh
    C = 0.75
    T = 18

    u, x, t, cpu = solver(I, V, f, c, L, Nx, C, T)
    u_e = exact_solution(x, t[-1])
    diff = abs(u - u_e).max()
    nt.assert_almost_equal(diff, 0, places=14)

```

2.3 Visualization: animating $u(x, t)$

Now that we have verified the implementation it is time to do a real computation where we also display the evolution of the waves on the screen.

Visualization via SciTools. The following viz function defines a `user_action` callback function for plotting the solution at each time level:

```

def viz(I, V, f, c, L, Nx, C, T, umin, umax, animate=True):
    """Run solver and visualize u at each time level."""
    import scitools.std as st, time, glob, os

    def plot_u(u, x, t, n):
        """user_action function for solver."""
        st.plot(x, u, 'r-',
                xlabel='x', ylabel='u',
                axis=[0, L, umin, umax],
                title='t=%f' % t[n], show=True)

```

```

    # Let the initial condition stay on the screen for 2
    # seconds, else insert a pause of 0.2 s between each plot
    time.sleep(2) if t[n] == 0 else time.sleep(0.2)
    st.savefig('frame_%04d.png' % n) # for movie making

# Clean up old movie frames
for filename in glob.glob('frame_*.png'):
    os.remove(filename)

user_action = plot_u if animate else None
u, x, t, cpu = solver(I, V, f, c, L, Nx, C, T, user_action)

# Make movie files
st.movie('frame_*.png', encoder='mencoder', fps=4,
        output_file='movie.avi')
st.movie('frame_*.png', encoder='html', fps=4,
        output_file='movie.html')

```

A function inside another function, like `plot_u` in the above code segment, has access to *and remembers* (!) all the local variables in the surrounding code inside the `viz` function. This is known in computer science as a *closure* and is convenient. For example, the `st` and `time` modules defined outside `plot_u` are accessible for `plot_u` the function is called (as `user_action`) in the `solver` function. Some think, however, that a class instead of a closure is a cleaner and easier-to-understand implementation of the user action function, see Section 7.

Making movie files. Two hardcopies of the animation are made from the `frame_*.png` files, using `movie` function from SciTools at the end of the `viz` function. The AVI file `movie1.avi` can be played in a standard movie player. The HTML file `movie.html` is essentially a movie player that can be loaded into a browser to display the individual `frame_*.png` files. Note that padding the frame counter in the `frame_*.png` files (`%04d` format) is essential so that the wildcard notation `frame_*.png` expands to the correct set of files.

Rather than using the simple `movie` function, one can run native commands in the terminal window. An AVI movie can be made by `mencoder`:

```

Terminal> mencoder "mf://frame_%04d.png" -mf fps=4:type=png \
          -ovc lavc -o movie.avi
Terminal> mplayer movie.avi

```

The HTML player can only be generated by

```

Terminal> scitools movie output_file=movie.html fps=4 frame_*.png

```

The `fps` parameter controls the number of *frames per second*, i.e., the speed of the movie.

Skipping frames for animation speed. Sometimes the time step is small and T is large, leading to an inconveniently large number of plot files and a slow animation on the screen. The solution to such a problem is to decide on a total number of frames in the animation, `num_frames`, and plot the solution only at every `every` frame. The total number of time levels (i.e., maximum possible number of frames) is the length of `t`, `t.size`, and if we want `num_frames`, we need to plot every `t.size/num_frames` frame:

```
every = int(t.size/float(num_frames))
if n % every == 0 or n == t.size-1:
    st.plot(x, u, 'r-', ...)
```

The initial condition (`n=0`) is natural to include, and as `n % every == 0` will very seldom be true for the very final frame, we also add that frame (`n == t.size-1`).

A simple choice of numbers may illustrate the formulas: say we have 801 frames in total (`t.size`) and we allow only 60 frames to be plotted. Then we need to plot every 801/60 frame, which with integer division yields 13 as `every`. Using the mod function, `n % every`, this operation is zero every time `n` can be divided by 13 without a remainder. That is, the `if` test is true when `n` equals 0, 13, 26, 39, ..., 780, 801. The associated code is included in the `plot_u` function in the file `wave1D_u0_sv.py`.

Visualization via Matplotlib. The previous code based on the `plot` interface from `scitools.std` can be run with Matplotlib as the visualization backend, but if one desires to program directly with Matplotlib, quite different code is needed. Matplotlib's interactive mode must be turned on:

```
import matplotlib.pyplot as plt
plt.ion() # interactive mode on
```

The most commonly used animation technique with Matplotlib is to update the data in the plot at each time level:

```
# Make a first plot
lines = plt.plot(t, u)
# call plt.axis, plt.xlabel, plt.ylabel, etc. as desired

# At later time levels
lines[0].set_ydata(u)
plt.legend('t=%g' % t[n])
plt.draw() # make updated plot
plt.savefig(...)
```

An alternative is to rebuild the plot at every time level:

```
plt.clf() # delete any previous curve(s)
plt.axis([...])
plt.plot(t, u)
# plt.xlabel, plt.legend and other decorations
plt.draw()
plt.savefig(...)
```

Many prefer to work with figure and axis objects as in MATLAB:

```
fig = plt.figure()
...
fig.clf()
ax = fig.gca()
ax.axis(...)
ax.plot(t, u)
# ax.set_xlabel, ax.legend and other decorations
plt.draw()
fig.savefig(...)
```

Running a case. The first demo of our 1D wave equation solver concerns vibrations of a string that is initially deformed to a triangular shape, like when picking a guitar string:

$$I(x) = \begin{cases} ax/x_0, & x < x_0, \\ a(L-x)/(L-x_0), & \text{otherwise} \end{cases} \quad (28)$$

We choose $L = 40$ cm, $x_0 = 0.8L$, $a = 5$ mm, $N_x = 50$, and a time frequency $\nu = 440$ Hz. The relation between the wave speed c and ν is $c = \nu\lambda$, where λ is the wavelength, taken as $2L$ because the longest wave on the string form half a wavelength. There is no external force, so $f = 0$, and the string is at rest initially so that $V = 0$. A function setting these physical parameters and calling `viz` for this case goes as follows:

```
def guitar(C):
    """Triangular wave (pulled guitar string)."""
    L = 0.4
    x0 = 0.8*L
    a = 0.005
    freq = 440
    wavelength = 2*L
    c = freq*wavelength
    omega = 2*pi*freq
    num_periods = 1
    T = 2*pi/omega*num_periods
    Nx = 50

    def I(x):
        return a*x/x0 if x < x0 else a/(L-x0)*(L-x)

    umin = -1.2*a; umax = -umin
    cpu = viz(I, 0, 0, c, L, Nx, C, T, umin, umax, animate=True)
```

The associated program has the name `wave1D_u0_s.py`. Run the program and watch the movie of the vibrating string².

The benefits of scaling. The previous example demonstrated that quite some work is needed with establishing relevant physical parameters for a case.

²http://hplgit.github.com/INF5620/doc/notes/mov-wave/guitar_C0.8/index.html

By *scaling* the mathematical problem we can often reduce the need to estimate physical parameters dramatically. A scaling consists of introducing new independent and dependent variables, with the aim that the absolute value of these vary between 0 and 1:

$$\bar{x} = \frac{x}{L}, \quad \bar{t} = \frac{L}{c}t, \quad \bar{u} = \frac{u}{a}.$$

Replacing old by new variables in the PDE, using $f = 0$, and dropping the bars, results in $u_{tt} = u_{xx}$. That is, the original equation with $c = 1$. The initial condition corresponds to (28) with $a = 1$, $L = 1$, and $x_0 \in [0, 1]$. This means that we only need to decide on x_0 , because the scaled problem corresponds to setting all other parameters to unity! In the code we can just set `a=c=L=1`, `x0=0.8`, and there is no need to calculate with wavelengths and frequencies to estimate c .

The only non-trivial parameter to estimate is T and how it relates to periods in periodic solutions. If $u \sim \sin(\omega t) = \sin(\omega \bar{t} L / c)$ in time, $\omega = 2\pi c / \lambda$, where $\lambda = 2L$ is the wavelength. One period T in dimensionless time means $\omega T L / c = 2\pi$, which gives $T = 2$.

3 Vectorization

The computational algorithm for solving the wave equation visits one mesh point at a time and evaluates a formula for the new value at that point. Technically, this is implemented by a loop over array elements in a program. Such loops may run slowly in Python (and similar interpreted languages such as R and MATLAB). One technique for speeding up loops is to perform operations on entire arrays instead of working with one element at a time. This is referred to as *vectorization*, *vector computing*, or *array computing*. Operations on whole arrays are possible if the computations involving each element is independent of each other and therefore can, at least in principle, be performed simultaneously. Vectorization not only speeds up the code on serial computers, but it also makes it easy to exploit parallel computing.

3.1 Operations on slices of arrays

Efficient computing with `numpy` arrays demands that we avoid loops and compute with entire arrays at once, or at least large portions of them. Consider this calculation of differences:

```
n = u.size
for i in range(0, n-1):
    d[i] = u[i+1] - u[i]
```

All the differences here are independent of each other. The computation of `d` can therefore alternatively be done by subtracting the array $(u_0, u_1, \dots, u_{n-1})$ from the array where the elements are shifted one index upwards: (u_1, u_2, \dots, u_n) ,

see Figure 3. The former subset of the array can be expressed by `u[0:n-1]`, `u[0:-1]`, or just `u[:-1]`, meaning from index 0 up to, but not including, the last element (-1). The latter subset is obtained by `u[1:n]` or `u[1:]`, meaning from index 1 and the rest of the array. The computation of `d` can now be done without an explicit Python loop:

```
d = u[1:] - u[:-1]
```

or with explicit limits if desired:

```
d = u[1:n] - u[0:n-1]
```

Indices with a colon, going from an index to (but not including) another index are called *slices*. With `numpy` arrays, the computations are still done by loops, but in efficient, compiled, highly optimized code in C or Fortran. Such array operations can also easily be distributed among many processors on parallel computers. We say that the *scalar code* above, working on an element (a scalar) at a time, has been replaced by an equivalent *vectorized code*. The process of vectorizing code is called *vectorization*.




Figure 3: Illustration of subtracting two (displaced) slices of two arrays.

Newcomers to vectorization are encouraged to choose a small array `u`, say with five elements, and simulate with pen and paper both the loop version and the vectorized version.

Finite difference schemes basically contains differences between array elements with shifted indices. Consider the updating formula

```
for i in range(1, n-1):
    u2[i] = u[i-1] - 2*u[i] + u[i+1]
```

The vectorization consists of replacing the loop by arithmetics on slices of arrays of length `n-2`:

```
u2 = u[:-2] - 2*u[1:-1] + u[1:]
u2 = u[0:n-2] - 2*u[1:n-1] + u[1:n]    # alternative
```

Note that `u2` here gets length `n-2`. If `u2` is already an array of length `n` and we want to use the formula to update all the "inner" elements of `u2`, as we will when solving a 1D wave equation, we can write

```
u2[1:-1] = u[:-2] - 2*u[1:-1] + u[1:]
u2[1:n-1] = u[0:n-2] - 2*u[1:n-1] + u[1:n] # alternative
```

Pen and paper calculations with a small array will demonstrate what is actually going on. The expression on the right-hand side are done in the following steps, involving temporary arrays with intermediate results, since we can only work with two arrays at a time in arithmetic expressions:

```
temp1 = 2*u[1:-1]
temp2 = u[0:-2] - temp1
temp3 = temp2 + u[1:]
u2[1:-1] = temp3
```

We can extend the example to a formula with an additional term computed by calling a function:

```
def f(x):
    return x**2 + 1

for i in range(1, n-1):
    u2[i] = u[i-1] - 2*u[i] + u[i+1] + f(x[i])
```

Assuming `u2`, `u`, and `x` all have length `n`, the vectorized version becomes

```
u2[1:-1] = u[:-2] - 2*u[1:-1] + u[1:] + f(x[1:-1])
```

3.2 Finite difference schemes expressed as slices

We now have the necessary tools to vectorize the algorithm for the wave equation. There are three loops: one for the initial condition, one for the first time step, and finally the loop that is repeated for all subsequent time levels. Since only the latter is repeated a potentially large number of times, we limit the efforts of vectorizing the code to this loop:

```
for i in range(1, Nx):
    u[i] = 2*u_1[i] - u_2[i] + \
        C2*(u_1[i-1] - 2*u_1[i] + u_1[i+1])
```

The vectorized version becomes

```
u[1:-1] = - u_2[1:-1] + 2*u_1[1:-1] + \
    C2*(u_1[:-2] - 2*u_1[1:-1] + u_1[2:])
```

or

```
u[1:Nx] = 2*u_1[1:Nx]- u_2[1:Nx] + \
          C2*(u_1[0:Nx-1] - 2*u_1[1:Nx] + u_1[2:Nx+1])
```

The program `wave1D_u0_sv.py` contains a new version of the function `solver` where both the scalar and the vectorized loops are included.

Verification. We may reuse the quadratic solution $u_e(x, t) = x(L-x)(1 + \frac{1}{2}t)$ for verifying also the vectorized code. A nose test can now test both the scalar and the vectorized version. Moreover, we may use a `user_action` function that compares the computed and exact solution at each time level and performs an assert:

```
def test_quadratic():
    """
    Check the scalar and vectorized versions work for
    a quadratic u(x,t)=x(L-x)(1+t) that is exactly reproduced.
    """
    exact_solution = lambda x, t: x*(L - x)*(1 + 0.5*t)
    I = lambda x: exact_solution(x, 0)
    V = lambda x: 0.5*exact_solution(x, 0)
    f = lambda x, t: 2*c**2*(1 + 0.5*t)
    L = 2.5
    c = 1.5
    Nx = 3 # very coarse mesh
    C = 1
    T = 18 # long time integration

    def assert_no_error(u, x, t, n):
        u_e = exact_solution(x, t[n])
        diff = abs(u - u_e).max()
        print diff
        nt.assert_almost_equal(diff, 0, places=13)

    solver(I, V, f, c, L, Nx, C, T,
          user_action=assert_no_error, version='scalar')
    solver(I, V, f, c, L, Nx, C, T,
          user_action=assert_no_error, version='vectorized')
```

Here, we also used the opportunity to demonstrate how to achieve very compact code with the use of lambda functions for the various input parameters that require a Python function.

Efficiency measurements. Running the `wave1D_u0_sv.py` code with the previous string vibration example for $N_x = 50, 100, 200, 400, 800$ and measuring the CPU time (see the `run_efficiency_experiments` function), shows that the speed-up of vectorization goes approximately like $5/N_x$, which is a substantial effect!

4 Exercises

Exercise 1: Add storage of solution in a user action function

Extend the `plot_u` function in the file `wave1D_u0_s.py` to also store the solutions `u` in a list. To this end, declare `all_u` as an empty list in the `viz` function, outside `plot_u`, and perform an append operation inside the `plot_u` function. Note that a function, like `plot_u`, inside another function, like `viz`, remembers all local variables in `viz` function, including `all_u`, even when `plot_u` is called (`user_action`) in the `solver` function. Test both `all_u.append(u)` and `all_u.append(u.copy())`. Why does one of these constructions fail to store the solution correctly? Let the `viz` function return the `all_u` list converted to a two-dimensional `numpy` array. Filename: `wave1D_u0_s2.py`.

Exercise 2: Use a class for the user action function

Redo Exercise 1 using a class for the user action function. That is, define a class `Action` where the `all_u` list is an attribute, and implement the user action function as a method (the special method `__call__` is a natural choice). The class versions avoids that the user action function depends on parameters defined outside the function (such as `all_u` in Exercise 1). Filename: `wave1D_u0_s2c.py`.

Exercise 3: Compare several Courant numbers in one movie

Use the program from Exercise 1 or 2 to compute the and store the solutions corresponding to different Courant numbers, say 1.0, 0.9, and 0.1. Make visualization where the three solutions are compared. That is, each frame in the animation shows three curves. The challenge in such a visualization is to ensure that the curves in one plot corresponds to the same time point. Use slicing of the array returned from the `viz` function to pick out the solution at the right time levels (provided that the Courant numbers are integer factors of the smallest Courant number such that the slicing works).

5 Generalization: reflecting boundaries

Now we shall generalize the boundary condition $u = 0$ from Section 1 to the condition $u_x = 0$, which is more complicated to express numerically and also implement.

5.1 Neumann boundary condition

When a wave hits a boundary and is to be reflected back, one applies the condition

$$\frac{\partial u}{\partial n} \equiv \mathbf{n} \cdot \nabla u = 0. \quad (29)$$

The derivative $\partial/\partial n$ is in the outward normal direction from a general boundary. For a 1D domain $[0, L]$, we have that $\partial/\partial n = \partial/\partial x$ at $x = L$ and $\partial/\partial n = -\partial/\partial x$ at $x = 0$. Boundary conditions specifying the value of $\partial u/\partial n$ are known as *Neumann conditions*, while *Dirichlet conditions* refer to specifications of u . When the values are zero ($\partial u/\partial n = 0$ or $u = 0$) we speak about *homogeneous* Neumann or Dirichlet conditions.

5.2 Discretization of derivatives at the boundary

How can we incorporate the condition (29) in the finite difference scheme? Since we have used central differences in all the other approximations to derivatives in the scheme, it is tempting to implement (29) at $x = 0$ and $t = t_n$ by the difference

$$\frac{u_{-1}^n - u_1^n}{2\Delta x} = 0. \quad (30)$$

The problem is that u_{-1}^n is not a u value that is being computed since the point is outside the mesh. However, if we combine (30) with the scheme for $i = 0$,

$$u_i^{n+1} = -u_i^{n-1} + 2u_i^n + C^2 (u_{i+1}^n - 2u_i^n + u_{i-1}^n), \quad (31)$$

we can eliminate the fictitious value u_{-1}^n . We see that $u_{-1}^n = u_1^n$ from (30), which can be used in (31) to arrive at a modified scheme for the boundary point u_0^{n+1} :

$$u_i^{n+1} = -u_i^{n-1} + 2u_i^n + 2C^2 (u_{i+1}^n - u_i^n), \quad i = 0. \quad (32)$$

Figure 4 illustrates how this equation looks like for computing u_0^3 in terms of u_0^2 , u_0^1 , and u_1^2 .

Similarly, (29) applied at $x = L$ is discretized by a central difference

$$\frac{u_{N_x+1}^n - u_{N_x-1}^n}{2\Delta x} = 0. \quad (33)$$

Combined with the scheme for $i = N_x$ we get a modified scheme for the boundary value $u_{N_x}^{n+1}$:

$$u_i^{n+1} = -u_i^{n-1} + 2u_i^n + 2C^2 (u_{i-1}^n - u_i^n), \quad i = N_x. \quad (34)$$

The modification of the scheme at the boundary is also required for the special formula for the first time step. How the stencil moves through the mesh and is modified at the boundary can be illustrated by an animation³

³http://hplgit.github.com/INF5620/doc/notes/mov-wave/wave1D_PDE_Neumann_stencil_gpl/index.html




Figure 4: Modified stencil at a boundary with a Neumann condition.

5.3 Implementation of Neumann conditions

The implementation of the special formulas for the boundary points can benefit from using the general formula for the interior points, but replacing u_{i-1}^n by u_{i+1}^n for $i = 0$ and u_{i+1}^n by u_{i-1}^n for $i = N_x$. This is achieved by just replacing the index $i - 1$ by $i + 1$ for $i = 0$ and $i + 1$ by $i - 1$ for $i = N_x$. In a program, we introduce variables to hold the value of the offset indices: `im1` for $i-1$ and `ip1` for $i+1$. It is now just a manner of defining `im1` and `ip1` properly for the internal points and the boundary points. The coding for the latter reads

```
i = 0
ip1 = i+1
im1 = ip1 # i-1 -> i+1
u[i] = u_1[i] + C2*(u_1[im1] - 2*u_1[i] + u_1[ip1])

i = Nx
im1 = i-1
ip1 = im1 # i+1 -> i-1
u[i] = u_1[i] + C2*(u_1[im1] - 2*u_1[i] + u_1[ip1])
```

We can in fact create one loop over both the internal and boundary points and use only one updating formula:

```
for i in range(0, Nx+1):
    ip1 = i+1 if i < Nx else i-1
    im1 = i-1 if i > 0 else i+1
    u[i] = u_1[i] + C2*(u_1[im1] - 2*u_1[i] + u_1[ip1])
```

The program `wave1D_dn.py` solves the 1D wave equation $u_{tt} = c^2 u_{xx} + f(x, t)$ with general boundary and initial conditions:

- $x = 0$: $u = U_0(t)$ or $u_x = 0$
- $x = L$: $u = U_L(t)$ or $u_x = 0$
- $t = 0$: $u = I(x)$
- $t = 0$: $u_t = I(x)$

5.4 Alternative implementation via ghost cells

Idea. Instead of modifying the scheme at the boundary, we can introduce extra points outside the mesh such that u_{-1}^n and $u_{N_x+1}^n$ are defined. Adding the intervals $[-\Delta x, 0]$ and $[L, L + \Delta x]$, often referred to as "ghost cells", to the mesh gives us all the needed mesh points $i = -1, 0, \dots, N_x, N_x + 1$. If we ensure that we always have

$$u_{-1}^n = u_1^n \text{ and } u_{N_x-1}^n = u_{N_x+1}^n,$$

the application of the standard scheme at a boundary point $i = 0$ or $i = N_x$ will be correct and ensure the solution is compatible with the boundary condition.

Implementation. We may keep `x` as the array for the physical mesh points, but just add extra elements to `u`,

```
u = zeros(Nx+3)
```

and similar for `u_1` and `u_2`.

A major indexing problem arises as Python indices *must* start at 0 (`u[-1]` will always mean the last element in `u`). This implies that we in the mathematics should write the scheme as

$$u_i^{n+1} = \dots, \quad i = 1, \dots, N_x + 1,$$

when i runs through all points in the physical mesh, and $i = 0$ corresponds to the fictitious point $x_{-1} = -\Delta x$. The old code for updating `u` at the inner mesh points must now be extended to cover all points:

```
for i in range(1, Nx+2):
    u[i] = - u_2[i] + 2*u_1[i] + \
        C2*(0.5*(q[i] + q[i+1])*(u_1[i+1] - u_1[i]) - \
            0.5*(q[i] + q[i-1])*(u_1[i] - u_1[i-1])) + \
        dt2*f(x[i], t[n])
```

Then the boundary points are updated, assuming the Neumann condition applies at both boundaries:

```

u[0] = u[2]
u[Nx+2] = u[Nx]

```

The ghost cell is only added to the boundary where we have a Neumann condition.

The solution returned to the user should preferably be an array with values corresponding to the physical mesh points in \mathbf{x} , meaning that we return `u[1:-1]` (if there are two ghost cells).

6 Generalization: variable wave velocity

Our next generalization of the 1D wave equation (1) or (18) is to allow for a variable wave velocity c : $c = c(x)$, usually motivated by wave motion in a domain composed of different physical media with different properties for propagating waves.

6.1 The model PDE with a variable coefficient

Instead of working with the squared quantity $c^2(x)$ we shall for notational convenience introduce $q(x) = c^2(x)$. A 1D wave equation with variable wave velocity often takes the form

$$\frac{\partial^2 u}{\partial t^2} = \frac{\partial}{\partial x} \left(q(x) \frac{\partial u}{\partial x} \right) + f(x, t). \quad (35)$$

This equation sampled at a mesh point (x_i, t_n) reads

$$\frac{\partial^2}{\partial t^2} u(x_i, t_n) = \frac{\partial}{\partial x} \left(q(x_i) \frac{\partial}{\partial x} u(x_i, t_n) \right) + f(x_i, t_n),$$

where the only new term is

$$\frac{\partial}{\partial x} \left(q(x_i) \frac{\partial}{\partial x} u(x_i, t_n) \right) = \left[\frac{\partial}{\partial x} \left(q(x) \frac{\partial u}{\partial x} \right) \right]_i^n.$$

6.2 Discretizing the variable coefficient

The principal idea is to first discretize the outer derivative. Define

$$\phi = q(x) \frac{\partial u}{\partial x},$$

and use a centered derivative around $x = x_i$ for the derivative of ϕ :

$$\left[\frac{\partial \phi}{\partial x} \right]_i^n \approx \frac{\phi_{i+\frac{1}{2}} - \phi_{i-\frac{1}{2}}}{\Delta x} = [D_x \phi]_i^n.$$

Then discretize

$$\phi_{i+\frac{1}{2}} = q_{i+\frac{1}{2}} \left[\frac{\partial u}{\partial x} \right]_{i+\frac{1}{2}}^n \approx q_{i+\frac{1}{2}} \frac{u_{i+1}^n - u_i^n}{\Delta x} = [q D_x u]_{i+\frac{1}{2}}^n.$$

Similarly,

$$\phi_{i-\frac{1}{2}} = q_{i-\frac{1}{2}} \left[\frac{\partial u}{\partial x} \right]_{i-\frac{1}{2}}^n \approx q_{i-\frac{1}{2}} \frac{u_i^n - u_{i-1}^n}{\Delta x} = [q D_x u]_{i-\frac{1}{2}}^n.$$

These intermediate results are now combined to

$$\left[\frac{\partial}{\partial x} \left(q(x) \frac{\partial u}{\partial x} \right) \right]_i^n \approx \frac{1}{\Delta x^2} \left(q_{i+\frac{1}{2}} (u_{i+1}^n - u_i^n) - q_{i-\frac{1}{2}} (u_i^n - u_{i-1}^n) \right). \quad (36)$$

With operator notation we can write the discretization as

$$\left[\frac{\partial}{\partial x} \left(q(x) \frac{\partial u}{\partial x} \right) \right]_i^n \approx [D_x q D_x u]_i^n. \quad (37)$$

Remark. Many are tempted to use the chain rule on the term $\frac{\partial}{\partial x} (q(x) \frac{\partial u}{\partial x})$, but this is not a good idea when discretizing such a term.

6.3 Computing the coefficient between mesh points

If q is a known function of x , we can easily evaluate $q_{i+\frac{1}{2}}$ simply as $q(x_{i+\frac{1}{2}})$ with $x_{i+\frac{1}{2}} = x_i + \frac{1}{2}\Delta x$. However, in many cases c , and hence q , is only known as a discrete function, often at the mesh points x_i . Evaluating q between two mesh points x_i and x_{i+1} can be done by averaging in three ways:

$$q_{i+\frac{1}{2}} \approx \frac{1}{2} (q_i + q_{i+1}) = [\bar{q}^x]_i, \quad (\text{arithmetic mean}) \quad (38)$$

$$q_{i+\frac{1}{2}} \approx 2 \left(\frac{1}{q_i} + \frac{1}{q_{i+1}} \right)^{-1}, \quad (\text{harmonic mean}) \quad (39)$$

$$q_{i+\frac{1}{2}} \approx (q_i q_{i+1})^{1/2}, \quad (\text{geometric mean}) \quad (40)$$

The arithmetic mean in (38) is by far the most commonly used averaging technique.

With the operator notation from (38) we can specify the discretization of the complete variable-coefficient wave equation in a compact way:

$$[D_t D_t u = D_x \bar{q}^x D_x u + f]_i^n. \quad (41)$$

From this notation we immediately see what kind of differences that each term is approximated with. The notation \bar{q}^x also specifies that the variable coefficient is approximated by an arithmetic mean. With the notation $[D_x q D_x u]_i^n$, we specify

that q is evaluated directly, as a function, between the mesh points: $q(x_{i-\frac{1}{2}})$ and $q(x_{i+\frac{1}{2}})$.

Before any implementation, it remains to solve (41) with respect to u_i^{n+1} :

$$u_i^{n+1} = -u_i^{n-1} + 2u_i^n + \left(\frac{\Delta x}{\Delta t}\right)^2 \left(\frac{1}{2}(q_i + q_{i+1})(u_{i+1}^n - u_i^n) - \frac{1}{2}(q_i + q_{i-1})(u_i^n - u_{i-1}^n)\right) + \Delta t^2 f_i^n. \quad (42)$$

6.4 How a variable coefficient affects the stability

The stability criterion derived in Section 11.2 reads $\Delta t \leq \Delta x/c$. If $c = c(x)$, the criterion will depend on the spatial location. We must therefore choose a Δt that is small enough such that no mesh cell has $\Delta x/c(x) > \Delta t$. That is, we must use the largest c value in the criterion:

$$\Delta t \leq \beta \frac{\Delta x}{\max_{x \in [0, L]} c(x)}. \quad (43)$$

The parameter β is included as a safety factor: in some problems with a significantly varying c it turns out that one must choose $\beta < 1$ to have stable solutions ($\beta = 0.9$ may act as an all-round value).

6.5 Implementation of variable coefficients

The implementation of the scheme with a variable wave velocity may assume that c is available as an array $c[i]$ at the mesh points. The following loop is a straightforward implementation of the scheme (42):

```
for i in range(1, Nx):
    u[i] = -u_2[i] + 2*u_1[i] + \
        C2*(0.5*(q[i] + q[i+1])*(u_1[i+1] - u_1[i]) - \
            0.5*(q[i] + q[i-1])*(u_1[i] - u_1[i-1])) + \
        dt2*f(x[i], t[n])
```

The coefficient $C2$ is now defined as $(dt/dx)**2$ and *not* as the squared Courant number since the wave velocity is variable and appears inside the parenthesis.

With Neumann conditions $\partial u / \partial x = 0$ at the boundary, we need to combine this scheme with the discrete version of the boundary condition, $u_{-1}^{n+1} = u_1^{n+1}$:

```
i = 0
ip1 = i+1
im1 = ip1
u[i] = -u_2[i] + 2*u_1[i] + \
    C2*(0.5*(q[i] + q[ip1])*(u_1[ip1] - u_1[i]) - \
        0.5*(q[i] + q[im1])*(u_1[i] - u_1[im1])) + \
    dt2*f(x[i], t[n])
```

A vectorized version of the scheme with a variable coefficient at internal points in the mesh becomes

```
u[1:-1] = - u_2[1:-1] + 2*u_1[1:-1] + \
    C2*(0.5*(q[1:-1] + q[2:])*(u_1[2:] - u_1[1:-1]) -
    0.5*(q[1:-1] + q[:-2])*(u_1[1:-1] - u_1[:-2])) + \
    dt2*f(x[1:-1], t[n])
```

6.6 A more general model PDE with variable coefficients

Sometimes a wave PDE has a variable coefficient also in front of the time-derivative term:

$$\varrho(x) \frac{\partial^2 u}{\partial t^2} = \frac{\partial}{\partial x} \left(q(x) \frac{\partial u}{\partial x} \right) + f(x, t). \quad (44)$$

A natural scheme is

$$[\varrho D_t D_t u = D_x \bar{q}^x D_x u + f]_i^n. \quad (45)$$

We realize that the ϱ coefficient poses no particular difficulty because the only value ϱ_i^n enters the formula above (when written out). There is hence no need for any averaging of ϱ . Often, ϱ will be moved to the right-hand side, also without any difficulty:

$$[D_t D_t u = \varrho^{-1} D_x \bar{q}^x D_x u + f]_i^n. \quad (46)$$

6.7 Generalization: including damping

Waves die out by two mechanisms. In 2D and 3D the energy of the wave spreads out in space, and energy conservation then requires the amplitude to decrease. This effect is not present in 1D. The other cause of amplitude reduction is by damping. For example, the vibrations of a string die out because of air resistance and non-elastic effects in the string.

The simplest way of including damping is to add a first-order derivative to the equation (in the same way as friction forces enter a vibrating mechanical system):

$$\frac{\partial^2 u}{\partial t^2} + b \frac{\partial u}{\partial t} = c^2 \frac{\partial^2 u}{\partial x^2} + f(x, t), \quad (47)$$

where $b \geq 0$ is a prescribed damping coefficient.

A typical discretization of (47) in terms of centered differences reads

$$[D_t D_t u + b D_{2t} u = c^2 D_x D_x u + f]_i^n. \quad (48)$$

Writing out the equation and solving for the unknown u_i^{n+1} gives the scheme

$$u_i^{n+1} = (1 + \frac{1}{2}b\Delta t)^{-1}((\frac{1}{2}b\Delta t - 1)u_i^{n-1} + 2u_i^n + C^2(u_{i+1}^n - 2u_i^n + u_{i-1}^n) + \Delta t^2 f_i^n), \quad (49)$$

New equations must be derived for u_i^1 , and for boundary points in case of Neumann conditions.

The damping is very small in many wave phenomena and then only evident for very long time simulations, so it is common to drop the bu_t term in the wave equation.

7 Building a general 1D wave equation solver

The program `wave1D_dn_vc.py` is a fairly general code for 1D wave propagation problems that targets the following initial-boundary value problem

$$u_t = (c^2(x)u_x)_x + f(x, t), \quad x \in (0, L), \quad t \in (0, T], \quad (50)$$

$$u(x, 0) = I(x), \quad x \in [0, L], \quad (51)$$

$$u_t(x, 0) = V(t), \quad x \in [0, L], \quad (52)$$

$$u(0, t) = U_0(t) \text{ or } u_x(0, t) = 0, \quad t \in (0, T], \quad (53)$$

$$u(L, t) = U_L(t) \text{ or } u_x(L, t) = 0, \quad t \in (0, T] \quad (54)$$

The `solver` function is a natural extension of the simplest `solver` function in the initial `wave1D_u0_s.py` program, extended with Neumann boundary conditions ($u_x = 0$), a possibly time-varying boundary condition on u ($U_0(t)$, $U_L(t)$), and a variable wave velocity. The different code segments needed to make these extensions are shown and commented upon in the preceding text.

The vectorization is only applied inside the time loop, not for the initial condition or the first time steps, since this initial work is negligible for long time simulations in 1D problems.

7.1 User action function as a class

A useful feature in the `wave1D_dn_vc.py` program is the specification of the `user_action` function as a class. Although the `plot_u` function in the `viz` function of previous `wave1D*.py` programs remembers the local variables in the `viz` function, it is a cleaner solution to store the needed variables together with the function, which is exactly what a class offers.

A class for flexible plotting, cleaning up files, and making a movie files like function `viz` and `plot_u` did can be coded as follows:

```
class PlotSolution:
    """
    Class for the user_action function in solver.
    Visualizes the solution only.
    """
```

```

def __init__(self,
              casename='tmp',      # prefix in filenames
              umin=-1, umax=1,    # fixed range of y axis
              pause_between_frames=None, # movie speed
              backend='matplotlib', # or 'gnuplot'
              screen_movie=True, # show movie on screen?
              every_frame=1):      # show every_frame frame
    self.casename = casename
    self.yaxis = [umin, umax]
    self.pause = pause_between_frames
    module = 'scitools.easyviz.' + backend + '_'
    exec('import %s as st' % module)
    self.st = st
    self.screen_movie = screen_movie
    self.every_frame = every_frame

    # Clean up old movie frames
    for filename in glob('frame_*.png'):
        os.remove(filename)

def __call__(self, u, x, t, n):
    if n % self.every_frame != 0:
        return
    self.st.plot(x, u, 'r-',
                 xlabel='x', ylabel='u',
                 axis=[x[0], x[-1],
                      self.yaxis[0], self.yaxis[1]],
                 title='t=%f' % t[n],
                 show=self.screen_movie)

    # pause
    if t[n] == 0:
        time.sleep(2) # let initial condition stay 2 s
    else:
        if self.pause is None:
            pause = 0.2 if u.size < 100 else 0
        time.sleep(pause)

    self.st.savefig('%s_frame_%04d.png' % (self.casename, n))

```

Understanding this class requires quite some familiarity with Python in general and class programming in particular.

The constructor shows how we can flexibly import the plotting engine as (typically) `scitools.easyviz.gnuplot_` or `scitools.easyviz.matplotlib_` (the trailing underscore is SciTool's way of avoiding name clash between its interface modules and plotting packages). With the `screen_movie` parameter we can suppress displaying each movie frame on the screen (`show=False` parameter to `plot`). Alternatively, for slow movies associated with fine meshes, one can set `every_frame` to, e.g., 10, causing every 10 frames to be shown.

The `__call__` method makes `PlotSolution` instances behave like functions, so we can just pass an instance, say `p`, as the `user_action` argument in the `solver` function, and any call to `user_action` will be a call to `p.__call__`.

7.2 Collection of initial conditions

The function `pulse` in `wave1D_dn_vc.py` demonstrates wave motion in heterogeneous media where c varies. One can specify an interval where the wave velocity is decreased by a factor `slowness_factor` (or increased by making this factor less than one). Four types of initial conditions are available: a square pulse (`plug`), a Gaussian function (`gaussian`), a cosine "hat" consisting of one period of the cosine function (`cosinehat`), and half a period of a cosine "hat" (`half-cosinehat`). These peak-shaped initial conditions can be placed in the middle (`loc='center'`) or at the left end (`loc='left'`) of the domain. The `pulse` function is a flexible tool for playing around with various wave shapes and location of a medium with a different wave velocity:

```
def pulse(C=1, Nx=200, animate=True, version='vectorized', T=2,
         loc='center', pulse_tp='gaussian', slowness_factor=2,
         medium=[0.7, 0.9], every_frame=1):
    """
    Various peaked-shaped initial conditions on [0,1].
    Wave velocity is decreased by the slowness_factor inside
    the medium specification. The loc parameter can be 'center'
    or 'left', depending on where the pulse is to be located.
    """
    L = 1.
    if loc == 'center':
        xc = L/2
    elif loc == 'left':
        xc = 0
    sigma = L/20. # width measure of I(x)
    c_0 = 1.0     # wave velocity outside medium

    if pulse_tp in ('gaussian', 'Gaussian'):
        def I(x):
            return exp(-0.5*((x-xc)/sigma)**2)
    elif pulse_tp == 'plug':
        def I(x):
            return 0 if abs(x-xc) > sigma else 1
    elif pulse_tp == 'cosinehat':
        def I(x):
            # One period of a cosine
            w = 2
            a = w*sigma
            return 0.5*(1 + cos(pi*(x-xc)/a)) \
                if xc - a <= x <= xc + a else 0

    elif pulse_tp == 'half-cosinehat':
        def I(x):
            # Half a period of a cosine
            w = 4
            a = w*sigma
            return cos(pi*(x-xc)/a) \
                if xc - 0.5*a <= x <= xc + 0.5*a else 0
    else:
        raise ValueError('Wrong pulse_tp="%s"' % pulse_tp)

    def c(x):
        return c_0/slowness_factor \
            if medium[0] <= x <= medium[1] else c_0
```

```

umin=-0.5; umax=1.5*I(xc)
casename = '%s_Nx%s_sf%s' % \
            (pulse_tp, Nx, slowness_factor)
action = PlotMediumAndSolution(
    medium, casename=casename, umin=umin, umax=umax,
    every_frame=every_frame, screen_movie=animate)

solver(I=I, V=None, f=None, c=c, U_0=None, U_L=None,
       L=L, Nx=Nx, C=C, T=T,
       user_action=action, version=version,
       dt_safety_factor=1)

```

The `PlotMediumAndSolution` class used here is a subclass of `PlotSolution` where the medium, specified by the `medium` interval, is indicated in the plots.

7.3 Calling functions from the command line

Experimenting with the `pulse` function, as suggested in Exercises 8, 9, and 10 can easily be done in a little program that imports the `wave1D_dn_vc` module and calls `pulse` with appropriate parameters. However, sometimes it is handy to perform such calls directly on the command line. A useful function `function_UI` from the `scitools.misc` module takes a list of functions in your program, along with `sys.argv`, and returns the right call to a function based on the command-line arguments. That is, by writing

```

Terminal> python progname.py funcname \
          arg1 arg2 ..., kwarg1=v1 kwarg2=v2

```

the `function_UI` function constructs string

```
'funcname(arg1, arg2, ..., kwarg1=v1, kwarg2=v2'
```

that you can send to `eval` to realize the call. For example, the code

```

from scitools.misc import function_UI
cmd = function_UI([pulse,], sys.argv)
eval(cmd)

```

combined with the command

```

Terminal> python wave1D_dn_vc.py pulse C=1 loc="'left'" \
      "medium=[0.5, 0.88]" slowness_factor=2 Nx=40 \
      pulse_tp="'cosinehat'" T=1 every_frame=10

```

makes `eval(cmd)` perform the call

```
pulse(C=1, loc='left', medium=[0.5, 0.8], slowness_factor=2,
      Nx=40, pulse_tp='cosinehat', T=1, every_frame=10)
```

Writing only the function name and no arguments on the command line makes `function_UI` figure out what the arguments are and write a help string. In short, `scitools.mist.function_UI` makes all your desired functions in a program callable from the command line with two statements. This is especially useful during program testing.

8 Exercises

Exercise 4: Use ghost cells to implement Neumann conditions

Modify the program `wave1D_dn.py` to incorporate the ghost cell technique outlined in Section 5.4. Add no, one, or two ghost cells according to the specified boundary conditions. Filename: `wave1D_dn_ghost.py`.

Exercise 5: Find a symmetry boundary condition

Consider the solution $u(x, t)$ in Exercise 7 that is symmetric around $x = 0$. This means that we can simulate the wave process in only the half of the domain $[0, L]$. What is the correct boundary condition to impose at $x = 0$? Filename: `wave1D_symmetric`.

Exercise 6: Prove symmetry of a 1D wave problem computationally

Perform simulations of the complete wave problem from Exercise 7 on $[-L, L]$. Thereafter, utilize the symmetry of the solution and run a simulation in half of the domain $[0, L]$, using the boundary condition at $x = 0$ as derived in Exercise 5. Compare the two solutions and make sure that they are the same. Filename: `wave1D_symmetric`.

Exercise 7: Prove symmetry of a 1D wave problem analytically

Consider the simple "plug" wave where $\Omega = [-L, L]$ and

$$I(x) = \begin{cases} 1, & x \in [-\delta, \delta], \\ 0, & \text{otherwise} \end{cases}$$

for some number $0 < \delta < L$. The boundary conditions can be set to $u = 0$. The solution to this problem is symmetric around $x = 0$. Prove this by setting up the complete initial-boundary value problem and showing that if $u(x, t)$ is a solution, then also $u(-x, t)$ is a solution. Filename: `wave1D_symmetric.py`.

Exercise 8: Send pulse waves through a layered medium

Use the `pulse` function in `wave1D_dn_vs.py` to investigate sending a pulse, located with its peak at $x = 0$, through the medium to the right where it hits another medium for $x \in [0.7, 0.9]$ where the wave velocity is decreased by a factor s_f . Report what happens with a Gaussian pulse, a "cohat" pulse, half a "cohat" pulse, and a plug pulse for resolutions $N_x = 40, 80, 160$, and $s_f = 2, 4$. Make a reference solution in the homogeneous case too with $s_f = 1$ and $N_x = 40$. Use $C = 1$ in the medium outside $[0.7, 0.9]$. Simulate until $T = 2$. Filename: `pulse1D.py`.

Exercise 9: Explain why numerical noise occurs

The experiments performed in Exercise 8 shows considerable numerical noise in the form of non-physical waves, especially for $s_f = 4$ and the plug pulse or the half a "cohat" pulse. The noise is much less visible for a Gaussian pulse. Run the case with the plug and half a "cohat" pulses for $s_f = 1$, $C = 0.9, 0.25$, and $N_x = 40, 80, 160$. Use the numerical dispersion relation to explain the observations. Filename: `pulse1D_analysis.pdf`.

Exercise 10: Investigate harmonic averaging in a 1D model

Will harmonic averaging of the wave velocity give better numerical results for the case $s_f = 4$ in Exercise 8? Filenames: `pulse1D_harmonic.pdf`, `pulse1D_harmonic.py`.

9 Finite difference methods for 2D and 3D wave equations

A natural next step is to consider extensions of the methods for various variants of the one-dimensional wave equation to two-dimensional (2D) and three-dimensional (3D) versions of the wave equation.

9.1 Multi-dimensional wave equations

The general wave equation in d space dimensions, with constant wave velocity c , can be written in the compact form

$$\frac{\partial^2 u}{\partial t^2} = c^2 \nabla^2 u \text{ for } \mathbf{x} \in \Omega \subset \mathbb{R}^d, \ t \in (0, T]. \quad (55)$$

In a 2D problem, $d = 2$, and

$$\nabla^2 u = \frac{\partial^2 u}{\partial x^2} + \frac{\partial^2 u}{\partial y^2},$$

while in three space dimensions, $d = 3$, and

$$\nabla^2 u = \frac{\partial^2 u}{\partial x^2} + \frac{\partial^2 u}{\partial y^2} + \frac{\partial^2 u}{\partial z^2}.$$

Many applications involve variable coefficients, and the general wave equation in d dimensions is in this case written as

$$\varrho \frac{\partial^2 u}{\partial t^2} = \nabla \cdot (q \nabla u) + f \text{ for } \mathbf{x} \in \Omega \subset \mathbb{R}^d, \ t \in (0, T], \quad (56)$$

which in 2D becomes

$$\varrho(x, y) \frac{\partial^2 u}{\partial t^2} = \frac{\partial}{\partial x} \left(q(x, y) \frac{\partial u}{\partial x} \right) + \frac{\partial}{\partial y} \left(q(x, y) \frac{\partial u}{\partial y} \right) + f(x, y, t). \quad (57)$$

To save some writing and space we may use the index notation, where subscript t , x , y , or z means differentiation with respect to that coordinate. For example,

$$\begin{aligned} \frac{\partial^2 u}{\partial t^2} &= u_{tt}, \\ \frac{\partial}{\partial y} \left(q(x, y) \frac{\partial u}{\partial y} \right) &= (qu_y)_y. \end{aligned}$$

The 3D versions of the two model PDEs, with and without variable coefficients, can with now with the aid of the index notation for differentiation be stated as

$$u_{tt} = c^2(u_{xx} + u_{yy} + u_{zz}) + f, \quad (58)$$

$$\varrho u_{tt} = (qu_x)_x + (qu_z)_z + (qu_z)_z + f. \quad (59)$$

At *each point* of the boundary $\partial\Omega$ of Ω we need *one* boundary condition involving the unknown u . The boundary conditions are of three principal types:

1. u is prescribed (usually for a known incoming wave),
2. $\partial u / \partial n = \mathbf{n} \cdot \nabla u$ prescribed (zero for reflecting boundaries),
3. an open boundary condition (also called radiation condition) is specified to let waves travel undisturbed out of the domain, see Exercise 14 for details.

All the listed wave equations with *second-order* derivatives in time need *two* initial conditions:

1. $u = I$,
2. $\partial u / \partial t = V$.

9.2 Mesh

We introduce a mesh in time and in space. The mesh in time consists of time points $t_0 = 0 < t_1 < \dots < t_N$, often with a constant spacing $\Delta t = t_{n+1} - t_n$, $n = 0, \dots, N-1$.

When using finite difference approximations, the domain shape in space is normally simple. We assume that Ω has the shape of a d -dimensional box shape. Mesh points are introduced separately in the various space directions: $x_0 < x_1 < \dots < x_{N_x}$ in x direction, $y_0 < y_1 < \dots < y_{N_y}$ in y direction, and $z_0 < z_1 < \dots < z_{N_z}$ in z direction. It is a very common choice to use constant mesh spacings: $\Delta x = x_{i+1} - x_i$, $i = 0, \dots, N_x - 1$, $\Delta y = y_{j+1} - y_j$, $j = 0, \dots, N_y - 1$, and $\Delta z = z_{k+1} - z_k$, $k = 0, \dots, N_z - 1$, but often with $\Delta x \neq \Delta y \neq \Delta z$. In case the mesh spacings are equal in the spatial directions, one often introduces the symbol h : $h = \Delta x = \Delta y = \Delta z$.

The unknown u at mesh point (x_i, y_j, z_k, t_n) is denoted $u_{i,j,k}^n$. In 2D problems we just skip the z coordinate (by assuming no variation in that direction: $\partial/\partial z = 0$) and write $u_{i,j}^n$.

9.3 Discretization

Two- and three-dimensional wave equations are easily discretized by assembling building blocks for discretization of 1D wave equations, because the multi-dimensional versions just contain terms of the same type that occur in 1D.

Discretizing the PDEs. For example, (58) can be discretized as

$$[D_t D_t u = c^2(D_x D_x u + D_y D_y u + D_z D_z u) + f]_{i,j,k}^n. \quad (60)$$

A 2D version might be instructive to write out in detail:

$$[D_t D_t u = c^2(D_x D_x u + D_y D_y u) + f]_{i,j,k}^n,$$

which becomes

$$\frac{u_{i,j}^{n+1} - 2u_{i,j}^n + u_{i,j}^{n-1}}{\Delta t^2} = c^2 \frac{u_{i+1,j}^n - 2u_{i,j}^n + u_{i-1,j}^n}{\Delta x^2} + c^2 \frac{u_{i,j+1}^n - 2u_{i,j}^n + u_{i,j-1}^n}{\Delta y^2} + f_{i,j}^n,$$

Assuming as usual that all values at the time levels n and $n-1$ are known, we can solve for the only unknown $u_{i,j}^{n+1}$.

As in the 1D case, we need to develop a special formula for $u_{i,j}^1$ where we combine the general scheme for $u_{i,j}^{n+1}$, when $n = 0$, with the discretization of the initial condition:

$$[D_{2t} u = V]_{i,j}^0 \Rightarrow u_{i,j}^{-1} = u_{i,j}^1 - 2\Delta t V_{i,j}.$$

The PDE (59) with variable coefficients is discretized term by term using the corresponding elements from the 1D case:

$$[\varrho D_t D_t u = (D_x \bar{q}^x D_x u + D_y \bar{q}^y D_y u + D_z \bar{q}^z D_z u) + f]_{i,j,k}^n. \quad (61)$$

When written out and solved for the unknown $u_{i,j,k}^{n+1}$, one gets the scheme

$$\begin{aligned} u_{i,j,k}^{n+1} &= -u_{i,j,k}^{n-1} + 2u_{i,j,k}^n + \\ &= \frac{1}{\varrho_{i,j,k}} \frac{1}{\Delta x^2} \left(\frac{1}{2} (q_{i,j,k} + q_{i+1,j,k}) (u_{i+1,j,k}^n - u_{i,j,k}^n) - \right. \\ &\quad \left. \frac{1}{2} (q_{i-1,j,k} + q_{i,j,k}) (u_{i,j,k}^n - u_{i-1,j,k}^n) \right) + \\ &= \frac{1}{\varrho_{i,j,k}} \frac{1}{\Delta x^2} \left(\frac{1}{2} (q_{i,j,k} + q_{i,j+1,k}) (u_{i,j+1,k}^n - u_{i,j,k}^n) - \right. \\ &\quad \left. \frac{1}{2} (q_{i,j-1,k} + q_{i,j,k}) (u_{i,j,k}^n - u_{i,j-1,k}^n) \right) + \\ &= \frac{1}{\varrho_{i,j,k}} \frac{1}{\Delta x^2} \left(\frac{1}{2} (q_{i,j,k} + q_{i,j,k+1}) (u_{i,j,k+1}^n - u_{i,j,k}^n) - \right. \\ &\quad \left. \frac{1}{2} (q_{i,j,k-1} + q_{i,j,k}) (u_{i,j,k}^n - u_{i,j,k-1}^n) \right) + \\ &\quad + \Delta t^2 f_{i,j,k}^n. \end{aligned}$$

Also here we need to develop a special formula for $u_{i,j,k}^1$ by combining the scheme for $n = 0$ with the discrete initial condition $u_{i,j,k}^{-1} = u_{i,j,k}^1 - 2\Delta t V_{i,j,k}$.

Handling boundary conditions where u is known. The schemes listed above are valid for the internal points in the mesh. After updating these, we need to visit all the mesh points at the boundaries and set the prescribed u value.

Discretizing the $\partial u / \partial n = 0$. The condition $\partial u / \partial n = 0$ was implemented in 1D by discretizing it with a $D_{2t}u$ centered difference, and thereafter eliminating the fictitious u point outside the mesh by using the general scheme at the boundary point. Exactly the same idea is reused in multi dimensions. Consider $\partial u / \partial n = 0$ at a boundary $y = 0$. The normal direction is then in $-y$ direction, so

$$\frac{\partial u}{\partial n} = -\frac{\partial u}{\partial y},$$

and we set

$$[-D_{2y}u = 0]_{i,0}^n \Rightarrow \frac{u_{i,1}^n - u_{i,-1}^n}{2\Delta y} = 0.$$

From this it follows that $u_{i,-1}^n = u_{i,1}^n$. The discretized PDE at the boundary point $(i, 0)$ reads

$$\frac{u_{i,0}^{n+1} - 2u_{i,0}^n + u_{i,0}^{n-1}}{\Delta t^2} = c^2 \frac{u_{i+1,0}^n - 2u_{i,0}^n + u_{i-1,0}^n}{\Delta x^2} + c^2 \frac{u_{i,1}^n - 2u_{i,0}^n + u_{i,-1}^n}{\Delta y^2} + f_{i,j}^n,$$

We can then just insert $u_{i,j}^1$ for $u_{i,-1}^n$ in this equation and then solve for the boundary value $u_{i,0}^{n+1}$.

From these calculations, we see a pattern: the general scheme applies at the boundary $j = 0$ too if we just replace $j - 1$ by $j + 1$. Such a pattern is particularly useful for implementations.

10 Implementation

We shall now describe in detail various Python implementations for solving a standard 2D, linear wave equation with constant wave velocity and $u = 0$ on the boundary. The wave equation is to be solved in the space-time domain $\Omega \times (0, T]$, where $\Omega = [0, L_x] \times [0, L_y]$ is a rectangular spatial domain. More precisely, the complete initial-boundary value problem is defined by

$$u_t = c^2(u_{xx} + u_{yy}) + f(x, y, t), \quad (x, y) \in \Omega, \quad t \in (0, T], \quad (62)$$

$$u(x, y, 0) = I(x, y), \quad (x, y) \in \Omega, \quad (63)$$

$$u_t(x, y, 0) = V(x, y), \quad (x, y) \in \Omega, \quad (64)$$

$$u = 0, \quad (x, y) \in \partial\Omega, \quad t \in (0, T], \quad (65)$$

where $\partial\Omega$ is the boundary of Ω , in this case the four sides of the rectangle $[0, L_x] \times [0, L_y]$: $x = 0$, $x = L_x$, $y = 0$, and $y = L_y$.

The PDE is discretized as

$$[D_t D_t u = c^2(D_x D_x u + D_y D_y u) + f]_{i,j}^n,$$

which leads to an explicit updating formula to be implemented in a program:

$$u^{n+1} = -u_{i,j}^{n-1} + 2u_{i,j}^n + C_x^2(u_{i+1,j}^n - 2u_{i,j}^n + u_{i-1,j}^n) + C_y^2(u_{i,j+1}^n - 2u_{i,j}^n + u_{i,j-1}^n) + \Delta t^2 f_{i,j}^n, \quad (66)$$

for all interior mesh points $i = 0, \dots, N_x - 1$ and $j = 0, \dots, N_y - 1$, and for $n = 1, \dots, N$. The constants C_x and C_y are defined as

$$C_x = c \frac{\Delta t}{\Delta x}, \quad C_y = c \frac{\Delta t}{\Delta y}.$$

At the boundary we simply set $u_{i,j}^{n+1} = 0$ for $i = 0, j = 0, \dots, N_y$; $i = N_x, j = 0, \dots, N_y$; $j = 0, i = 0, \dots, N_x$; and $j = N_y, i = 0, \dots, N_x$. For the first step, $n = 0$, (67) is combined with the discretization of the initial condition $u_t = V$, $[D_{2t} u = V]_{i,j}^0$ to obtain a special formula for $u_{i,j}^1$ at the interior mesh points:

$$u^1 = u_{i,j}^0 + \Delta t V_{i,j} + \frac{1}{2} C_x^2 (u_{i+1,j}^0 - 2u_{i,j}^0 + u_{i-1,j}^0) + \frac{1}{2} C_y^2 (u_{i,j+1}^0 - 2u_{i,j}^0 + u_{i,j-1}^0) + \frac{1}{2} \Delta t^2 f_{i,j}^n, \quad (67)$$

10.1 Scalar computations

The `solver` function for a 2D case with constant wave velocity and $u = 0$ as boundary condition follows the setup from the similar function for the 1D case in `wave1D_u0_s.py`, but there are a few necessary extensions. The code is in the program `wave2D_u0.py`.

Domain and mesh. The spatial domain is now $[0, L_x] \times [0, L_y]$, specified by the arguments `Lx` and `Ly`. Similarly, the number of mesh points in the x and y directions, N_x and N_y , become the arguments `Nx` and `Ny`. In multi-dimensional problems it makes less sense to specify a Courant number as the wave velocity is a vector and the mesh spacings may differ in the various spatial directions. We therefore give Δt explicitly. The signature of the `solver` function is then

```
def solver(I, V, f, c, Lx, Ly, Nx, Ny, dt, T,
          user_action=None, version='scalar',
          dt_safety_factor=1):
```

Key parameters used in the calculations are created as

```
x = linspace(0, Lx, Nx+1) # mesh points in x dir
y = linspace(0, Ly, Ny+1) # mesh points in y dir
dx = x[1] - x[0]
dy = y[1] - y[0]
N = int(round(T/float(dt)))
t = linspace(0, N*dt, N+1) # mesh points in time
Cx2 = (c*dt/dx)**2; Cy2 = (c*dt/dy)**2 # help variables
dt2 = dt**2
```

Stability limit. Specifying a negative `dt` parameter makes us use the stability limit with a safety factor:

```
if dt <= 0:
    stability_limit = (1/float(c))*(1/sqrt(1/dx**2 + 1/dy**2))
    dt = dt_safety_factor*stability_limit
```

Solution arrays. We store $u_{i,j}^{n+1}$, $u_{i,j}^n$, and $u_{i,j}^{n-1}$ in three two-dimensional arrays,

```
u = zeros((Nx+1,Ny+1)) # solution array
u_1 = zeros((Nx+1,Ny+1)) # solution at t-dt
u_2 = zeros((Nx+1,Ny+1)) # solution at t-2*dt
```

where $u_{i,j}^{n+1}$ corresponds to `u[i,j]`, $u_{i,j}^n$ to `u_1[i,j]`, and $u_{i,j}^{n-1}$ to `u_2[i,j]`

Computing the solution. Inserting the initial condition I in `u_1` and making a callback to the user in terms of the `user_action` function should be straightforward generalization of the 1D code:

```
for i in range(0, Nx+1):
    for j in range(0, Ny+1):
        u_1[i,j] = I(x[i], y[j])

if user_action is not None:
    user_action(u_1, x, xv, y, yv, t, 0)
```

The `user_action` function has more arguments which will be commented upon in the section on vectorization.

The key finite difference formula for updating the solution at a time level is implemented as

```
for i in range(1, Nx):
    for j in range(1, Ny):
        u[i,j] = 2*u_1[i,j] - u_2[i,j] + \
            Cx2*(u_1[i-1,j] - 2*u_1[i,j] + u_1[i+1,j]) + \
            Cy2*(u_1[i,j-1] - 2*u_1[i,j] + u_1[i,j+1]) + \
            dt2*f(x[i], y[j], t[n])
```

We must thereafter ensure that u is zero on the four boundaries:

```
j = 0
for i in range(0, Nx+1): u[i,j] = 0
j = Ny
for i in range(0, Nx+1): u[i,j] = 0
i = 0
for j in range(0, Ny+1): u[i,j] = 0
i = Nx
for j in range(0, Ny+1): u[i,j] = 0
```

The special formula for the first step ($u_{i,j}^1$) is implemented in a very similar manner.

The code segments with the loop over internal points and boundary points at each time level are put in a separate function `advance_scalar`. Below, we will make many alternative implementations to speed up the code since most of the CPU time is spent in this function.

10.2 Vectorized computations

The scalar code above turns out to be extremely slow for large 2D meshes, and probably useless beyond initial debugging in 3D. Vectorization is therefore a must for multi-dimensional finite difference computations in Python. For example, with a mesh consisting of 30×30 cells, vectorization brings down the CPU time by a factor of 70 (!).

In the vectorized case we must be able to evaluate user-given functions like $I(x, y)$ and $f(x, y, t)$, provided as Python functions `I(x,y)` and `f(x,y,t)`, for the entire mesh in one array operation. Having the one-dimensional coordinate arrays `x` and `y` is not sufficient: these must be extended to vectorized versions,

```
xv = x[:,newaxis]
yv = y[newaxis,:]
# or
xv = x.reshape((x.size, 1))
yv = y.reshape((1, y.size))
```

This is a standard required technique when evaluating functions over a 2D mesh, say $\sin(xv)*\cos(yv)$, which then gives a results with shape $(Nx+1, Ny+1)$. With the xv and yv arrays for vectorized computing, setting the initial condition is just a matter of

```
u_1[:, :] = I(xv, yv)
```

One could also have written $u_1 = I(xv, yv)$ and let u_1 point to a new object, but vectorized operations often makes use of direct insertion in the original array through $u_1[:, :]$ because sometimes not all of the array is to be filled by such a function evaluation. This is the case with the computational scheme for $u_{i,j}^{n+1}$:

```
u[1:-1,1:-1] = 2*u_1[1:-1,1:-1] - u_2[1:-1,1:-1] + \
    Cx2*(u_1[:-2,1:-1] - 2*u_1[1:-1,1:-1] + u_1[2:,1:-1]) + \
    Cy2*(u_1[1:-1,:-2] - 2*u_1[1:-1,1:-1] + u_1[1:-1,2:]) + \
    dt2*f_a[1:-1,1:-1]
```

Array slices in 2D are more complicated to understand than those in 1D, but the logic from 1D applies to each dimension separately. For example, when doing $u_{i,j}^n - u_{i-1,j}^n$ for $i = 1, \dots, N_x$, we just keep j constant and make a slice in the first index: $u_1[1:, j] - u_1[:-1, j]$, exactly as in 1D. The $1:$ slice specifies all the indices $i = 1, \dots, N_x$ for the term $u_{i,j}^n$, while $:-1$ specifies the relevant indices for the second term, where $i - 1 = 1, \dots, N_x$ and the indices therefore becomes $0, 1, \dots, N_x - 1$. In the above code segment, the situation is slightly more complicated, because each displaced slice in one direction is accompanied by a $1:-1$ slice in the other direction. The reason is that we only work with the internal points for the index that is kept constant in a difference.

The f function is in the above vectorized update of u first computed as an array over all mesh points:

```
f_a = f(xv, yv, t[n])
```

(We could, alternatively, used the call $f(xv, yv, t[n])[1:-1, 1:-1]$ in the last term of the update statement, but other implementations in compiled languages benefit from having f available in an array rather than calling our Python function $f(x, y, t)$ for every point.)

The boundary conditions along the four sides makes use of a slice consisting of all indices along a boundary:

```
u[:, 0] = 0
u[:, Ny] = 0
u[0, :] = 0
u[Nx, :] = 0
```


The shown snippets appear in a function `advance_vectorized`.

The callback function now has the arguments `u`, `x`, `xv`, `y`, `yv`, `t`, `n`. The inclusion of `xv` and `yv` makes it easy to, e.g., compute an exact 2D solution in the callback function and compute errors, through something like `u - exact_solution(xv, yv, t[n])`.

10.3 Verification

Testing a quadratic solution. The 1D solution from Section 18 can be generalized to multi dimensions and provides a test case where the exact solution also fulfills the discrete equations. In 2D we have

$$u_e(x, y, t) = x(L_x - x)y(L_y - y)(1 + \frac{1}{2}t). \quad (68)$$

This solution fulfills the PDE problem if $I(x, y) = u_e(x, y, 0)$, $V = \frac{1}{2}u_e(x, y, 0)$, and $f = 2c^2(1 + \frac{1}{2}t)(y(L_y - y) + x(L_x - x))$. With the results $[D_t D_t t]^n = 0$ and $[D_t D_t t^2] = 2$, and

$$[D_x D_x u_e]_{i,j}^n = [y(L_y - y)(1 + \frac{1}{2}t)D_x D_x x(L_x - x)]_{i,j}^n = y_j(L_y - y_j)(1 + \frac{1}{2}t_n)2,$$

with similar calculations for the other terms, we can easily show that u_e fulfills the discrete equation

$$[D_t D_t u_e = c^2(D_x D_x u_e + D_y D_y u_e + f)]_{i,j}^n.$$

With an exact solution of the discrete equation we should expect a numerical error at the level of the machine precision (as long as $|u_e| \sim \mathcal{O}(1)$). The `test_quadratic` function in the `wave2D_u0.py` program implements this verification as a nose test.

10.4 Migrating loops to Cython

Although vectorization can bring down the CPU time dramatically compared with scalar code, there is still some factor 5-10 to win in these types of applications by implementing the finite difference scheme in compiled code, typically in Fortran, C, or C++. This can quite easily be done by adding a little extra code to our program. Cython is an extension of Python that offers the easiest way to nail our Python loops in the scalar code down to machine code and the efficiency of C.

Cython can be viewed as an extension of Python where variables are declared with types and functions are marked to be implemented in C. Migrating Python code to Cython is done by copying the desired code segments to functions (or classes) and placing them in one or more separate files with extension `.pyx`.

Declaring variables and annotating the code. Our starting point is the plain `advance_scalar` function. We simply take a copy of this code and put it in a file `wave2D_u0_loop_cy.pyx`. The relevant Cython implementation arises from declaring variables with types and adding some important annotations to speed up array computing in Cython. Let us first list the code:

```
import numpy as np
cimport numpy as np
cimport cython
ctypedef np.float64_t DT      # data type

@cython.boundscheck(False) # turn off array bounds check
@cython.wraparound(False)  # turn off negative indices (u[-1,-1])
cpdef advance(
    np.ndarray[DT, ndim=2, mode='c'] u,
    np.ndarray[DT, ndim=2, mode='c'] u_1,
    np.ndarray[DT, ndim=2, mode='c'] u_2,
    np.ndarray[DT, ndim=2, mode='c'] f,
    double Cx2, double Cy2, double dt2):

    cdef int Nx, Ny, i, j
    Nx = u.shape[0]-1
    Ny = u.shape[1]-1
    for i in xrange(1, Nx):
        for j in xrange(1, Ny):
            u[i,j] = 2*u_1[i,j] - u_2[i,j] + \
                Cx2*(u_1[i-1,j] - 2*u_1[i,j] + u_1[i+1,j]) + \
                Cy2*(u_1[i,j-1] - 2*u_1[i,j] + u_1[i,j+1]) + \
                dt2*f[i,j]
    # Boundary condition u=0
    j = 0
    for i in range(0, Nx+1): u[i,j] = 0
    j = Ny
    for i in range(0, Nx+1): u[i,j] = 0
    i = 0
    for j in range(0, Ny+1): u[i,j] = 0
    i = Nx
    for j in range(0, Ny+1): u[i,j] = 0
    return u
```

This example may act as a recipe on how to transform array-intensive code with loops into Cython.

1. Variables are declared with types: for example, `double v` in the argument list and `cdef double v` for a variable `v` in the body of the function. A Python `float` variable is declared as `double` for translation to C by Cython, while an `int` object is declared by `int`.
2. Arrays need a comprehensive type declaration involving
 - the type `np.ndarray`,
 - the data type of the elements, here 64-bit floats, abbreviated as `DT` through `ctypedef np.float64_t DT` (instead of `DT` we could use the full name of the data type: `np.float64_t`, which is a Cython-defined type),

- the dimensions of the array, here `ndim=2` and `ndim=1`,
 - specification of contiguous memory for the array (`mode='c'`)
3. Functions declared with `cpdef` are translated to C but also accessible from Python.
 4. In addition to the standard `numpy` import we also need a special Cython import of `numpy`: `cimport numpy as np`, to appear *after* the standard import.
 5. By default, array indices are checked to be within their legal limits. To speed up the code one should turn off this feature for a specific function by placing `@cython.boundscheck(False)` above the function header.
 6. Also by default, array indices can be negative (counting from the end), but this feature has a performance penalty and is therefore here turned off by writing `@cython.wraparound(False)` right above the function header.

Visual inspection of the C translation. Cython can visually explain how successfully it can translate a code from Python to C. The command

```
Terminal> cython -a wave2D_u0_loop_cy.pyx
```

produces an HTML file `wave2D_u0_loop_cy.html`, which can be loaded into a web browser to illustrate which lines of the code that have been translated to C. Figure 5 shows the results. Yellow lines indicate the lines that Cython did not manage to translate to efficient C code and that remain in Python. For the present code we see that Cython is able to translate all the loops with array computing to C, which is our primary goal.

You can also inspect the generated C code directly, in the file `wave2D_u0_loop_cy.c`, but understanding this C code requires familiarity with writing Python extension modules in C by hand. However, deep down in the file we can see in detail how the compute-intensive statements are translated some complex C that is quite different from what we a human would write, at least with a direct corresponding to the mathematics in mind.

Building the extension module. Cython code must be translated to C, compiled, and linked to form what is known in the Python world as a *C extension module*. This is usually done by making a `setup.py` script, which is the standard way of building and installing Python software:

Raw output: [wave2D_u0_loop_cy.c](#)

```
1: import numpy as np
2: cimport numpy as np
3: cimport cython
4: ctypedef np.float64_t DT # data type
5:
6: @cython.boundscheck(False) # turn off array bounds check
7: @cython.wraparound(False) # turn off negative indices (u[-1,-1])
8: cdef advance(
9:     np.ndarray[DT, ndim=2, mode='c'] u,
10:     np.ndarray[DT, ndim=2, mode='c'] u_1,
11:     np.ndarray[DT, ndim=2, mode='c'] u_2,
12:     np.ndarray[DT, ndim=2, mode='c'] f,
13:     np.ndarray[DT, ndim=1, mode='c'] x,
14:     np.ndarray[DT, ndim=1, mode='c'] y,
15:     np.ndarray[DT, ndim=1, mode='c'] t,
16:     double Cx2, double Cy2, double dt2):
17:
18:     cdef int Nx, Ny, i, j
19:     Nx = u.shape[0]-1
20:     Ny = u.shape[1]-1
21:     for i in xrange(1, Nx):
22:         for j in xrange(1, Ny):
23:             u[i,j] = 2*u_1[i,j] - u_2[i,j] + \
24:                 Cx2*(u_1[i-1,j] - 2*u_1[i,j] + u_1[i+1,j]) + \
25:                 Cy2*(u_1[i,j-1] - 2*u_1[i,j] + u_1[i,j+1]) + \
26:                 dt2*f[i,j]
27:     # Boundary condition u=0
28:     j = 0
29:     for i in range(0, Nx+1): u[i,j] = 0
30:     j = Ny
31:     for i in range(0, Nx+1): u[i,j] = 0
32:     i = 0
33:     for j in range(0, Ny+1): u[i,j] = 0
34:     i = Nx
35:     for j in range(0, Ny+1): u[i,j] = 0
36:     return u
```

Figure 5: Visual illustration of Cython's ability to translate Python to C.

```
from distutils.core import setup
from distutils.extension import Extension
from Cython.Distutils import build_ext

cymodule = 'wave2D_u0_loop_cy'
setup(
    name=cymodule
    ext_modules=[Extension(cymodule, [cymodule + '.pyx'],)],
    cmdclass={'build_ext': build_ext},
)
```

To translate to C, compile, and build the module `wave2D_u0_loop_cy`, run

```
Terminal> python setup.py build_ext --inplace
```

The `--inplace` option makes the extension module available in the current directory as the file `wave2D_u0_loop_cy.so`. This file acts as a normal Python module:

```
>>> import wave2D_u0_loop_cy
>>> dir(wave2D_u0_loop_cy)
['__builtins__', '__doc__', '__file__', '__name__',
 '__package__', '__test__', 'advance', 'np']
```

The `setup.py` file makes use of the `distutils` package in Python and Cython's extension of this package. These tools know how Python was built on the computer and will use compatible compiler(s) and options when building other code in Cython, C, or C++. Quite some experience with building large program systems is needed to do the build process manually.

When there is no need to link the C code with special libraries, Cython offers a shortcut for generating and importing the extension module:

```
import pyximport; pyximport.install()
```

This makes the `setup.py` script redundant. However, in the `wave2D_u0.py` code we do not use `pyximport` and require an explicit build process of this and many other modules.

Calling the Cython function. The `wave2D_u0_loop_cy` module contains our `advance` function, which we now may call from the Python program for the wave equation:

```
import wave2D_u0_loop_cy
cython_advance = wave2D_u0_loop_cy.advance
...
for n in range(1, N):      # time loop
    f_a = f(xv, yv, t[n])  # precompute, size as u
    u = cython_advance(u, u_1, u_2, f_a, x, y, t,
                      Cx2, Cy2, dt2)
```

Efficiency. For a mesh consisting of 120×120 cells, the scalar Python code require 1370 CPU time units, the vectorized version requires 5.5, while the Cython version requires only 1! For a mesh with 60×60 cells Cython is about 1000 times faster than the scalar Python code, and the vectorized version is about 6 times slower than the Cython version.

10.5 Migrating loops to Fortran

Instead of relying on Cython's (excellent) ability to translate Python to C, we can invoke a compiled language directly and write the loops ourselves. Let us start with Fortran 77, because this is a language with more convenient array handling than C (or plain C++). Or more precisely, we can with ease program with multi-dimensional indices in the `numpy` arrays that are transferred to Fortran, while in C these arrays are one-dimensional. Let us also mention that the Fortran compilers build on 60 years of intensive research on how to optimize loops with array computations.

The Fortran subroutine. We write a Fortran subroutine `advance` in a file `wave2D_u0_loop_f77.f` for implementing the updating formula (67) and setting the solution to zero at the boundaries:

```

subroutine advance(u, u_1, u_2, f, Cx2, Cy2, dt2, Nx, Ny)
  integer Nx, Ny
  real*8 u(0:Nx,0:Ny), u_1(0:Nx,0:Ny), u_2(0:Nx,0:Ny)
  real*8 f(0:Nx, 0:Ny), Cx2, Cy2, dt2
  integer i, j
Cf2py intent(in, out) u

C    Scheme at interior points
do j = 1, Ny-1
  do i = 1, Nx-1
    u(i,j) = 2*u_1(i,j) - u_2(i,j) +
&      Cx2*(u_1(i-1,j) - 2*u_1(i,j) + u_1(i+1,j)) +
&      Cy2*(u_1(i,j-1) - 2*u_1(i,j) + u_1(i,j+1)) +
&      dt2*f(i,j)
  end do
end do

C    Boundary conditions
j = 0
do i = 0, Nx
  u(i,j) = 0
end do
j = Ny
do i = 0, Nx
  u(i,j) = 0
end do
i = 0
do j = 0, Ny
  u(i,j) = 0
end do
i = Nx
do j = 0, Ny
  u(i,j) = 0
end do
return
end

```

This code is plain Fortran 77, except for the special `Cf2py` comment line, which here specifies that `u` is both an input argument *and* an object to be returned from the `advance` routine. Or more precisely, Fortran is not able return an array from a function, but we need a *wrapper code* in C for the Fortran subroutine to enable calling it from Python, and in this wrapper code one can return `u` to the calling Python code. It is not strictly necessary to return `u` to the calling Python code since the `advance` function will modify the elements of `u`, but the convention in Python is to get all output from a function as returned values.

Building the Fortran module with `f2py`. The nice feature of writing loops in Fortran is that the tool `f2py` can with very little work produce a C extension module such that we can call the Fortran version of `advance` from Python. The necessary commands to run are

```

Terminal> f2py -m wave2D_u0_loop_f77 -h wave2D_u0_loop_f77.pyf \
--overwrite-signature wave2D_u0_loop_f77.f
Terminal> f2py -c wave2D_u0_loop_f77.pyf --build-dir build_f77 \
-DF2PY_REPORT_ON_ARRAY_COPY=1 wave2D_u0_loop_f77.f

```

First, `f2py` interprets the Fortran code and makes a Fortran 90 specification of the extension module in the file `wave2D_u0_loop_f77.pyf`. Second, `f2py` generates all necessary wrapper code, compiles our Fortran file and the wrapper code, and finally builds the module. The build process takes place in the specified subdirectory `build_f77` so that files can be inspected if something goes wrong. The option `-DF2PY_REPORT_ON_ARRAY_COPY=1` makes `f2py` write a message for every array that is copied in the communication between Fortran and Python, which is very useful for avoiding unnecessary array copying (see below). The name of the module file is `wave2D_u0_loop_f77.so`, and this file can be imported as any other Python module:

```
>>> import wave2D_u0_loop_f77
>>> dir(wave2D_u0_loop_f77)
['__doc__', '__file__', '__name__', '__package__',
 '__version__', 'advance']
>>> print wave2D_u0_loop_f77.__doc__
This module 'wave2D_u0_loop_f77' is auto-generated with f2py...
Functions:
    u = advance(u,u_1,u_2,f,x,y,t,cx2,cy2,dt2,
               nx=(shape(u,0)-1),ny=(shape(u,1)-1),n=(len(t)-1))
```

Examining doc strings. Printing the doc strings of the module and its functions is extremely important after having created a module with `f2py`, because `f2py` makes Python interfaces to the Fortran functions that are different from how the functions are declared in the Fortran code (!). The rationale for this behavior is that `f2py` creates *Pythonic* interfaces to that Fortran routines can be called in the same way as one calls Python functions. Since output data from Python functions are always returned to the calling code and this is technically impossible in Fortran. Also, arrays in Python are passed to Python functions without their dimensions because that information is packed with the array data in the array objects, but this is not possible in Fortran. Therefore, `f2py` removes array dimensions from the argument list to a function like `advance` (!), and it makes it possible to return objects back to Python. Let us take a closer look at the documentation `f2py` has generated for our Fortran `advance` subroutine:

```
>>> print wave2D_u0_loop_f77.advance.__doc__
This module 'wave2D_u0_loop_f77' is auto-generated with f2py
Functions:
    u = advance(u,u_1,u_2,f,cx2,cy2,dt2,
               nx=(shape(u,0)-1),ny=(shape(u,1)-1))
.
advance - Function signature:
    u = advance(u,u_1,u_2,f,cx2,cy2,dt2,[nx,ny])
Required arguments:
    u : input rank-2 array('d') with bounds (nx + 1,ny + 1)
    u_1 : input rank-2 array('d') with bounds (nx + 1,ny + 1)
    u_2 : input rank-2 array('d') with bounds (nx + 1,ny + 1)
    f : input rank-2 array('d') with bounds (nx + 1,ny + 1)
```

```

cx2 : input float
cy2 : input float
dt2 : input float
Optional arguments:
  nx := (shape(u,0)-1) input int
  ny := (shape(u,1)-1) input int
Return objects:
  u : rank-2 array('d') with bounds (nx + 1,ny + 1)

```

Here we see that the `nx`, `ny`, and `n` parameters declared in Fortran are optional arguments that can be omitted when calling `advance` from Python. We strongly recommend to print out the documentation of *every* Fortran function to be called from Python and make sure the call syntax is exactly as listed in the documentation.

How to avoid array copying. Multi-dimensional arrays are stored as a stream of numbers in memory. For a two-dimensional array consisting of rows and columns there are two ways of creating such a stream: *row-major ordering*, which means that rows are stored consecutively in memory, or *column-major ordering*, which means that the columns are stored after each other. All programming languages inherited from C, including Python, apply the row-major ordering, but Fortran uses column-major storage. Thinking of a two-dimensional array as a matrix, it means that Fortran works with the transposed matrix.

Fortunately, `f2py` creates extra code so that accessing `u(i,j)` in the Fortran subroutine corresponds to the element `u[i,j]` in the underlying `numpy` array (without the extra code, `u(i,j)` in Fortran would access `u[j,i]` in the `numpy` array). Technically, `f2py` takes a copy of our `numpy` arrays and reorders the data before sending the arrays to Fortran. Such copying can be costly. For 2D wave simulations on a 60×60 grid the overhead of copying is a factor of 5, which means that almost the whole performance gain of Fortran over vectorized `numpy` code is lost!

To avoid having `f2py` to copy arrays with C storage to the corresponding Fortran storage, we can declare the arrays with Fortran storage:

```

order = 'Fortran' if version == 'f77' else 'C'
u = zeros((Nx+1,Ny+1), order=order) # solution array
u_1 = zeros((Nx+1,Ny+1), order=order) # solution at t-dt
u_2 = zeros((Nx+1,Ny+1), order=order) # solution at t-2*dt

```

In the compile and build step of using `f2py`, it is recommended to add an extra option for making `f2py` report on array copying:

```

Terminal> f2py -c wave2D_u0_loop_f77.pyf --build-dir build_f77 \
          -DF2PY_REPORT_ON_ARRAY_COPY=1 wave2D_u0_loop_f77.f

```

It can sometimes be a challenge to track down which array that causes a copying. There are two principal reasons for copying array data: either the array does not have Fortran storage or the element types do not match those

declared in the Fortran code. The latter cause is usually effectively eliminated by using `real*8` data in the Fortran code and `float64` (the default `float` type in `numpy`) in the arrays on the Python side. The former reason is more common, and to check whether an array before a Fortran call has the right storage one can print the result of `isfortran(a)`, which is `True` if the array `a` has Fortran storage. A typical problem in the `wave2D_u0.py` code is to set

```
f_a = f(xv, yv, t[n])
```

before the call to the Fortran `advance` routine. This computation creates a new array with C storage. There are two remedies, either direct insertion of data in an array with Fortran storage,

```
f_a = zeros((Nx+1, Ny+1), order='Fortran')
...
f_a[:, :] = f(xv, yv, t[n])
```

or remaking the `f(xv, yv, t[n])` array,

```
f_a = asarray(f(xv, yv, t[n]), order='Fortran')
```

Efficiency. The efficiency of this Fortran code is very similar to the Cython code. There is usually nothing more to gain, from a computational efficiency point of view, by implementing the complete Python program in Fortran or C. That will just be a lot more code for all administering work that is needed in scientific software, especially if we extend our sample program `wave2D_u0.py` to handle a real scientific problem. Then only a small portion will consist of loops with intensive array calculations. These can be migrated to Cython or Fortran as explained, while the rest of the programming can be more conveniently done in Python.

10.6 Migrating loops to C via Cython

The computationally intensive loops can alternatively be implemented in C code. Just as Fortran calls for care regarding the storage of two-dimensional arrays, working with two-dimensional arrays in C is a bit tricky. The reason is that `numpy` arrays are viewed as one-dimensional arrays when transferred to C, while C programmers will think of `u`, `u_1`, and `u_2` as two dimensional arrays and index them like `u[i][j]`. The C code must declare `u` as `double* u` and translate an index pair `[i][j]` to a corresponding single index when `u` is viewed as one-dimensional. This translation requires knowledge of how the numbers in `u` are stored in memory.

Translating index pairs to single indices. Two-dimensional `numpy` arrays with the default C storage are stored row by row. In general, multi-dimensional arrays with C storage are stored such that the last index has the fastest variation,

then the next last index, and so on, ending up with the slowest variation in the first index. For a two-dimensional `u` declared as `zeros((Nx+1,Ny+1))` in Python, the individual elements are stored in the following order:

```
u[0,0], u[0,1], u[0,2], ..., u[0,Ny], u[1,0], u[1,1], ...,
u[1,Ny], u[2,0], ..., u[Nx,0], u[Nx,1], ..., u[Nx, Ny]
```

Viewing `u` as one-dimensional, the index pair (i, j) translates to $i*(Ny+1) + j$. Where a C programmer would naturally write an index `u[i][j]`, the indexing must read `u[i*(Ny+1) + j]`. This is tedious to write, so it can be handy to define a C macro,

```
#define idx(i,j) (i)*(Ny+1) + j
```

so that we can write `u[idx(i,j)]`, which reads much better and is easier to debug. Macros perform simple text substitutions: `idx(hello,world)` is expanded to `(hello)*(Ny+1) + world`. The parenthesis in `(i)` are essential - with natural mathematical formula $i*(Ny+1) + j$ in the macro definition, `idx(i-1,j)` would expand to `i-1*(Ny+1) + j`, which is the wrong formula. Macros are handy, but requires careful use.

The complete C code. The C version of function `advance` can be coded as follows.

```
#define idx(i,j) (i)*(Ny+1) + j

void advance(double* u, double* u_1, double* u_2, double* f,
             double Cx2, double Cy2, double dt2,
             int Nx, int Ny)
{
    int i, j;
    /* Scheme at interior points */
    for (i=1; i<=Nx-1; i++) {
        for (j=1; j<=Ny-1; j++) {
            u[idx(i,j)] = 2*u_1[idx(i,j)] - u_2[idx(i,j)] +
                Cx2*(u_1[idx(i-1,j)] - 2*u_1[idx(i,j)] + u_1[idx(i+1,j)]) +
                Cy2*(u_1[idx(i,j-1)] - 2*u_1[idx(i,j)] + u_1[idx(i,j+1)]) +
                dt2*f[idx(i,j)];
        }
    }
    /* Boundary conditions */
    j = 0; for (i=0; i<=Nx; i++) u[idx(i,j)] = 0;
    j = Ny; for (i=0; i<=Nx; i++) u[idx(i,j)] = 0;
    i = 0; for (j=0; j<=Ny; j++) u[idx(i,j)] = 0;
    i = Nx; for (j=0; j<=Ny; j++) u[idx(i,j)] = 0;
}
```

The Cython interface file. All the code above appears in a file `wave2D_u0_loop_c.c`. We need to compile this file together with C wrapper code such that `advance` can be called from Python. Cython can be used to interface our

C code. The Cython code, placed in a file with extension `.pyx`, here `wave2D_u0_loop_c.pyx`, looks like

```
import numpy as np
cimport numpy as np
cimport cython

cdef extern from "wave2D_u0_loop_c.h":
    void advance(double* u, double* u_1, double* u_2, double* f,
                double Cx2, double Cy2, double dt2,
                int Nx, int Ny)

@cython.boundscheck(False)
@cython.wraparound(False)
def advance_cwrap(
    np.ndarray[double, ndim=2, mode='c'] u,
    np.ndarray[double, ndim=2, mode='c'] u_1,
    np.ndarray[double, ndim=2, mode='c'] u_2,
    np.ndarray[double, ndim=2, mode='c'] f,
    double Cx2, double Cy2, double dt2):
    advance(&u[0,0], &u_1[0,0], &u_2[0,0], &f[0,0],
          Cx2, Cy2, dt2,
          u.shape[0]-1, u.shape[1]-1)
    return u
```

Here, we first declare the C functions to be interfaced. These must also appear in a C header file, `wave2D_u0_loop_c.h`,

```
extern void advance(double* u, double* u_1, double* u_2, double* f,
                  double Cx2, double Cy2, double dt2,
                  int Nx, int Ny);
```

The next step is to write a Cython function with Python objects as arguments. The name `advance` is already used for the C function so the function to be called from Python is named `advance_cwrap`. The contents of this function is simply a call to the `advance` version in C. To this end, the right information from the Python objects must be passed on as arguments to `advance`. Arrays are sent with their C pointers to the first element, obtained in Cython as `&u[0,0]` (the `&` takes the address of a C variable). The `Nx` and `Ny` arguments in `advance` are easily obtained from the shape of the `numpy` array `u`. Finally, `u` must be returned such that we can set `u = advance(...)` in Python.

Building the extension module. It remains to build the extension module. An appropriate `setup.py` file is

```
from distutils.core import setup
from distutils.extension import Extension
from Cython.Distutils import build_ext

sources = ['wave2D_u0_loop_c.c', 'wave2D_u0_loop_c.pyx']
module = 'wave2D_u0_loop_c_cy'
setup(
    name=module,
    ext_modules=[Extension(module, sources,
```

```

        libraries=[], # C libs to link with
    ],
    cmdclass={'build_ext': build_ext},
)

```

All we need to specify is the `.c` file(s) and the `.pyx` interface file. Cython is automatically run to generate the necessary wrapper code. Files are then compiled and linked to an extension module residing in the file `wave2D_u0_loop_c_cy.so`. This module can be imported in Python,

```

>>> import wave2D_u0_loop_c_cy as m
>>> dir(m)
['__builtins__', '__doc__', '__file__', '__name__', '__package__',
 '__test__', 'advance_cwrap', 'np']

```

The call to the C version of `advance` can go like this in Python:

```

advance = wave2D_u0_loop_c_cy.advance_cwrap
f_a[:, :] = f(xv, yv, t[n])
u = advance(u, u_1, u_2, f_a, Cx2, Cy2, dt2)

```

Efficiency. In this example, the C and Fortran code runs at the same speed, and there are no significant differences in the efficiency of the wrapper code. The overhead implied by the wrapper code is negligible as long as we do not work with very small meshes and consequently little numerical work in the `advance` function.

10.7 Migrating loops to C via `f2py`

As an alternative to using Cython for interfacing C code, we can apply `f2py`. The C code is the same, just the details of specifying how it is to be called from Python differ. The `f2py` tool requires the call specification to be a Fortran 90 module defined in a `.pyf` file. This file was automatically generated when we interfaced a Fortran subroutine. With a C function we need to write this module ourselves, or we can use a trick and let `f2py` generate it for us. The trick consists in writing the signature of the C function with Fortran syntax and place it in a Fortran file, here `wave2D_u0_loop_c_f2py_signature.f`:

```

subroutine advance(u, u_1, u_2, f, Cx2, Cy2, dt2, Nx, Ny)
Cf2py intent(c) advance
    integer Nx, Ny, N
    real*8 u(0:Nx,0:Ny), u_1(0:Nx,0:Ny), u_2(0:Nx,0:Ny)
    real*8 f(0:Nx, 0:Ny), Cx2, Cy2, dt2
Cf2py intent(in, out) u
Cf2py intent(c) u, u_1, u_2, f, Cx2, Cy2, dt2, Nx, Ny
    return
end

```

Since `f2py` is just concerned with the signature and not the complete contents of the function body, it can generate the Fortran 90 module specification:

```
Terminal> f2py -m wave2D_u0_loop_c_f2py \
             -h wave2D_u0_loop_c_f2py.pyf --overwrite-signature \
             wave2D_u0_loop_c_f2py_signature.f
```

The compile and build step is as for the Fortran code, except that we list C file(s) instead of Fortran file(s):

```
Terminal> f2py -c wave2D_u0_loop_c_f2py.pyf \
            --build-dir tmp_build_c \
            -DF2PY_REPORT_ON_ARRAY_COPY=1 wave2D_u0_loop_c.c
```

As when interfacing Fortran code with `f2py`, we need to print out the doc string to see the exact call syntax from the Python side. This doc string is identical for the C and Fortran versions of `advance`.

10.8 Migrating loops to C via Instant

10.9 Migrating loops to C++ via `f2py`

C++ is a much more versatile language than C or Fortran and has over the last two decades become very popular for numerical computing. Many will therefore prefer to migrate compute-intensive Python code to C++. This is, in principle, easy: just write the desired C++ code and use some tool for interfacing it from Python. A tool like SWIG⁴ can interpret the C++ code and generate interfaces for Python, Perl, Ruby, Java, and a wide range of languages. However, SWIG is a comprehensive tool with a correspondingly non-trivial learning curve. Alternative tools, such as Boost Python and SIP, are similarly comprehensive.

A technically much easier way of interfacing C++ code is to drop the possibility to use C++ classes directly from Python, but instead make a C interface to the C++ code. The C interface can be handled by `f2py` as shown in the example with pure C code. Such a solution means that classes in Python and C++ cannot be mixed and that only primitive data types like numbers, strings, and arrays can be transferred between Python and C++. Actually, this is often a very good solution because it forces the C++ code to work on array data, which usually gives faster code than if fancy data structures with classes are used. The arrays coming from Python, and looking like plain C/C++ arrays, can be efficiently wrapped in more user-friendly C++ array classes in the C++ code, if desired.

Remaining. Use some array class. Key issue: `extern "C"` declaration of C++ function in the C file with the interface we want to wrap.

⁴<http://swig.org/>

10.10 Using classes to implement a simulator

- class Mesh, Function, Problem, Solver, Visualizer, File
- communicate with compiled code by ensuring that classes work with arrays

10.11 Callbacks to Python from Fortran or C

11 Analysis of the continuous and discrete solutions

11.1 Properties of the solution of the wave equation

The wave equation

$$\frac{\partial^2 u}{\partial t^2} = c^2 \frac{\partial^2 u}{\partial x^2}$$

has solutions of the form

$$u(x, t) = g_R(x - ct) + g_L(x + ct), \quad (69)$$

for any functions g_R and g_L sufficiently smooth to be differentiated twice. The result follows from inserting (69) in the wave equation. A function of the form $g_R(x - ct)$ represents a signal moving to the right in time with constant velocity c . This feature can be explained as follows. At time $t = 0$ the signal looks like $g_R(x)$. Introducing a moving x axis with coordinates $\xi = x - ct$, we see the function $g_R(\xi)$ is "at rest" in the ξ coordinate system, and the shape is always the same. Say the $g_R(\xi)$ function has a peak at $\xi = 0$. This peak is located at $x = ct$, which means that it moves with the velocity $dx/dt = c$ in the x coordinate system. Similarly, $g_L(x + ct)$ is a function initially with shape $g_L(x)$ that moves in the negative x direction with constant velocity c .

With the particular initial conditions

$$u(x, 0) = I(x), \quad \frac{\partial}{\partial t} u(x, 0) = 0,$$

we get, with u as in (69),

$$g_R(x) + g_L(x) = I(x), \quad -cg'_R(x) + cg'_L(x) = 0,$$

which have the solution $g_R = g_L = I/2$, and consequently

$$u(x, t) = \frac{1}{2}I(x - ct) + \frac{1}{2}I(x + ct). \quad (70)$$

The interpretation of (70) is that the initial shape of u is split into two parts, each with the same shape as I but half of the initial amplitude. One part is traveling to the left and the other one to the right.

The solution (70) will be influenced by boundary conditions when the parts $\frac{1}{2}I(x-ct)$ and $\frac{1}{2}I(x+ct)$ hit the boundaries and get, e.g., reflected back into the domain. However, when $I(x)$ is nonzero only in a small part in the middle of the spatial domain $[0, L]$, which means that the boundaries are placed far away from the initial disturbance of u , the solution (70) is very clearly observed in a simulation.

A useful representation of solutions of wave equations is a linear combination of sine and/or cosine waves. Such a sum of waves is a solution if the governing PDE is linear and each wave fulfills the equation. To ease analytical calculations by hand we shall work with complex exponential functions instead of real-valued sine or cosine functions. The real part of complex expressions will typically be taken as the physical relevant quantity (whenever a physical relevant quantity is strictly needed). The idea now is to build $I(x)$ of complex wave components e^{ikx} :

$$I(x) \approx \sum_{k \in K} b_k e^{ikx}. \quad (71)$$

Here, k is the frequency of a component, K is some set of all the discrete k values needed to approximate $I(x)$ well, and b_k are constants that must be determined. We will very seldom need to compute the b_k coefficients: most of the insight we look for, and understanding of the numerical methods, comes from investigating how the PDE and the scheme treat a single component e^{ikx} of the solution.

Letting the number of k values in K tend to infinity makes the sum (71) converge to $I(x)$, and this sum is known as a *Fourier series* representation of $I(x)$. Looking at (70), we see that the solution $u(x, t)$, when $I(x)$ is represented as in (71), is also built of basic complex exponential wave components of the form $e^{ik(x \pm ct)}$ according to

$$u(x, t) = \frac{1}{2} \sum_{k \in K} b_k e^{ik(x-ct)} + \frac{1}{2} \sum_{k \in K} b_k e^{ik(x+ct)}. \quad (72)$$

It is common to introduce the frequency in time $\omega = kc$ and assume that $u(x, t)$ is a sum of basic wave components written as $e^{ikx - \omega t}$. Observe that inserting such a wave component in the governing PDE reveals that $\omega^2 = k^2 c^2$, or $\omega \pm kc$, reflecting the two solutions: one $(+kc)$ traveling to the right and the other $(-kc)$ traveling to the left.

11.2 Analysis of the finite difference scheme

The scheme

$$[D_t D_t u = c^2 D_x D_x u]_p^n \quad (73)$$

for the wave equation $u_t = c^2 u_{xx}$ allows basic wave components $e^{i(kx - \tilde{\omega}t)}$ as solution, but it turns out that the frequency in time, $\tilde{\omega}$, is not equal to the exact $\omega = kc$. The idea now is to study how the scheme treats an arbitrary wave component with a given k . How accurate is $\tilde{\omega}$ compared to ω ? And does

the amplitude of such a wave component preserve its (unit) amplitude, as it should, or does it get amplified or damped in time (due to a complex $\tilde{\omega}$)? These are the questions we aim to answer in the following analysis. Note the need for using p as counter for the mesh point in x direction since i is already used as the imaginary unit (in this analysis).

A key result needed in the investigations is the finite difference approximation of a second-order derivative acting on a complex wave component:

$$[D_t D_t e^{i\omega t}]^n = -\frac{4}{\Delta t^2} \sin^2\left(\frac{\omega \Delta t}{2}\right) e^{i\omega n \Delta t}.$$

Similarly, by just changing symbols ($\omega \rightarrow k$, $t \rightarrow x$, $n \rightarrow p$) it follows that

$$[D_x D_x e^{ikx}]_p = -\frac{4}{\Delta x^2} \sin^2\left(\frac{k \Delta x}{2}\right) e^{ikp \Delta x}.$$

Inserting a basic wave component $u = e^{i(kx - \tilde{\omega}t)}$ in (73) results in the need to evaluate two expressions:

$$\begin{aligned} [D_t D_t e^{ikx} e^{-i\tilde{\omega}t}]_p^n &= [D_t D_t e^{-i\tilde{\omega}t}]^n e^{ikp \Delta x} \\ &= -\frac{4}{\Delta t^2} \sin^2\left(\frac{\tilde{\omega} \Delta t}{2}\right) e^{-i\tilde{\omega} n \Delta t} e^{ikp \Delta x} \end{aligned} \quad (74)$$

$$\begin{aligned} [D_x D_x e^{ikx} e^{-i\tilde{\omega}t}]_p^n &= [D_x D_x e^{ikx}]_p e^{-i\tilde{\omega} n \Delta t} \\ &= -\frac{4}{\Delta x^2} \sin^2\left(\frac{k \Delta x}{2}\right) e^{ikp \Delta x} e^{-i\tilde{\omega} n \Delta t}. \end{aligned} \quad (75)$$

Then the complete scheme,

$$[D_t D_t e^{ikx} e^{-i\tilde{\omega}t}] = c^2 [D_x D_x e^{ikx} e^{-i\tilde{\omega}t}]_p^n$$

leads to the following equation for the unknown numerical frequency $\tilde{\omega}$ (after dividing by $-e^{ikx} e^{-i\tilde{\omega}t}$):

$$\frac{4}{\Delta t^2} \sin^2\left(\frac{\tilde{\omega} \Delta t}{2}\right) = c^2 \frac{4}{\Delta x^2} \sin^2\left(\frac{k \Delta x}{2}\right),$$

or

$$\sin^2\left(\frac{\tilde{\omega} \Delta t}{2}\right) = C^2 \sin^2\left(\frac{k \Delta x}{2}\right), \quad (76)$$

where

$$C = \frac{c \Delta t}{\Delta x} \quad (77)$$

is the Courant number. Taking the square root of (76) yields

$$\sin\left(\frac{\tilde{\omega} \Delta t}{2}\right) = C \sin\left(\frac{k \Delta x}{2}\right), \quad (78)$$

Since the exact ω is real it is reasonable to look for a real solution $\tilde{\omega}$ of (78). The right-hand side of (78) must then be in $[-1, 1]$ because the sine function on the left-hand side has values in $[-1, 1]$ for real $\tilde{\omega}$. The sine function on the right-hand side can attain the value 1 when

$$\frac{k\Delta x}{2} = q\frac{\pi}{2}, \quad q \in \mathbb{N}.$$

With $q = 1$ we have $k\Delta x = \pi$, which means that the wavelength $\lambda = 2\pi/k$ becomes $2\Delta x$. This is the absolutely shortest wavelength that can be represented on the mesh: the wave jumps up and down between each mesh point. Higher values of q are irrelevant since these waves are too short to be represented on a mesh with spacing Δx . For the shortest possible wave in the mesh, $\sin(k\Delta x/2) = 1$ and we must require

$$C \leq 1. \quad (79)$$

For smoother wave components with longer wave lengths per length Δx , (79) can in theory be relaxed. However, small round-off errors are always present in a numerical solution and these vary arbitrarily from mesh point to mesh point and can be viewed as unavoidable noise with wavelength $2\Delta x$. As we shall see below, $C > 1$ will for this very small noise lead to exponential growth of the shortest possible wave component in the mesh. This noise will therefore grow with time and destroy the whole solution.

Consider a right-hand side in (78) of magnitude larger than unity. The solution $\tilde{\omega}$ of (78) must then be a complex number $\tilde{\omega} = \tilde{\omega}_r + i\tilde{\omega}_i$. One can show that for any ω_i there will also be a corresponding solution with $-\omega_i$. The component with $\omega_i > 0$ gives an amplification factor $e^{\omega_i t}$ that grows exponentially in time. We cannot allow this and must therefore require $C \leq 1$ as a *stability criterion*.

Equation (78) can be solved with respect to $\tilde{\omega}$:

$$\tilde{\omega} = \frac{2}{\Delta t} \sin^{-1} \left(C \sin \left(\frac{k\Delta x}{2} \right) \right). \quad (80)$$

The relation between the numerical frequency $\tilde{\omega}$ and the other parameters k , c , Δx , and Δt is called a *numerical dispersion relation*. Correspondingly, $\omega = kc$ is the *analytical dispersion relation*.

The special case $C = 1$ deserves attention since then the right-hand side of (80) reduces to

$$\frac{2}{\Delta t} \frac{k\Delta x}{2} = \frac{1}{\Delta t} \frac{\omega\Delta x}{c} = \frac{\omega}{C} = \omega.$$

That is, $\tilde{\omega} = \omega$ and the numerical solution is exact at all mesh points regardless of Δx and Δt ! This implies that the numerical solution method is also an analytical solution method, at least for computing u at discrete points (the numerical method says nothing about the variation of u *between* the mesh points, and employing the common linear interpolation for extending the discrete solution gives a curve that deviates from the exact one).

For a closer examination of the error in the numerical dispersion relation when $C < 1$, we can study $\tilde{\omega} - \omega$, $\tilde{\omega}/\omega$, or the similar error measures in wave velocity: $\tilde{c} - c$ and \tilde{c}/c , where $c = \omega/k$ and $\tilde{c} = \tilde{\omega}/k$. It appears that the most convenient expression to work with is \tilde{c}/c :

$$\frac{\tilde{c}}{c} = \frac{1}{Cp} \sin^{-1}(C \sin p),$$

with $p = k\Delta x/2$ as a non-dimensional measure of spatial frequency. In essence, p tells how many spatial mesh points we have per spatial period of the wave component with frequency k (the period is $2\pi/k$), i.e., how well the spatial variation of the wave component is resolved in the mesh. Wave components with wave length less than $2\Delta x$ ($2\pi/k < 2\Delta x$) are not visible in the mesh, so it does not make sense to have $p > \pi/2$.

Note that the symbol p , previously used as index, is here reused with another meaning. We may introduce the function $r(C, p) = \tilde{c}/c$:

$$r(C, p) = \frac{1}{Cp} \sin^{-1}(C \sin p), \quad C \in (0, 1], \quad p \in (0, \pi/2]. \quad (81)$$

This function is very well suited for plotting since it combines several parameters in the problem into a dependence on two non-dimensional numbers C and p .




Figure 6: The fractional error in the wave velocity for different Courant numbers.

Defining

```
def r(C, p):
    return 2/(C*p)*asin(C*sin(p))
```

we can easily plot $r(C, p)$ as a function of p for various values of C , see Figure 6. Note that the shortest waves have the most erroneous velocity, and that short waves move more slowly than they should.

With `sympy` we can also easily make a Taylor series expansion in the discretization parameter p :

```
>>> C, p = symbols('C p')
>>> rs = r(C, p).series(p, 0, 7)
>>> print rs
1 - p**2/6 + p**4/120 - p**6/5040 + C**2*p**2/6 -
C**2*p**4/12 + 13*C**2*p**6/720 + 3*C**4*p**4/40 -
C**4*p**6/16 + 5*C**6*p**6/112 + O(p**7)
>>> rs_factored = [factor(term) for term in rs.lseries(p)]
>>> rs_factored = sum(rs_factored)
>>> print rs_factored
p**6*(C - 1)*(C + 1)*(225*C**4 - 90*C**2 + 1)/5040 +
p**4*(C - 1)*(C + 1)*(3*C - 1)*(3*C + 1)/120 +
p**2*(C - 1)*(C + 1)/6 + 1
```

We see that $C = 1$ makes all the terms in `rs_factored` vanish, except the last one. Since we already know that the numerical solution is exact for $C = 1$, the remaining terms in the Taylor series expansion will also contain factors of $C - 1$ and cancel for $C = 1$.

From `rs_factored` expression above we also see that the leading order terms in the error of this series expansion are

$$\frac{k^2}{24} (c^2 \Delta t^2 - \Delta x^2) \quad (82)$$

pointing to an error $\mathcal{O}(\Delta t^2, \Delta x^2)$, which is compatible with the errors in the difference approximations ($D_t D_t$ and $D_x D_x$).

11.3 Extending the analysis to 2D and 3D

The typical analytical solution of a 2D wave equation

$$u_{tt} = c^2(u_{xx} + u_{yy}),$$

is a wave traveling in the direction of $\mathbf{k} = k_x \mathbf{i} + k_y \mathbf{j}$, where \mathbf{i} and \mathbf{j} are unit vectors in the x and y directions, respectively. Such a wave can be expressed by

$$u(x, y, t) = g(k_x x + k_y y - kct)$$

or

$$u(x, y, t) = g(k_x x + k_y y - \omega t)$$

We can in particular build a solution from complex Fourier components of the form

$$\exp(i(k_x x + k_y y - \omega t)).$$

A discrete 2D wave equation can be written as

$$[D_t D_t u = c^2(D_x D_x u + D_y D_y u)]_{p,q}^n. \quad (83)$$

This equation admits a Fourier component

$$u_{p,q}^n = \exp(i(k_x p \Delta x + k_y q \Delta y - \tilde{\omega} n \Delta t)), \quad (84)$$

as solution. Letting the operators $D_t D_t$, $D_x D_x$, and $D_y D_y$ act on $u_{p,q}^n$ results in

$$\frac{4}{\Delta t^2} \sin^2\left(\frac{\tilde{\omega} \Delta t}{2}\right) = c^2 \frac{4}{\Delta x^2} \sin^2\left(\frac{k_x \Delta x}{2}\right) + c^2 \frac{4}{\Delta y^2} \sin^2\left(\frac{k_y \Delta y}{2}\right). \quad (85)$$

or

$$\sin^2\left(\frac{\tilde{\omega} \Delta t}{2}\right) = \frac{c^2 \Delta t^2}{\Delta x^2} \sin^2\left(\frac{k_x \Delta x}{2}\right) + \frac{c^2 \Delta t^2}{\Delta y^2} \sin^2\left(\frac{k_y \Delta y}{2}\right). \quad (86)$$

For a real-valued $\tilde{\omega}$ we must have that the right-hand side is less than or equal to unity in absolute value, requiring in general that

$$\frac{c^2 \Delta t^2}{\Delta x^2} + \frac{c^2 \Delta t^2}{\Delta y^2} \leq 1. \quad (87)$$

This gives the stability criterion, more commonly expressed directly in an inequality for the time step:

$$\Delta t \leq \frac{1}{c} \left(\frac{1}{\Delta x^2} + \frac{1}{\Delta y^2} \right)^{-1/2} \quad (88)$$

A similar, straightforward analysis for the 3D case implies

$$\Delta t \leq \frac{1}{c} \left(\frac{1}{\Delta x^2} + \frac{1}{\Delta y^2} + \frac{1}{\Delta z^2} \right)^{-1/2} \quad (89)$$

In the case of variable coefficient $c^2 = c^2(\mathbf{x})$, we must use the worst-case value

$$\bar{c} = \sqrt{\max_{\mathbf{x} \in \Omega} c^2(\mathbf{x})} \quad (90)$$

in the stability criteria. Often, especially in the variable wave velocity case, it is wise to introduce a safety factor $\beta \in (0, 1]$ too:

$$\Delta t \leq \beta \frac{1}{\bar{c}} \left(\frac{1}{\Delta x^2} + \frac{1}{\Delta y^2} + \frac{1}{\Delta z^2} \right)^{-1/2} \quad (91)$$

The exact numerical dispersion relation in 3D becomes, for constant c ,

$$\begin{aligned} \tilde{\omega} = \frac{2}{\tilde{\omega} \Delta t} \sin^{-1} & (c^2 \Delta t^2 (\Delta x^2 \sin^2(\frac{1}{2} k_x \Delta x) + \Delta y^2 \sin^2(\frac{1}{2} k_y \Delta y) \\ & + \Delta z^2 \sin^2(\frac{1}{2} k_z \Delta z))). \end{aligned} \quad (92)$$

12 Applications of wave equations

This section presents a range of wave equation models for different physical phenomena. Although many wave motion problems in physics can be modeled by the standard linear wave equation, or a similar formulation with a system of first-order equations, there are some exceptions. Perhaps the most important is water waves: these are modeled by the Laplace equation with time-dependent boundary conditions at the water surface (long water waves, however, can be approximated by a standard wave equation, see Section 12.7). Quantum mechanical waves constitute another example where the waves are governed by the Schrödinger equation and not a standard wave equation. Many wave phenomena also need to take nonlinear effects into account when the wave amplitude is significant. Shock waves in the air is a primary example.

The derivations in the following are very brief. Those with a firm background in continuum mechanics will probably have enough information to fill in the details, while other readers will hopefully get some impression of the physics and approximations involved when establishing wave equation models.

12.1 Waves on a string

Figure 7 shows a model we may use to derive the equation for waves on a string. The string is modeled as a set of discrete point masses (at mesh points) with elastic strings in between. The strings are at a high constant tension T . We let the mass at mesh point x_i be m_i . The displacement of this mass point in y direction is denoted by $u_i(t)$.

The motion of mass m_i is governed by Newton's second law of motion. The position of the mass at time t is $x_i\mathbf{i} + u_i(t)\mathbf{j}$, where \mathbf{i} and \mathbf{j} are unit vectors in the x and y direction, respectively. The acceleration is then $u_i''(t)\mathbf{j}$. Two forces are acting on the mass as indicated in Figure 7. The force \mathbf{T}^- acting toward the point x_{i-1} can be decomposed as

$$\mathbf{T}^- = -T \sin \phi \mathbf{i} - T \cos \phi \mathbf{j},$$

where ϕ is the angle between the force and the line $x = x_i$. Let $\Delta u_i = u_i - u_{i-1}$ and let $\Delta s_i = \sqrt{\Delta u_i^2 + (x_i - x_{i-1})^2}$ be the distance from mass m_{i-1} to mass m_i . It is seen that $\cos \phi = \Delta u_i / \Delta s_i$ and $\sin \phi = (x_i - x_{i-1}) / \Delta s_i$ or $\Delta x / \Delta s_i$ if we introduce a constant mesh spacing $\Delta x = x_i - x_{i-1}$. The force can then be written

$$\mathbf{T}^- = -T \frac{\Delta x}{\Delta s_i} \mathbf{i} - T \frac{\Delta u_i}{\Delta s_i} \mathbf{j}.$$

The force \mathbf{T}^+ acting toward x_{i+1} can be calculated in a similar way:

$$\mathbf{T}^+ = T \frac{\Delta x}{\Delta s_{i+1}} \mathbf{i} + T \frac{\Delta u_{i+1}}{\Delta s_{i+1}} \mathbf{j}.$$

Newton's second law becomes




Figure 7: Discrete string model with point masses connected by elastic strings.

$$m_i u_i''(t) \mathbf{j} = \mathbf{T}^+ + \mathbf{T}^-,$$

which gives the component equations

$$T \frac{\Delta x}{\Delta s_i} = T \frac{\Delta x}{\Delta s_{i+1}}, \quad (93)$$

$$m_i u_i''(t) = T \frac{\Delta u_{i+1}}{\Delta s_{i+1}} - T \frac{\Delta u_i}{\Delta s_i}. \quad (94)$$

A basic assumption now is small displacements u_i and small displacement gradients $\Delta u_i / \Delta x$. For small $g = \Delta u_i / \Delta x$ we have that

$$\Delta s_i = \sqrt{\Delta u_i^2 + \Delta x^2} = \Delta x \sqrt{1 + g^2} + \Delta x(1 + \frac{1}{2}g^2 + \mathcal{O}(g^4)) \approx \Delta x.$$

Equation (93) is then simply the identity $T = T$, while (94) can be written as

$$m_i u_i''(t) = T \frac{\Delta u_{i+1}}{\Delta x} - T \frac{\Delta u_i}{\Delta x},$$

which upon division by Δx and introducing the density $\varrho_i = m_i/\Delta x$ becomes

$$\varrho_i u_i''(t) = T \frac{1}{\Delta x^2} (u_{i+1} - 2u_i + u_{i-1}). \quad (95)$$

We can now choose to approximate u_i'' by a finite difference in time and get the discretized wave equation,

$$\varrho_i \frac{1}{\Delta t^2} (u_i^{n+1} - 2u_i^n + u_i^{n-1}) = T \frac{1}{\Delta x^2} (u_{i+1} - 2u_i + u_{i-1}), \quad (96)$$

or we may go to the continuum limit $\Delta x \rightarrow 0$ and replace $u_i(t)$ by $u(x, t)$, ϱ_i by $\varrho(x)$, and recognize that the right-hand side of (95) approaches $\partial^2 u / \partial x^2$ as $\Delta x \rightarrow 0$. We end up with the continuous model for waves on a string:

$$\varrho \frac{\partial^2 u}{\partial t^2} = T \frac{\partial^2 u}{\partial x^2}. \quad (97)$$

Note that the density ϱ may change along the string, while the tension T is a constant. With variable wave velocity $c(x) = \sqrt{T/\varrho(x)}$ we can write the wave equation in the more standard form

$$\frac{\partial^2 u}{\partial t^2} = c^2(x) \frac{\partial^2 u}{\partial x^2}. \quad (98)$$

Because of the way ϱ enters the equations, the variable wave velocity does *not* appear inside the derivatives as in many other versions of the wave equation. However, most strings of interest have constant ϱ .

Normally, the end point of a string are fixed so that the displacement u is zero. The boundary conditions are therefore $u = 0$.

External forcing. It is easy to include an external force acting on the string. Say we have a vertical force $\tilde{f}_i \mathbf{j}$ acting on mass m_i . This force affects the vertical component of Newton's law and gives rise to an extra term $\tilde{f}(x, t)$ on the right-hand side of (97). In the model (98) we would add a term $f(x, t) = \tilde{f}(x, y)/\varrho(x)$.

Modeling the tension via springs. We assumed, in the derivation above, that the tension in the string, T , was constant. It is easy to check this assumption by modeling the string segments between the masses as standard springs, where the force (tension T) is proportional to the elongation of the spring segment. Let k be the spring constant, and set $T_i = k\Delta\ell$ for the tension in the

spring segment between x_{i-1} and x_i , where $\Delta\ell$ is the elongation of this segment from the tension-free state. A basic feature of a string is that it has high tension in the equilibrium position $u = 0$. Let the string segment have an elongation $\Delta\ell_0$ in the equilibrium position. After deformation of the string, the elongation is $\Delta\ell = \Delta\ell_0 + \Delta s_i$: $T_i = k(\Delta\ell_0 + \Delta s_i) \approx k(\Delta\ell_0 + \Delta x)$. This shows that T_i is independent of i . Moreover, the extra approximate elongation Δx is very small compared to $\Delta\ell_0$, so we may well set $T_i = T = k\Delta\ell_0$. This means that the tension is completely dominated by the initial tension determined by the tuning of the string. The additional deformations of the spring during the vibrations do not introduce significant changes in the tension.

12.2 Waves on a membrane

12.3 Elastic waves in a rod

Consider an elastic rod subject to a hammer impact at the end. This experiment will give rise to an elastic deformation pulse that travels through the rod. A mathematical model for longitudinal waves along an elastic rod starts with the general equation for deformations and stresses in an elastic medium,

$$\varrho \mathbf{u}_{tt} = \nabla \cdot \boldsymbol{\sigma} + \varrho \mathbf{f}, \quad (99)$$

where ϱ is the density, \mathbf{u} the displacement field, $\boldsymbol{\sigma}$ the stress tensor, and \mathbf{f} body forces. The latter has normally no impact on elastic waves.

For stationary deformation of an elastic rod, one has that $\sigma_{xx} = Eu_x$ and all other stress components are zero. Moreover, $\mathbf{u} = u(x)\mathbf{i}$. The parameter E is known as Young's modulus. Assuming that this simple stress and deformation field, which is exact in the stationary case, is a good approximation in the transient case with wave motion, (99) simplifies to

$$\varrho \frac{\partial^2 u}{\partial t^2} = \frac{\partial}{\partial x} \left(E \frac{\partial u}{\partial x} \right). \quad (100)$$

The associated boundary conditions are u or σ_{xx} known, typically $u = 0$ for a clamped end and $\sigma_{xx} = 0$ for a free end.

12.4 The acoustic model for seismic waves

Seismic waves are used to infer properties of subsurface geological structures. The physical model is a heterogeneous elastic medium where sound is propagated by small elastic vibrations. The general mathematical model for deformations in an elastic medium is based on Newton's second law,

$$\varrho \mathbf{u}_{tt} = \nabla \cdot \boldsymbol{\sigma} + \varrho \mathbf{f}, \quad (101)$$

and a constitutive law relating $\boldsymbol{\sigma}$ to \mathbf{u} , often Hooke's generalized law,

$$\boldsymbol{\sigma} = K \nabla \cdot \mathbf{u} \mathbf{I} + G(\nabla \mathbf{u} + (\nabla \mathbf{u})^T - \frac{2}{3} \nabla \cdot \mathbf{u} \mathbf{I}). \quad (102)$$

Here, \mathbf{u} is the displacement field, $\boldsymbol{\sigma}$ is the stress tensor, \mathbf{I} is the identity tensor, ϱ is the medium's density, \mathbf{f} are body forces (such as gravity), K is the medium's bulk modulus and G is the associated shear modulus. All these quantities may vary in space, while \mathbf{u} and $\boldsymbol{\sigma}$ will also show significant variation in time during wave motion.

The acoustic approximation to elastic waves arises from a basic assumption that the second term in Hooke's law, representing the deformations that give rise to shear stresses, can be neglected. This assumption can be interpreted as approximating the geological medium by a fluid. Neglecting also the body forces \mathbf{f} , (101) becomes

$$\varrho \mathbf{u}_{tt} = \nabla(K \nabla \cdot \mathbf{u}) \quad (103)$$

Introducing p as a pressure via

$$p = -K \cdot \nabla \mathbf{u}, \quad (104)$$

and dividing (103) by ϱ , we get

$$\mathbf{u}_{tt} = -\frac{1}{\varrho} \nabla p. \quad (105)$$

Taking the divergence of this equation, using $\nabla \cdot \mathbf{u} = -p/K$, gives the *acoustic approximation to elastic waves*:

$$p_{tt} = K \nabla \cdot \left(\frac{1}{\varrho} \nabla p \right). \quad (106)$$

This is a standard, linear wave equation with variable coefficients. It is common to add a source term $s(x, y, z, t)$ to model the generation of sound waves:

$$p_{tt} = K \nabla \cdot \left(\frac{1}{\varrho} \nabla p \right) + s. \quad (107)$$

A common further approximation of (107) is based on using the chain rule on the right-hand side,

$$K \nabla \cdot \left(\frac{1}{\varrho} \nabla p \right) = \frac{K}{\varrho} \nabla^2 p + K \nabla \left(\frac{1}{\varrho} \right) \cdot \nabla p \approx \frac{K}{\varrho} \nabla^2 p,$$

in that one assumes the relative spatial gradient $\nabla \varrho^{-1} = -\varrho^{-2} \nabla \varrho$ to be small. This results in

$$p_{tt} = \frac{K}{\varrho} \nabla^2 p + s. \quad (108)$$

The acoustic approximations to seismic waves are used for sound waves in the ground, and the Earth's surface is then a boundary where p equals the atmospheric pressure p_0 such that the boundary condition becomes $p = p_0$.

Anisotropy. Quite often in geological materials, the effective wave velocity $c = \sqrt{K/\varrho}$ is different in different spatial directions because geological layers are compacted such that the properties in the horizontal and vertical direction differ. With z as the vertical coordinate, we can introduce a vertical wave velocity c_z and a horizontal wave velocity c_h , and generalize (108) to

$$p_{tt} = c_z^2 p_{zz} + c_h^2 (p_{xx} + p_{yy}) + s. \quad (109)$$

12.5 Sound waves in liquids and gases

Sound waves arise from pressure and density variations in fluids. The starting point of modeling sound waves is the basic equations for a compressible fluid where we omit viscous (frictional) forces, body forces (gravity, for instance), and temperature effects:

$$\varrho_t + \nabla \cdot (\varrho \mathbf{u}) = 0, \quad (110)$$

$$\varrho \mathbf{u}_t + \varrho \mathbf{u} \cdot \nabla \mathbf{u} = -\nabla p, \quad (111)$$

$$\varrho = \varrho(p). \quad (112)$$

These equations are often referred to as the Euler equations for the motion of a fluid. The parameters involved are the density ϱ , the velocity \mathbf{u} , and the pressure p . Equation (111) reflects mass balance, (110) is Newton's second law for a fluid, with frictional and body forces omitted, and (112) is a constitutive law relating density to pressure by thermodynamics considerations. A typical model for (112) is the so-called isentropic relation⁵, valid for adiabatic processes where there is not heat transfer:

$$\varrho = \varrho_0 \left(\frac{p}{p_0} \right)^{1/\gamma}. \quad (113)$$

Here, p_0 and ϱ_0 are reference values for p and ϱ , for instances when the fluid is at rest, and γ is the ratio of specific heat at constant pressure and constant volume ($\gamma = 5/3$ for air).

The key approximation in a mathematical model for sound waves is to assume that those waves are small perturbations to the density, pressure, and velocity. We therefore write

$$p = p_0 + \hat{p},$$

$$\varrho = \varrho_0 + \hat{\varrho},$$

$$\mathbf{u} = \hat{\mathbf{u}},$$

where we have decomposed the fields in a constant equilibrium value, corresponding to $\mathbf{u} = 0$, and a small perturbation marked with a hat symbol. By

⁵http://en.wikipedia.org/wiki/Isentropic_process

inserting these decompositions in (110) and (111, neglecting all product terms of small perturbations and/or their derivatives, and dropping the hat symbols, one gets the following linearized PDE system for the small perturbations in density, pressure, and velocity:

$$\varrho_t + \varrho_0 \nabla \cdot \mathbf{u} = 0, \quad (114)$$

$$\varrho_0 \mathbf{u}_t = -\nabla p. \quad (115)$$

Eliminating ϱ with the aid of a relation $\varrho = \varrho(p)$, expanded as $\varrho = \varrho'(p_0)p = \varrho_0/(\gamma p_0)$ for small perturbations, yields a PDE system for p and \mathbf{u} :

$$p_t + \gamma \frac{p_0}{\varrho_0} \nabla \cdot \mathbf{u} = 0, \quad (116)$$

$$\mathbf{u}_t = -\frac{1}{\varrho_0} \nabla p, \quad (117)$$

Taking the divergence of (117) and differentiating (116) with respect to time gives the possibility to easily eliminate $\nabla \cdot \mathbf{u}_t$ and arrive at a standard, linear wave equation for p :

$$p_{tt} = c^2 \nabla^2 p, \quad (118)$$

where $c = \sqrt{\gamma p_0 / \varrho_0}$ is the speed of sound in the fluid.

12.6 Spherical waves

Spherically symmetric three-dimensional waves propagate in the radial direction r only so that $u = u(r, t)$. The fully three-dimensional wave equation

$$\frac{\partial^2 u}{\partial t^2} = \nabla \cdot (c^2 \nabla u) + f$$

then reduces to the spherically symmetric wave equation

$$\frac{\partial^2 u}{\partial t^2} = \frac{1}{r^2} \frac{\partial}{\partial r} \left(c^2(r) r^2 \frac{\partial u}{\partial r} \right) + f(r), \quad r \in (0, R), \quad t > 0. \quad (119)$$

Assume that the wave velocity c is constant. One can easily show that the function $v(r, t) = ru(r, t)$ fulfills a standard wave equation in Cartesian coordinates. To this end, insert $u = v/r$ in

$$\frac{1}{r^2} \frac{\partial}{\partial r} \left(c^2(r) r^2 \frac{\partial u}{\partial r} \right)$$

to obtain

$$r \left(\frac{dc^2}{dr} \frac{\partial v}{\partial r} + c^2 \frac{\partial^2 v}{\partial r^2} \right) - \frac{dc^2}{dr} v.$$

The two terms in the parenthesis can be combined to

$$r \frac{\partial}{\partial r} \left(c^2 \frac{\partial v}{\partial r} \right),$$

which is recognized as the variable-coefficient Laplace operator in one Cartesian coordinate. The spherically symmetric wave equation in terms of $v(r, t)$ now becomes

$$\frac{\partial^2 u}{\partial t^2} = \frac{\partial}{\partial r} \left(c^2(r) \frac{\partial v}{\partial t} - \frac{1}{r} \frac{dc^2}{dr} v \right) + r f(r), \quad r \in (0, R), \quad t > 0. \quad (120)$$

In the case of constant wave velocity c , this equation reduces to the wave equation in a single Cartesian coordinate:

$$\frac{\partial^2 u}{\partial t^2} = \frac{\partial}{\partial r} \left(c^2(r) \frac{\partial v}{\partial t} \right) + r f(r), \quad r \in (0, R), \quad t > 0. \quad (121)$$

That is, any program for solving the one-dimensional wave equation in a Cartesian coordinate system can be used to solve (121), provided the source term is multiplied by the coordinate, and that we divide the Cartesian mesh solution by r to get the spherically symmetric solution. Moreover, if $r = 0$ is included in the domain, spherical symmetry demands that $\partial u / \partial r = 0$ at $r = 0$, which means that

$$\frac{\partial u}{\partial r} = \frac{1}{r^2} \left(r \frac{\partial v}{\partial r} - v \right) = 0, \quad r = 0,$$

implying $v(0, t) = 0$ as a necessary condition. For practical applications, we exclude $r = 0$ from the domain and assume that some boundary condition is assigned at $r = \epsilon$, for some $\epsilon > 0$.

12.7 The linear shallow water equations

The next example considers water waves whose wavelengths are much larger than the depth and whose wave amplitudes are small. This class of waves may be generated by catastrophic geophysical events, such as earthquakes at the sea bottom, landslides moving into water, or underwater slides (or a combination, as earthquakes frequently release avalanches of masses). For example, a subsea earthquake will normally have an extension of many kilometers but lift the water only a few meters. The wave length will have a size dictated by the earthquake area, which is much larger than the water depth, and compared to this wave length, an amplitude of a few meters is very small. The water is essentially a thin film, and mathematically we can average the problem in the vertical direction and approximate the 3D wave phenomenon by 2D PDEs. Instead of a moving water domain in three space dimensions, we get a horizontal 2D domain with an unknown function for the surface elevation and the water depth as a variable coefficient in the PDEs.

Let $\eta(x, y, t)$ be the elevation of the water surface, $H(x, y)$ the water depth corresponding to a flat surface ($\eta = 0$), $u(x, y, t)$ and $v(x, y, t)$ the depth-averaged horizontal velocities of the water. Mass and momentum balance of the water volume give rise to the PDEs involving these quantities:

$$\eta_t = -(Hu)_x - (Hv)_x \quad (122)$$

$$u_t = -g\eta_x, \quad (123)$$

$$v_t = -g\eta_y, \quad (124)$$

where g is the acceleration of gravity. Equation (122) corresponds to mass balance while the other two are derived from momentum balance (Newton's second law).

The initial conditions associated with (122)-(124) are η , u , and v prescribed at $t = 0$. A common condition is to have some water elevation $\eta = I(x, y)$ and assume that the surface is at rest: $u = v = 0$. A subsea earthquake usually means a sufficiently rapid motion of the bottom and the water volume to say that the bottom deformation is mirrored at the water surface as an initial lift $I(x, y)$ and that $u = v = 0$.

Boundary conditions may be η prescribed for incoming, known waves, or zero normal velocity at reflecting boundaries (steep mountains, for instance): $un_x + vn_y = 0$, where (n_x, n_y) is the outward unit normal to the boundary. More sophisticated boundary conditions are needed when waves run up at the shore, and at open boundaries where we want the waves to leave the computational domain undisturbed.

Equations (122), (123), and (124) can be transformed to a standard, linear wave equation. First, multiply (123) and (124) by H , differentiate (123) with respect to x and (124) with respect to y . Second, differentiate (122) with respect to t and use that $(Hu)_{xt} = (Hu_t)_x$ and $(Hv)_{yt} = (Hv_t)_y$ when H is independent of t . Third, eliminate $(Hu_t)_x$ and $(Hv_t)_y$ with the aid of the other two differentiated equations. These manipulations results in a standard, linear wave equation for η :

$$\eta_{tt} = (gH\eta_x)_x + (gH\eta_y)_y = \nabla \cdot (gH\nabla\eta). \quad (125)$$

In the case we have an initial non-flat water surface at rest, the initial conditions become $\eta = I(x, y)$ and $\eta_t = 0$. The latter follows from (122) if $u = v = 0$, or simply from the fact that the vertical velocity of the surface is η_t , which is zero for a surface at rest.

The system (122)-(124) can be extended to handle a time-varying bottom topography, which is relevant for modeling long waves generated by underwater slides. In such cases the water depth function H is also a function of t , due to the moving slide, and one must add a time-derivative term H_t to the left-hand side of (122). A moving bottom is best described by introducing $z = H_0$ as the still-water level, $z = B(x, y, t)$ as the time- and space-varying bottom topography, so that $H = H_0 - B(x, y, t)$. In the elimination of u and v one may

assume that the dependence of H on t can be neglected in the terms $(Hu)_{xt}$ and $(Hv)_{yt}$. We then end up with a source term in (125), because of the moving (accelerating) bottom:

$$\eta_{tt} = \nabla \cdot (gH\nabla\eta) + B_{tt}. \quad (126)$$

The reduction of (126) to 1D, for long waves in a straight channel, or for approximately plane waves in the ocean, is trivial by assuming no change in y direction ($\partial/\partial y = 0$):

$$\eta_t = (gH\eta_x)_x + B_{tt}. \quad (127)$$

Wind drag on the surface. Surface waves are influenced by the drag of the wind, and if the wind velocity some meters above the surface is (U, V) , the wind drag gives contributions $C_V\sqrt{U^2 + V^2}U$ and $C_V\sqrt{U^2 + V^2}V$ to (123) and (124), respectively, on the right-hand sides.

Bottom drag. The waves will experience a drag from the bottom, often roughly modeled by a term similar to the wind drag: $C_B\sqrt{u^2 + v^2}u$ on the right-hand side of (123) and $C_B\sqrt{u^2 + v^2}v$ on the right-hand side of (124). Note that in this case the PDEs (123) and (124) become nonlinear and the elimination of u and v to arrive at a 2nd-order wave equation for η is not possible anymore.

Effect of the Earth's rotation. Long geophysical waves will often be affected by the rotation of the Earth because of the Coriolis force. This force gives rise to a term fv on the right-hand side of (123) and $-fu$ on the right-hand side of (124). Also in this case one cannot eliminate u and v to work with a single equation for η . The Coriolis parameter is $f = 2\Omega \sin \phi$, where Ω is the angular velocity of the earth and ϕ is the latitude.

12.8 Waves in blood vessels

The flow of blood in our bodies is basically fluid flow in a network of pipes. Unlike rigid pipes, the walls in the blood vessels are elastic and will increase their diameter when the pressure rises. The elastic forces will then push the wall back and accelerate the fluid. This interaction between the flow of blood and the deformation of the vessel wall results in waves traveling along our blood vessels.

A model for one-dimensional waves along blood vessels can be derived from averaging the fluid flow over the cross section of the blood vessels. Let x be a coordinate along the blood vessel and assume that all cross sections are circular, though with different radius $R(x, t)$. The main quantities to compute is the cross section area $A(x, t)$, the averaged pressure $P(x, t)$, and the total volume flux $Q(x, t)$. The area of this cross section is

$$A(x, t) = 2\pi \int_0^{R(x, t)} r dr, \quad (128)$$

Let $v_x(x, t)$ be the velocity of blood averaged over the cross section at point x . The volume flux, being the total volume of blood passing a cross section per time unit, becomes

$$Q(x, t) = A(x, t)v_x(x, t) \quad (129)$$

Mass balance and Newton's second law lead to the PDEs

$$\frac{\partial A}{\partial t} + \frac{\partial Q}{\partial x} = 0, \quad (130)$$

$$\frac{\partial Q}{\partial t} + \frac{\gamma + 2}{\gamma + 1} \frac{\partial}{\partial x} \left(\frac{Q^2}{A} \right) + \frac{A}{\varrho} \frac{\partial P}{\partial x} = -2\pi(\gamma + 2) \frac{\mu}{\varrho} \frac{Q}{A}, \quad (131)$$

where γ is a parameter related to the velocity profile, ϱ is the density of blood, and μ is the dynamic viscosity of blood.

We have three unknowns A , Q , and P , and two equations (130) and (131). A third equation is needed to relate the flow to the deformations of the wall. A common form for this equation is

$$\frac{\partial P}{\partial t} + \frac{1}{C} \frac{\partial Q}{\partial x} = 0, \quad (132)$$

where C is the compliance of the wall, given by the constitutive relation

$$C = \frac{\partial A}{\partial P} + \frac{\partial A}{\partial t}, \quad (133)$$

which require a relationship between A and P . One common model is to view the vessel wall, locally, as a thin elastic tube subject to an internal pressure. This gives the relation

$$P = P_0 + \frac{\pi h E}{(1 - \nu^2) A_0} (\sqrt{A} - \sqrt{A_0}),$$

where P_0 and A_0 are corresponding reference values when the wall is not deformed, h is the thickness of the wall, and E and ν are Young's modulus and Poisson's ratio of the elastic material in the wall. The derivative becomes

$$C = \frac{\partial A}{\partial P} = \frac{2(1 - \nu^2) A_0}{\pi h E} \sqrt{A_0} + 2 \left(\frac{(1 - \nu^2) A_0}{\pi h E} \right)^2 (P - P_0). \quad (134)$$

Another (nonlinear) deformation model of the wall, which has a better fit with experiments, is

$$P = P_0 \exp(\beta(A/A_0 - 1)),$$

where β is some parameter to be estimated. This law leads to

$$C = \frac{\partial A}{\partial P} = \frac{A_0}{\beta P}. \quad (135)$$

Reduction to standard wave equation. It is not uncommon to neglect the viscous term on the right-hand side of (131) and also the quadratic term with Q^2 on the left-hand side. The reduced equations (131) and (132) form a first-order linear wave equation system:

$$C \frac{\partial P}{\partial t} = -\frac{\partial Q}{\partial x}, \quad (136)$$

$$\frac{\partial Q}{\partial t} = -\frac{A}{\rho} \frac{\partial P}{\partial x}. \quad (137)$$

These can be combined into standard 1D wave equation PDE by differentiating the first equation with respect t and the second with respect to x ,

$$\frac{\partial}{\partial t} \left(C C \frac{\partial P}{\partial t} \right) = \frac{\partial}{\partial x} \left(\frac{A}{\rho} \frac{\partial P}{\partial x} \right),$$

which can be approximated by

$$\frac{\partial^2 Q}{\partial t^2} = c^2 \frac{\partial^2 Q}{\partial x^2}, \quad c = \sqrt{\frac{A}{\rho C}}, \quad (138)$$

where the A and C in the expression for c are taken as constant reference values.

12.9 Electromagnetic waves

Light and radio waves are governed by standard wave equations arising from Maxwell's general equations. When there are no charges and no currents, as in a vacuum, Maxwell's equations take the form

$$\nabla \cdot \mathbf{E} = 0,$$

$$\nabla \cdot \mathbf{B} = 0,$$

$$\nabla \times \mathbf{E} = -\frac{\partial \mathbf{B}}{\partial t},$$

$$\nabla \times \mathbf{B} = \mu_0 \epsilon_0 \frac{\partial \mathbf{E}}{\partial t},$$

where $\epsilon_0 = 8.854187817620 \cdot 10^{-12}$ (F/m) is the permittivity of free space, also known as the electric constant, and $\mu_0 = 1.2566370614 \cdot 10^{-6}$ (H/m) is the permeability of free space, also known as the magnetic constant. Taking the curl of the two last equations and using the identity

$$\nabla \times (\nabla \times \mathbf{E}) = \nabla(\nabla \cdot \mathbf{E}) - \nabla^2 \mathbf{E} = -\nabla^2 \mathbf{E} \text{ when } \nabla \cdot \mathbf{E} = 0,$$

immediately gives the wave equation governing the electric and magnetic field:

$$\frac{\partial^2 \mathbf{E}}{\partial t^2} = c^2 \frac{\partial^2 \mathbf{E}}{\partial x^2}, \quad (139)$$

$$\frac{\partial^2 \mathbf{B}}{\partial t^2} = c^2 \frac{\partial^2 \mathbf{B}}{\partial x^2}, \quad (140)$$

with $c = 1/\sqrt{\mu_0\epsilon_0}$ as the velocity of light. Each component of \mathbf{E} and \mathbf{B} fulfills a wave equation and can hence be solved independently.

13 Exercises

Exercise 11: Simulate elastic waves in a rod

A hammer hits the end of an elastic rod. The exercise is to simulate the resulting wave motion using the model (100) from Section 12.3. Let the rod have length L and let the boundary $x = L$ be stress free so that $\sigma_{xx} = 0$, implying that $\partial u / \partial x = 0$. The left end $x = 0$ is subject to a strong stress pulse (the hammer), modeled as

$$\sigma_{xx}(t) = \begin{cases} S, & 0 < t \leq t_s, \\ 0, & t > t_s \end{cases}$$

The corresponding condition on u becomes $u_x = S/E$ for $t \leq t_s$ and zero afterwards (recall that $\sigma_{xx} = Eu_x$). This is a non-homogeneous Neumann condition, and you will need to approximate this condition and combine it with the scheme (the ideas and manipulations follow closely the handling of a non-zero initial condition $u_t = V$ in wave PDEs or the corresponding second-order ODEs for vibrations). Filename: `rod.py`.

Exercise 12: Test the efficiency of compiled loops in 3D

Extend the `wave2D_u0.py` code and the Cython, Fortran, and C versions to 3D. Set up an efficiency experiment to see the relative efficiency of pure scalar Python code, vectorized code, Cython-compiled loops, Fortran-compiled loops, and C-compiled loops. Normalize the CPU time for each mesh by the fastest version. Filename: `wave3D_u0.py`.

Exercise 13: Earthquake-generated tsunami in a 1D model

A subsea earthquake leads to an immediate lift of the surface, see Figure 8. The initial surface shape $I(x)$ is symmetric around $x = 0$ and will split into two tsunamis, one traveling to the right and one to the left, as depicted in Figure 9. Since the water surface will remain symmetric with respect to $x = 0$, given that the outgoing wave to the left does not come back due to reflection, we insert a boundary $x = 0$ and impose a symmetry condition there: $\partial \eta / \partial x = 0$, where $\eta(x, t)$ is the elevation of the water surface. We are not interested in what

happens with the right-going wave after it hits the right boundary, so whether we impose $\eta = 0$ or $\partial\eta/\partial x = 0$ at that boundary is not of importance.

The shape of the initial surface can be taken as a Gaussian function,

$$I(x; I_0, I_a, I_m, I_s) = I_0 + I_a \exp\left(-\left(\frac{x - I_m}{I_s}\right)^2\right), \quad (141)$$

with $I_m = 0$ reflecting the location of the peak of $I(x)$ and I_s being a measure of the width of the function $I(x)$ (I_s is $\sqrt{2}$ times the standard deviation of the familiar normal distribution curve).

Set up the relevant one-dimensional, linear, wave equation for η , assuming long waves of small amplitude in comparison with the depth, as described in Section 12.7. Import the `viz` method from the `wave1D.py` program and make a call to it to solve the present tsunami problem. The constant speed of the right-going wave is $c = \sqrt{gH}$ and use this quantity to determine a suitable time T for when the wave hits the right boundary and the simulation is to be stopped. An alternative is to check in `plot_u` if `u[-2]` is significantly different from 0 and then return `True` to stop the simulation.




Figure 8: Sketch of initial water surface due to a subsea earthquake.

Filename: `tsunami1D_flat.py`.

Exercise 14: Implement an open boundary condition

To enable the right-going wave in Exercise 13 to leave the computational domain and travel undisturbed through the boundary, one can in a one-dimensional problem impose the following condition, called a *radiation condition* or *open boundary condition*:

$$\frac{\partial\eta}{\partial t} + c \frac{\partial\eta}{\partial x} = 0, \quad (142)$$




Figure 9: An initial surface elevation is split into two waves.

at the right boundary $x = x_R$. The parameter c is the wave velocity, which for the model in Exercise 13 is $c = \sqrt{gH(x_R)}$.

Show that (142) accepts a solution $\eta = g_R(x - ct)$, but not $\eta = g_L(x + ct)$. This means that (142) will allow any right-going wave $g_R(x - ct)$ pass through the boundary.

The condition (142) can be discretized by centered differences at the spatial end point $i = N_x$, corresponding to $x = x_R$:

$$[D_{2t}\eta + cD_{2x}\eta = 0]_{N_x}^n. \quad (143)$$

Eliminate the fictitious value $\eta_{N_x+1}^n$ by using the discrete equation at the same point (n, N_x) . The equation for the first step, η_i^1 , is in principal affected, but we can then use the condition $\eta_{N_x}^1 = 0$ since the wave has not yet reached the right boundary.

Modify the `solver` function in the `wave1D.py` program to incorporate the condition (143). Demonstrate that the tsunami travels through the domain and out of the right boundary without leaving any reflections behind. Make a nose test for checking that after a certain time T , the surface is flat.

Remark 1. The condition (142) works perfectly in 1D when c is known. In 2D and 3D, however, the condition reads $\eta_t + c_x\eta_x + c_y\eta_y = 0$, where c_x and c_y are the wave speeds in the x and y directions, and estimating these components (i.e., the direction of the wave) is often challenging. Other methods are normally used in 2D and 3D to let waves move out of a computational domain.

Remark 2. A radiation or open boundary condition at the left boundary takes the same form as (142), except that there is a minus sign in front of the $c\eta_x$

term. One can easily show that with this sign, the condition accepts left-going waves of the form $\eta = g_L(x + ct)$. Filename: `wave1D_Ropen.py`.

Exercise 15: Earthquake-generated tsunami over a subsea hill

We consider the same problem as in Exercise 13, but now there is a hill at the sea bottom, see Figure 10. The wave speed $c = \sqrt{gH(x)} = \sqrt{g(H_0 - B(x))}$ will then be reduced in the shallow water above the hill.




Figure 10: Sketch of an earthquake-generated tsunami passing over a subsea hill.

One possible form of the hill is a Gaussian function,

$$B(x; B_0, B_a, B_m, B_s) = B_0 + B_a \exp\left(-\left(\frac{x - B_m}{B_s}\right)^2\right), \quad (144)$$

but many other shapes are also possible, e.g., a "cosine hat" where

$$B(x; B_0, B_a, B_m, B_s) = B_0 + B_a \cos\left(\pi \frac{x - B_m}{2B_s}\right), \quad (145)$$

when $x \in [B_m - B_s, B_m + B_s]$ while $B = B_0$ outside this interval.

Also an abrupt construction may be tried:

$$B(x; B_0, B_a, B_m, B_s) = B_0 + B_a, \quad (146)$$

for $x \in [B_m - B_s, B_m + B_s]$ while $B = B_0$ outside this interval.

Visualize both the bottom topography and the water surface elevation (this requires modifying `plot_u`). Allow for a flexible choice of bottom shape, (144), (145), (146), or $B(x) = B_0$ (flat) and see if the waves become qualitatively different. Also investigate the amount of numerical noise that is triggered by rapid

changes in the bottom function and a small water gap at the top of the hill, and how this noise varies with the mesh resolution Δx . Use either the open boundary condition from Exercise 14, or set $\eta = 0$ at the right boundary and stop the simulation when the wave hits this boundary. Filename: `tsunami1D_hill.py`.

Exercise 16: Implement Neumann conditions in 2D

Modify the `wave2D_u0.py` program, which solves the 2D wave equation $u_{tt} = c^2(u_{xx} + u_{yy})$ with constant wave velocity c and $u = 0$ on the boundary, to have Neumann boundary conditions: $\partial u / \partial n = 0$. Include both scalar code (for debugging and reference) and vectorized code (for speed).

To test the code, use $u = 1.2$ as solution ($I(x, y) = 1.2$, $V = f = 0$, and c arbitrary), which should be exactly reproduced with any mesh as long as the stability criterion is satisfied. Another test is to use the plug-shaped pulse in the `pulse` function from Section 7 and the `wave1D_dn_vc.py` program. This pulse is exactly propagated in 1D if $c\Delta t / \Delta x = 1$. Check that also the 2D program can propagate this pulse exactly in x direction ($c\Delta t / \Delta x = 1$, Δy arbitrary) and y direction ($c\Delta t / \Delta y = 1$, Δx arbitrary). Filename: `wave2D_n.py`.

Exercise 17: Implement a convergence test for a 2D code

Use the following manufactured solution to verify a 2D code for $u_{tt} = c^2(u_{xx} + u_{yy})$ in the spatial domain $[0, L_x] \times [0, L_y]$, with $\partial u / \partial n$ on the boundary (cf. Exercise 16):

$$u_e(x, y, t) = \cos(m_x x \pi / L_x) \cos(m_y y \pi / L_y) \cos(\omega t), \quad (147)$$

Here, m_x and m_y are freely chosen integers such that the wave lengths in the x and y directions become $2L_x/m_x$ and $2L_y/m_y$, respectively. The parameter ω is calculated by inserting (147) in the wave equation. The solution (147) is a *standing wave* with $\partial u / \partial n = 0$. This u_e is not an exact solution of the discrete equations so the test must be based on empirical analysis of the convergence. The error E is assumed to behave like

$$E = C_t \Delta t^2 + C_x \Delta x^2 + C_y \Delta y^2,$$

for some constants C_t , C_x , and C_y . Choose $\Delta t = F_t h$, $\Delta x = F_x h$, and $\Delta y = F_y h$, where h is a common discretization parameter to be varied ($h \rightarrow 0$) and F_t , F_x , and F_y are freely chosen constant factors compatible with the stability criterion in 2D. The error can then be expressed as

$$E = C h^2,$$

where $C = C_x F_t^2 + C_y F_x^2 + C_t F_t^2$. Perform experiments with decreasing h , compute E , and verify that E/h^2 is approximately constant. Filename: `wave2D_n2.py`.

Exercise 18: Earthquake-generated tsunami over a 3D hill

This exercise extends Exercise 15 to a three-dimensional wave phenomenon, governed by the 2D PDE (125). We assume that the earthquake arise from a fault along the line $x = 0$ in the xy -plane so that the initial lift of the surface can be taken as $I(x)$ in Exercise 15. That is, a plan wave is propagating to the right, but will experience bending because of the bottom.

The bottom shape is now a function of x and y . An "elliptic" Gauss function in two dimensions, with its peak at (B_{mx}, B_{my}) , generalizes (144):

$$B(x; B_0, B_a, B_{mx}, B_{my}, B_s, b) = B_0 + B_a \exp \left(- \left(\frac{x - B_{mx}}{B_s} \right)^2 - \left(\frac{y - B_{my}}{bB_s} \right)^2 \right), \quad (148)$$

where b is a scaling parameter: $b = 1$ gives a circular Gaussian function with circular contour lines, while $b \neq 1$ gives an elliptic shape with elliptic contour lines.

The "cosine hat" (145) can also be generalized to

$$B(x; B_0, B_a, B_{mx}, B_{my}, B_s) = B_0 + B_a \cos \left(\pi \frac{x - B_{mx}}{2B_s} \right) \cos \left(\pi \frac{y - B_{my}}{2B_s} \right), \quad (149)$$

when $0 \leq \sqrt{x^2 + y^2} \leq B_s$ and $B = B_0$ outside this circle.

A box-shaped obstacle means that

$$B(x; B_0, B_a, B_m, B_s, b) = B_0 + B_a \quad (150)$$

for x and y inside a rectangle

$$B_{mx} - B_s \leq x \leq B_{mx} + B_s, \quad B_{my} - bB_s \leq y \leq B_{my} + bB_s,$$

and $B = B_0$ outside this rectangle. The b parameter controls the rectangular shape of the cross section of the box.

Note that the initial condition and the listed bottom shapes are symmetric around the line $y = B_{my}$. We therefore expect the surface elevation also to be symmetric with respect to this line. This means that we can halve the computational domain by working with $[0, L_x] \times [0, B_{my}]$. Along the upper boundary, $y = B_{my}$, we must impose the symmetry condition $\partial\eta/\partial n = 0$. Such a symmetry condition ($-\eta_x = 0$) is also needed at the $x = 0$ boundary because the initial condition has a symmetry here. At the lower boundary $y = 0$ we also set a Neumann condition (which becomes $-\eta_y = 0$). At the right boundary $x = L_x$ one can either implement a radiation (or open boundary) condition as in Exercise 14 or just set $\eta = 0$ or use a reflecting condition $\partial\eta/\partial n = \eta_x = 0$.

Visualize the surface elevation. Investigate how different hill shapes, different sizes of the water gap above the hill, and different resolutions Δx and Δt influence the numerical quality of the solution. Filename: `tsunami2D_hill.py`.

Exercise 19: Implement loops in compiled languages

Extend the program from Exercise 18 such that the loops over mesh points, inside the time loop, are implemented in compiled languages. Consider implementations in Cython, Fortran via `f2py`, C via Cython, C via `f2py`, C/C++ via Instant, and C/C++ via `scipy.weave`. Perform efficiency experiments to investigate the relative performance of the various implementations. It is often advantageous to normalize CPU times by the fastest method on a given mesh. Filename: `tsunami3D_hill_compiled.py`.

Exercise 20: Write a complete program in Fortran or C

As an extension of Exercise 20, write the whole `tsunami3D_hill.py` code in Fortran or C to check if there is more to be won with respect to efficiency in large-scale problems. This exercise will also illustrate the difference in program development with Fortran/C and Python. Filename: `tsunami3D_hill.f`.

Exercise 21: Investigate Matplotlib for visualization

Play with native Matplotlib code for visualizing $\eta(x, y, t)$ from Exercise 18 and see if there are effective ways to visualize both $\eta(x, y, t)$ and $B(x, y)$. Filename: `tsunami2D_hill_mpl.py`.

Exercise 22: Investigate Mayavi for visualization

Explore the program Mayavi⁶ for visualizing cases from Exercise 18. Try to visualize both the surface and the subsea hill. Filename: `tsunami2D_hill_mayavi.py`.

Exercise 23: Investigate Paraview for visualization

Explore the program Paraview⁷ for visualizing cases from Exercise 18. Try to visualize both the surface and the subsea hill. Filename: `tsunami2D_hill_paraview.py`.

Exercise 24: Investigate OpenDX for visualization

Investigate the OpenDX⁸ program for visualizing η (and preferably B simultaneously) from Exercise 18. One may start with using OpenDX as backend for the SciTools code that is in `wave2D_v1.py`. Filename: `tsunami2D_hill_OpenDX.py`.

Exercise 25: Investigate harmonic vs arithmetic mean

In Exercise 15, investigate if there is significant difference between the harmonic mean (39) and the standard arithmetic mean (38). Use extreme cases for the

⁶<http://code.enthought.com/projects/mayavi/>

⁷<http://www.paraview.org/>

⁸<http://www.opendx.org/>

investigations where the subsea hill is close to the flat surface. Pay particular attention to the box-shaped obstruction (146) since it has discontinuities.

Repeat the investigations for the case of a box-shaped obstruction in the 2D model from Exercise 18.

Remark. With a small gap between the obstruction and the free surface, and with abrupt changes in the bottom shape, the model PDE does not necessarily describe the wave motion in an accurate or qualitatively correct way. Filename: `tsunami2D_hill_harmonic.py`.

Exercise 26: Simulate seismic waves in 2D

The goal of this exercise is to simulate seismic waves using the PDE model (109) in a 2D xz domain with geological layers. Introduce m horizontal layers of thickness h_i , $i = 0, \dots, m-1$. Inside layer number i we have a vertical wave velocity $c_{z,i}$ and a horizontal wave velocity $c_{h,i}$. Make a program for simulating such 2D waves. Test it on a case with 3 layers where

$$c_{z,0} = c_{z,1} = c_{z,2}, \quad c_{h,0} = c_{h,2}, \quad c_{h,1} \ll c_{h,0}.$$

Let s be a localized point source at the middle of the Earth's surface (the upper boundary) and investigate how the resulting wave travels through the medium. The source can be a localized Gaussian peak that oscillates in time for some time interval. Place the boundaries far enough from the expanding wave so that the boundary conditions do not disturb the wave. Then the type of boundary condition does not matter, except that we physically need to have $p = p_0$, where p_0 is the atmospheric pressure, at the upper boundary.

Filename: `seismic2D.py`.

Index

- array slices, 15
- C extension module, 41
- C/Python array storage, 46
- column-major ordering, 46
- Courant number, 54
- Cython, 39
- `cython -a` (Python-C translation in HTML), 41
- 1D, stability, 55
- 2D, implementation, 35
- waves
 - on a string, 2
- wrapper code, 43
- declaration of variables in Cython, 40
- Dirichlet conditions, 19
- `distutils`, 41
- Fortran array storage, 46
- Fortran subroutine, 43
- homogeneous Dirichlet conditions, 19
- homogeneous Neumann conditions, 19
- lambda function (Python), 17
- Neumann conditions, 19
- open boundary condition, 72
- radiation condition, 72
- row-major ordering, 46
- scalar code, 15
- `setup.py`, 41
- slice, 15
- stability criterion, 55
- stencil
 - 1D wave equation, 3
 - Neumann boundary, 19
- vectorization, 14
- wave equation
 - 1D, 2
 - 1D, analytical properties, 52
 - 1D, exact numerical solution, 53
 - 1D, finite difference method, 2
 - 1D, implementation, 9