Discretizing first-order ODEs by finite difference methods

Hans Petter Langtangen[1,2]

[1]Center for Biomedical Computing, Simula Research Laboratory
[2]Department of Informatics, University of Oslo

Sep 9, 2012

Note: **VERY PRELIMINARY VERSION!**

Contents

1

List of exercises

Finite difference methods for partial differential equations (PDEs) employ a range of concepts and tools that can be introduced and illustrated in the context of simple ordinary differential equation (ODE) examples. By first working with ODEs, we keep the mathematical problems to be solved as simple as possible (but no simpler), thereby allowing full focus on understanding the concepts and tools that will be reused and futher extended when addressing finite difference methods for time-dependent PDEs. The forthcoming treatment of ODEs is therefore solely dominated by reasoning and methods that directly carry over to numerical methods for PDEs.

We study two model problems: an ODE for a decaying phenomena, which will be relevant for PDEs of diffusive nature, and an ODE for oscillating phenomena, which will be relevant for PDEs of wave nature. Both problems are linear with known analytical solutions such that we can easily assess the quality of various numerical methods and analyze their behavior.

1 Finite difference methods for an ODE

The purpose of this module is to explain finite difference methods in detail for a simple ordinary differential equation (ODE). Emphasis is put on the reasoning when discretizing the problem, various ways of programming the methods, how to verify that the implementation is correct, experimental investigations of the numerical behavior of the methods, and theoretical analysis of the methods to explain the observations.

1.1 Exponential decay

Our model problem is perhaps the simplest ODE:

$$u'(t) = -au(t),$$

Here, $a > 0$ is a constant and $u'(t)$ means differentiation with respect to time t. This type of equation arises in a number of widely different phenomena where some quantity u undergoes exponential reduction. Examples include radioactive decay, population decay, investment decay, cooling of an object, pressure decay in the atmosphere, and retarded motion in fluids (for some of these models, a can be negative as well). Studying numerical solution methods for this simple ODE gives imporant insight that can be reused for diffusion PDEs.

The analytical solution of the ODE is found by the method of separation of variables, resulting in

$$u(t) = Ce^{-at},$$

for any arbitrary constant C. To formulate a mathematical problem for which there is a unique solution, we need a condition to fix the value of C. This condition is known as the *initial condition* and stated as $u(0) = I$. That is, we know the value I of u when the process starts at $t = 0$. The exact solution is then $u(t) = I \exp{(-at)}$.

We seek the solution $u(t)$ of the ODE for $t \in (0, T]$. The point $t = 0$ is not included since we know u here and assume that the equation governs u for $t > 0$. The complete ODE problem then reads: find $u(t)$ such that

$$u' = -au, \ t \in (0, T], \quad u(0) = I. \tag{1}$$

This is known as a *continuous problem* because the parameter t varies continuously from 0 to T. For each t we have a corresponding $u(t)$. There are hence infinitely many values of t and $u(t)$. The purpose of a numerical method is to

formulate a corresponding *discrete* problem whose solution is characterized by a finite number of values, which can be computed in a finite number of steps on a computer.

1.2   The Forward Euler scheme

Solving an ODE like (1) by a finite difference method consists of the following four steps:

1. discretizing the domain,

2. fulfilling the equation at discrete time points,

3. replacing derivatives by finite differences,

4. formulating a recursive algorithm.

Step 1: Discretizing the domain.   The time domain $[0, T]$ is represented by a finite number of $N + 1$ points

$$0 = t_0 < t_1 < t_2 < \cdots < t_{N-1} < t_N = T \,. \tag{2}$$

The collection of points $t_0, t_1, \ldots, t_N$ constitutes a *mesh* or *grid*. Often the mesh points will be uniformly spaced in the domain $[0, T]$, which means that the spacing $t_{n+1} - t_n$ is the same for all n. This spacing is then often denoted by Δt, in this case $t_n = n\Delta t$.

We seek the solution u at the mesh points: $u(t_n)$, $n = 1, 2, \ldots, N$ (note that u^0 is already known as I). A notational short-form for $u(t_n)$, which will be used extensively, is u^n. More precisely, we let u^n be the *numerical approximation* to the exact solution at $t = t_n$, $u(t_n)$. When we need to clearly distinguish the numerical and the exact solution, we often place a subscript e on the exact solution, as in $u_e(t_n)$. Figure 1 shows the t_n and u_n points for $n = 0, 1, \ldots, N = 7$ as well as $u_e(t)$ as the dashed line.

Since finite difference methods produce solutions at the mesh points only, it is an open question what the solution is between the mesh points. One can use methods for interpolation to compute the value of u between mesh points. The simplest (and most widely used) interpolation method is to assume that u varies linearly between the mesh points, see Figure 2. Given u^n and u^{n+1}, the value of u at some $t \in [t_n, t_{n+1}]$ is by linear interpolation

$$u(t) \approx u^n + \frac{u^{n+1} - u^n}{t_{n+1} - t_n}(t - t_n) \,. \tag{3}$$

Figure 1: Time mesh with discrete solution values.

Figure 2: Linear interpolation between the discrete solution values (dashed curve is exact solution).

Step 2: Fulfilling the equation at discrete time points. The ODE is supposed to hold for all $t \in (0, T]$, i.e., at an infinite number of points. Now we relax that requirement and require that the ODE is fulfilled at a finite set of discrete points in time. The mesh points $t_1, t_2, \ldots, t_N$ are a natural choice of points. The original ODE is then reduced to the following N equations:

6

$$u'(t_n) = -au(t_n), \quad n = 1, \ldots, N. \tag{4}$$

Step 3: Replacing derivatives by finite differences.   The next and most essential step of the method is to replace the derivative u' by a finite difference approximation. Let us first try a one-sided difference approximation (see Figure 3),

$$u'(t_n) \approx \frac{u^{n+1} - u^n}{t_{n+1} - t_n}. \tag{5}$$

Inserting this approximation in (4) results in

$$\frac{u^{n+1} - u^n}{t_{n+1} - t_n} = -au^n, \quad n = 0, 1, \ldots, N-1. \tag{6}$$

This equation is the discrete counterpart to the original ODE problem (1), and often known as a *finite difference scheme*, which yields a straightforward way to compute the solution at the mesh points $(u(t_n),\ n = 1, 2, \ldots, N)$ as shown next.

Figure 3:   Illustration of a forward difference.

Step 4: Formulating a recursive algorithm.   The final step is to identify the computational algorithm to be implemented in a program. The key observation here is to realize that (6) can be used to compute u^{n+1} if u^n is known. Starting with $n = 0$, u^0 is known since $u^0 = u(0) = I$, and (6) gives an equation

for u^1. Knowing u^1, u^2 can be found from (6). In general, u^n in (6) can be assumed known, and then we can easily solve for the unknown u^{n+1}:

$$u^{n+1} = u^n - a(t_{n+1} - t_n)u^n. \tag{7}$$

We shall refer to (7) as the Forward Euler (FE) scheme for our model problem. From a mathematical point of view, equations of the form (7) are known as *difference equations* since they express how differences in u, like $u^{n+1} - u^n$, evolve with n. The finite difference method can be viewed as a method for turning a differential equation into a difference equation.

Computation with (7) is straightforward:

$$u_0 = I,$$
$$u_1 = u^0 - a(t_1 - t_0)u^0 = I(1 - a(t_1 - t_0)),$$
$$u_2 = u^1 - a(t_2 - t_1)u^1 = I(1 - a(t_1 - t_0))(1 - a(t_2 - t_1)),$$
$$u^3 = u^2 - a(t_3 - t_2)u^2 = I(1 - a(t_1 - t_0))(1 - a(t_2 - t_1))(1 - a(t_3 - t_2)),$$

and so on until we reach u^N. In the case $t_{n+1} - t_n$ is a constant, denoted by Δt, we realize from the above calculations that

$$u_0 = I,$$
$$u_1 = I(1 - a\Delta t),$$
$$u_2 = I(1 - a\Delta t)^2,$$
$$u^3 = I(1 - a\Delta t)^3,$$
$$\vdots$$
$$u^N = I(1 - a\Delta t)^N.$$

This means that we have found a closed formula for u^n, and there is no need to let a computer generate the sequence $u^1, u^2, u^3, \ldots$. However, finding such a formula for u^n is possible only for a few very simple problems.

As the next sections will show, the scheme (7) is just one out of many alternative finite difference (and other) schemes for the model problem (1).

1.3   The Backward Euler scheme

There are many choices of difference approximations in step 3 of the finite difference method as presented in the previous section. Another alternative is

$$u'(t_n) \approx \frac{u^n - u^{n-1}}{t_n - t_{n-1}}. \tag{8}$$

Since this difference is based on going backward in time (t_{n-1}) for information, it is known as the Backward Euler difference. Figure 5 explains the idea.

Figure 4: Illustration of a backward difference.

Inserting (8) in (4) yields the Backward Euler (BE) scheme:

$$\frac{u^n - u^{n-1}}{t_n - t_{n-1}} = -au^n .$$

(9)

We assume, as explained under step 4 in Section 1.2, that we have computed $u^0, u^1, \ldots, u^{n-1}$ such that (9) can be used to compute u^n. For direct similarity with the Forward Euler scheme (7) we replace n by $n + 1$ in (9) and solve for the unknown value u^{n+1}:

$$u^{n+1} = \frac{1}{1 + a(t_{n+1} - t_n)} u^n .$$

(10)

1.4 The Crank-Nicolson scheme

The finite difference approximations used to derive the schemes (7) and (10) are both one-sided differences, known to be less accurate than central (or midpoint) differences. We shall now construct a central difference at $t_{n+1/2} = \frac{1}{2}(t_n + t_{n+1})$, or $t_{n+1/2} = (n + \frac{1}{2})\Delta t$ if the mesh spacing is uniform in time. The approximation reads

$$u'(t_{n+\frac{1}{2}}) \approx \frac{u^{n+1} - u^n}{t_{n+1} - t_n} .$$

(11)

Note that the fraction on the right-hand side is the same as for the Forward Euler approximation (5) and the Backward Euler approximation (8) (with n replaced by $n + 1$). The accuracy of this fraction as an approximation to the

9

derivative of u depends on *where* we seek the derivative: in the center of the interval $[t_{n+1}, t_n]$ or at the end points.

With the formula (11), where u' is evaluated at $t_{n+1/2}$, it is natural to demand the ODE to be fulfilled at the time points between the mesh points:

$$u'(t_{n+\frac{1}{2}}) = -au(t_{n+\frac{1}{2}}), \quad n = 0, \ldots, N-1. \tag{12}$$

Using (11) in (12) results in

$$\frac{u^{n+1} - u^n}{t_{n+1} - t_n} = -au^{n+\frac{1}{2}}, \tag{13}$$

where $u^{n+\frac{1}{2}}$ is a short form for $u(t_{n+\frac{1}{2}})$. The problem is that we aim to compute u^n for integer n, implying that $u^{n+\frac{1}{2}}$ is not a quantity computed by our method. It must be expressed by the quantities that we actually produce, i.e., u at the mesh points. One possibility is to approximate $u^{n+\frac{1}{2}}$ as an average of the u values at the neighboring mesh points:

$$u^{n+\frac{1}{2}} \approx \frac{1}{2}(u^n + u^{n+1}). \tag{14}$$

Using (14) in (13) results in

$$\frac{u^{n+1} - u^n}{t_{n+1} - t_n} = -a\frac{1}{2}(u^n + u^{n+1}). \tag{15}$$

Figure 5 sketches the geometric interpretation of such a centered difference.

Figure 5: Illustration of a centered difference.

We assume that u^n is already computed so that u^{n+1} is the unknown, which we can solve for:

$$u^{n+1} = \frac{1 - \frac{1}{2}a(t_{n+1} - t_n)}{1 + \frac{1}{2}a(t_{n+1} - t_n)} u^n. \tag{16}$$

The finite difference scheme (16) is known as the midpoint scheme or the Crank-Nicolson (CN) scheme. We shall use the latter name.

1.5 The unifying θ-rule

Let us reconsider the derivation of the Forward Euler, Backward Euler, and Crank-Nicolson schemes. In all the mentioned schemes we replace u' by the fraction

$$\frac{u^{n+1} - u^n}{t_{n+1} - t_n},$$

and the difference between the methods lies in which point this fraction approximates the derivative; i.e., in which point we sample the ODE. So far this has been the end points or the midpoint of $[t_n, t_{n+1}]$. However, we may choose any point $\tilde{t} \in [t_n, t_{n+1}]$. The difficulty is that evaluating the right-hand side $-au$ at an arbitrary point faces the same problem as in Section 1.4: the point value must be expressed by the discrete u quantities that we compute by the scheme, i.e., u^n and u^{n+1}. Following the averaging idea from Section 1.4, the value of u at an arbitrary point $\tilde{t}$ can be calculated as a *weighted average*, which generalizes the arithmetic average $\frac{1}{2}u^n + \frac{1}{2}u^{n+1}$. If we express $\tilde{t}$ as a weighted average

$$t_{n+\theta} = \theta t_{n+1} + (1 - \theta)t_n,$$

where $\theta \in [0, 1]$ is the weighting factor, we can write

$$u(\tilde{t}) = u(\theta t_{n+1} + (1 - \theta)t_n) \approx \theta u^{n+1} + (1 - \theta)u^n. \tag{17}$$

We can now let the ODE hold at the point $\tilde{t} \in [t_n, t_{n+1}]$, approximate u' by the fraction $(u^{n+1} - u^n)/(t_{n+1} - t_n)$, and approximate the right-hand side $-au$ by the weighted average (17). The result is

$$\frac{u^{n+1} - u^n}{t_{n+1} - t_n} = -a(\theta u^{n+1} + (1 - \theta)u^n). \tag{18}$$

This is a generalized scheme for our model problem: $\theta = 0$ gives the Forward Euler scheme, $\theta = 1$ gives the Backward Euler scheme, and $\theta = 1/2$ gives the Crank-Nicolson scheme. In addition, we may choose any other value of θ in $[0, 1]$.

As before, u^n is considered known and u^{n+1} unknown, so we solve for the latter:

$$u^{n+1} = \frac{1 - (1 - \theta)a(t_{n+1} - t_n)}{1 + \theta a(t_{n+1} - t_n)}. \tag{19}$$

This scheme is known as the θ-rule, or alternatively written as the "theta-rule".

1.6 Constant time step

All schemes up to now have been formulated for a general non-uniform mesh in time: $t_0, t_1, \ldots, t_N$. Non-uniform meshes are highly relevant since one can use many points in regions where u varies rapidly, and save points in regions where u is slowly varying. This is the key idea of *adaptive* methods where the spacing of the mesh points are determined as the computations proceed.

However, a uniformly distributed set of mesh points is very common and sufficient for many applications. It therefore makes sense to present the finite difference schemes for a uniform point distribution $t_n = n\Delta t$, where Δt is the constant spacing between the mesh points, also referred to as the *time step*. The resulting formulas look simpler and are perhaps more well known:

$$u^{n+1} = (1 - a\Delta t)u^n \qquad \text{FE} \qquad (20)$$

$$u^{n+1} = \frac{1}{1 + a\Delta t}u^n \qquad \text{BE} \qquad (21)$$

$$u^{n+1} = \frac{1 - \frac{1}{2}a\Delta t}{1 + \frac{1}{2}a\Delta t}u^n \qquad \text{CN} \qquad (22)$$

$$u^{n+1} = \frac{1 - (1 - \theta)a\Delta t}{1 + \theta a\Delta t}u^n \theta - \text{rule} \qquad (23)$$

Not surprisingly, we present alternative schemes because they have different pros and cons, both for the simple ODE in question (which can easily be solved as accurately as desired), and for more advanced differential equation problems.

1.7 Compact operator notation for finite differences

Finite difference formulas can be tedious to write and read, especially for differential equations with many terms and many derivatives. To save space and help the reader of the scheme to quickly see the nature of the difference approximations, we introduce a compact notation. A forward difference approximation is denoted by the D_t^+ operator:

$$[D_t^+ u]^n = \frac{u^{n+1} - u^n}{\Delta t} \approx \frac{d}{dt}u(t_n)\,. \qquad (24)$$

The notation consists of an operator that approximates differentiation with respect to an independent variable, here t. The operator is built of the symbol D, with the variable as subscript and a superscript denoting the type of difference. The superscript $^+$ indicates a forward difference. We place square brackets around the operator and the function it operates on and specify the mesh point, where the operator is acting, by a superscript.

The corresponding operator notation for a centered difference and a backward difference reads

$$[D_t u]^n = \frac{u^{n+\frac{1}{2}} - u^{n-\frac{1}{2}}}{\Delta t} \approx \frac{d}{dt} u(t_n), \tag{25}$$

and

$$[D_t^- u]^n = \frac{u^n - u^{n-1}}{\Delta t} \approx \frac{d}{dt} u(t_n). \tag{26}$$

Note that the superscript $^-$ denotes the backward difference, while no superscript implies a central difference.

An averaging operator is also convenient to have:

$$[\overline{u}^t]^n = \frac{1}{2}(u^{n-\frac{1}{2}} + u^{n+\frac{1}{2}}) \approx u(t_n) \tag{27}$$

The superscript t indicates that the average is taken along the time coordinate. The common average $(u^n + u^{n+1})/2$ can now be expressed as $[\overline{u}^t]^{n+1/2}$.

The Backward Euler finite difference approximation to $u' = -au$ can be written as follows utilizing the compact notation:

$$[D_t^- u]^n = -au^n.$$

In difference equations we often place the square brackets around the whole equation, to indicate at which mesh point the equation applies, since each term is supposed to be approximated at the same point:

$$[D_t^- u = -au]^n. \tag{28}$$

The Forward Euler scheme takes the form

$$[D_t^+ u = -au]^n, \tag{29}$$

while the Crank-Nicolson scheme is written as

$$[D_t u = -a\overline{u}^t]^{n+\frac{1}{2}}. \tag{30}$$

Just apply (25) and (27) and write out the expressions to see that (30) is indeed the Crank-Nicolson scheme.

The θ-rule can be specified by

$$[\bar{D}_t u = -a\overline{u}^{t,\theta}]^{n+\theta}, \tag{31}$$

if we define a new time difference and a *weighted averaging operator*:

$$[\bar{D}_t u]^{n+\theta} = \frac{u^{n+1} - u^n}{t^{n+1} - t^n}, \tag{32}$$

$$[\overline{u}^{t,\theta}]^{n+\theta} = (1-\theta)u^n + \theta u^{n+1} \approx u(t_{n+\theta}), \tag{33}$$

where $\theta \in [0, 1]$. Note that for $\theta = 1/2$ we recover the standard centered difference and the standard arithmetic average. The idea in (31) is to sample

13

the equation at $t_{n+\theta}$, use a skew difference at that point $[\bar{D}_t u]^{n+\theta}$, and a shifted mean value. An alternative notation is

$$[D_t u]^{n+1/2} = \theta[-au]^{n+1} + (1-\theta)[-au]^n.$$

Looking at the various examples above and comparing them with the underlying differential equations, we see immediately which difference approximations that have been used and at which point they apply. Therefore, the compact notation efficiently communicates the reasoning behind turning a differential equation into a difference equation.

2 Implementation

The purpose now is to make a computer program for solving

$$u'(t) = -au(t), \quad t \in (0, T], \quad u(0) = I,$$

and display the solution on the screen, preferably together with the exact solution. We shall also be concerned with how we can test that the implementation is correct.

All programs referred to in this section are found in the `src/decay` directory.

Mathematical problem. We want to explore the Forward Euler scheme, the Backward Euler, and the Crank-Nicolson schemes applied to our model problem. From an implementational points of view, it is advantageous to implement the θ-rule

$$u^{n+1} = \frac{1 - (1-\theta)a\Delta t}{1 + \theta a\Delta t} u^n,$$

since it can generate the three other schemes by various of choices of θ: $\theta = 0$ for Forward Euler, $\theta = 1$ for Backward Euler, and $\theta = 1/2$ for Crank-Nicolson. Given a, $u^0 = I$, T, and Δt, our task is to use the θ-rule to compute $u^1, u^2, \ldots, u^N$, where $t_N = N\Delta t$, and N the closest integer to $T/\Delta t$.

Computer Language: Python. Any programming language can be used to generate the u^{n+1} values from the formula above. However, in this document we shall mainly make use of Python of several reasons:

- Python has a very clean, readable syntax (often known as "executable pseudo-code").

- Python code is very similar to MATLAB code (and MATLAB has a particularly widespread use for scientific computing).

- Python is similar to, but much simpler to work with and results in more reliable code than C++.

- Python is a full-fledged, very powerful programming language.

- Python has a rich set of modules for scientific computing, and its popularity in scientific computing is rapidly growing.

- Python was made for being combined with compiled languages (C, C++, Fortran) to reuse existing numerical software and to reach high computational performance of new implementations.

- Python has extensive support for administrative task needed when doing large-scale computational investigations.

- Python has extensive support for graphics (visualization, user interfaces, web applications).

- FEniCS, a very powerful tool for solving PDEs by the finite element method, is most human-efficient to operate from Python.

Learning Python is easy. Many newcomers to the language will probably learn enough from the examples to perform their own computer experiments. The examples start with simple Python code and gradually make use of more powerful constructs as we proceed. As long as it is not inconvenient for the problem at hand, our Python code is made as close as possible to MATLAB code for easy transition between the two languages.

2.1 Making a program

We choose to have an array u for storing the u^n values, $n = 0, 1, \ldots, N$. The algorithmic steps are

1. initialize u^0

2. for $t = t_n$, $n = 1, 2, \ldots, N$: compute u_n using the θ-rule formula

Function for computing the numerical solution. The following Python function takes the input data of the problem $(I, a, T, \Delta t, \theta)$ as arguments and returns two arrays with the solution $u^0, \ldots, u^N$ and the mesh points $t_0, \ldots, t_N$, respectively:

```python
from numpy import *

def solver(I, a, T, dt, theta):
    """Solve u'=-a*u, u(0)=I, for t in (0,T] with steps of dt."""
    N = int(T/dt)            # no of time intervals
    T = N*dt                 # adjust T to fit time step dt
    u = zeros(N+1)           # array of u[n] values
    t = linspace(0, T, N+1)  # time mesh

    u[0] = I                 # assign initial condition
    for n in range(0, N):    # n=0,1,...,N-1
        u[n+1] = (1 - (1-theta)*a*dt)/(1 + theta*dt*a)*u[n]
    return u, t
```

The `numpy` library contains a lot of functions for array computing. Most of the function names are similar to what is found in the alternative scientific computing language MATLAB. Here we make use of

- `zeros(N+1)` for creating an array of a size `N+1` and initializing the elements to zero

- `linspace(0, T, N+1)` for creating an array with `N+1` coordinates uniformly distributed between `0` and `T`

The `for` loop deserves a comment, especially for newcomers to Python. The construction `range(0, N, s)` generates all integers from `0` to `N` in steps of `s`, *but not including* `N`. Omitting `s` means `s=1`. For example, `range(0, 6, 3)` gives `0` and `3`, while `range(0, N)` generates `0`, `1`, ..., `N-1`. In our loop, `n` takes on the values generated by `range(0, N)`, implying the following assignments `u[n+1]`: `u[1]`, `u[2]`, ..., `u[N]`, which is what we want since `u` has length `N+1`. The first index in Python arrays or lists is *always* `0` and the last is then `len(u)-1`.

To compute with the `solver` function, we need to *call* it. Here is a sample call:

```
u, t = solver(I=1, a=2, T=8, dt=0.8, theta=1)
```

Integer division. The shown implementation of the `solver` may face problems and wrong results if `T`, `a`, `dt`, and `theta` are given as integers, see Exercises 1 and 2. The problem is related to *integer division* in Python (as well as in Fortran, C, and C++): `1/2` becomes `0`, while `1.0/2`, `1/2.0`, or `1.0/2.0` all become are `0.5`. It is enough that at least the nominator or the denominator is a real number (i.e., a `float` object) to ensure correct mathematical division. Inserting a conversion `dt = float(dt)` guarantees that `dt` is `float` and avoids problems in Exercise 2.

Another problem with computing $N = T/\Delta t$ is that we should round N to the nearest integer. With `N = int(T/dt)` the `int` operation picks the largest integer smaller than `T/dt`. Correct rounding is obtained by

```
N = int(round(T/dt))
```

The complete version of our improved, safer `solver` function then becomes

```
from numpy import *

def solver(I, a, T, dt, theta):
    """Solve u'=-a*u, u(0)=I, for t in (0,T] with steps of dt."""
    dt = float(dt)            # avoid integer division
    N = int(round(T/dt))      # no of time intervals
    T = N*dt                  # adjust T to fit time step dt
    u = zeros(N+1)            # array of u[n] values
    t = linspace(0, T, N+1)   # time mesh
```

```
    u[0] = I                    # assign initial condition
    for n in range(0, N):      # n=0,1,...,N-1
        u[n+1] = (1 - (1-theta)*a*dt)/(1 + theta*dt*a)*u[n]
    return u, t
```

Doc strings. Right below the header line in the `solver` function there is a Python string enclosed in triple double quotes `"""`. The purpose of this string object is to document what the function does and what the arguments are. In this case the necessary documentation do not span more than one line, but with triple double quoted strings the text may span several lines:

```
def solver(I, a, T, dt, theta):
    """
    Solve

        u'(t) = -a*u(t),

    with initial condition u(0)=I, for t in the time interval
    (0,T]. The time interval is divided into time steps of
    length dt.

    theta=1 corresponds to the Backward Euler scheme, theta=0
    to the Forward Euler scheme, and theta=0.5 to the Crank-
    Nicolson method.
    """
    ...
```

Such documentation strings appearing right after the header of a function are called *doc strings*. There are tools that can automatically produce nicely formatted documentation by extracting the definition of functions and the contents of doc strings.

It is strongly recommended to equip any function whose purpose is not obvious with a doc string. Nevertheless, the forthcoming text deviates from this rule if the function is explained in the text.

Formatting of numbers. Having computed the discrete solution `u`, it is natural to look at the numbers:

```
# Write out a table of t and u values:
for i in range(len(t)):
    print t[i], u[i]
```

This compact `print` statement gives unfortunately quite ugly output because the `t` and `u` values are not aligned in nicely formatted columns. To fix this problem, we recommend to use the *printf format*, supported most programming languages inherited from C. Another choice is Python's recent *format string syntax*.

Writing `t[i]` and `u[i]` in two nicely formatted columns is done like this with the printf format:

```
print 't=%6.3f u=%g' % (t[i], u[i])
```

The percentage signs signify "slots" in the text where the variables listed at the end of the statement are inserted. For each "slot" one must specify a format for how the variable is going to appear in the string: `s` for pure text, `d` for an integer, `g` for a real number written as compactly as possible, `9.3E` for scientific notation with three decimals in a field of width 9 characters (e.g., `-1.351E-2`), or `.2f` for a standard decimal notation, here with two decimals, formatted with minimum width. The printf syntax provides a quick way of formatting tabular output of numbers with full control of the layout.

The alternative *format string syntax* looks like

```
print 't={t:6.3f} u={u:g}'.format(t=t[i], u=u[i])
```

As seen, this format allows logical names in the "slots" where `t[i]` and `u[i]` are to be inserted. The "slots" are surrounded by curly braces, and the logical name is followed by a colon and then the printf-like specification of how to format real numbers, integers, or strings.

Running the program. The function and main program shown above must be placed in a file, say with name `dc_v1.py`. Make sure you write the code with a suitable text editor (Gedit, Emacs, Vim, Notepad++, or similar). The program is run by executing the file this way:

<div align="center">Terminal</div>

```
Terminal> python dc_v1.py
```

The text `Terminal>` just signifies a prompt in a Unix/Linux or DOS terminal window. After this prompt (which will look different in your terminal window, depending on the terminal application and how it is set up), commands like `python dc_v1.py` can be issued. These commands are interpreted by the operating system.

We strongly recommend to run Python programs within the IPython shell. First start IPython by typing `ipython` in the terminal window. Inside the IPython shell, our program `dc_v1.py` is run by the command `run dc_v1.py`. The advantage of running programs in IPython are many: previous commands are easily recalled, `%pdb` turns on debugging so that variables can be examined if the program aborts due to an exception, output of commands are stored in variables, programs and statements can be profiled, any operating system command can be executed, modules can be loaded automatically and other customizations can be performed when starting IPython – to mention a few of the most useful features.

Although running programs in IPython is strongly recommended, most execution examples in the forthcoming text simply use a minimal text like `Terminal> python programname`.

<div align="center">18</div>

2.2 Verifying the implementation

It is easy to make mistakes while deriving and implementing numerical algorithms, so we should never believe in the printed u values before they have been thoroughly verified. The most obvious idea is to compare the computed solution with the exact solution, when that exists, but there will always be a discrepancy between these two solutions because of the numerical approximations. The challenging question is whether we have the mathematically correct discrepancy or if we have another, maybe small, discrepancy due to both an approximation error and an error in the implementation.

The purpose of *verifying* a program is to bring evidence for the fact that there are no errors in the implementation. To avoid mixing unavoidable approximation errors and undesired implementation errors, we should try to make tests where we have some exact computation of the discrete solution or at least parts of it.

Running a few algorithmic steps by hand. The simplest approach to produce a correct reference for the discrete solution u of finite difference equations is to compute a few steps of the algorithm by hand. Then we can compare the hand calculations with numbers produced by the program.

A straightforward approach is to use a calculator and compute u^1, u^2, and u^3. However, the chosen values of I and θ given in the execution example above are not good, because the numbers 0 and 1 can easily simplify formulas too much for test purposes. For example, with $\theta = 1$ the nominator in the formula for u^n will be the same for all a and Δt values. One should therefore choose more "arbitrary" values, say $\theta = 0.8$ and $I = 0.1$. Hand calculations with the aid of a calculator gives

$$A \equiv \frac{1 - (1 - \theta)a\Delta t}{1 + \theta a \Delta t} = 0.298245614035$$

$$u^1 = AI = 0.0298245614035,$$
$$u^2 = Au^1 = 0.00889504462912,$$
$$u^3 = Au^2 = 0.00265290804728$$

Comparison of these manual calculations with the result of the `solver` function is carried out in the function

```
def verify_three_steps():
    """Compare three steps with known manual computations."""
    theta = 0.8; a = 2; I = 0.1; dt = 0.8
    u_by_hand = array([I,
                       0.0298245614035,
                       0.00889504462912,
                       0.00265290804728])

    N = 3  # number of time steps
    u, t = solver(I=I, a=a, T=N*dt, dt=dt, theta=theta)
```

```
    tol = 1E-15  # tolerance for comparing floats
    difference = abs(u - u_by_hand).max()
    success = difference <= tol
    return success
```

The main program, where we call the `solver` function and print `u`, is now put in a separate function `main`:

```
def main():
    u, t = solver(I=1, a=2, T=8, dt=0.8, theta=1)
    # Write out a table of t and u values:
    for i in range(len(t)):
        print 't=%6.3f u=%g' % (t[i], u[i])
        # or print 't={t:6.3f} u={u:g}'.format(t=t[i], u=u[i])
```

The main program in the file may now first run the verification test and then go on with the real simulation (`main()`) only if the test is passed:

```
if verify_three_steps():
    main()
else:
    print 'Bug in the implementation!'
```

Since the verification test is always done, future errors introduced accidentally in the program have a good chance of being detected.

It is essential that verification tests can be automatically run at *any* time. For this purpose, there are test frameworks and corresponding programming rules that allow us to request running through a suite of test cases, but in this very early stage of program development we just implement and run the verification in our own code so that every detail is visible and understood.

The complete program including the `verify_three_steps*` functions is found in the file `dc_verf1.py`.

Comparison with an exact discrete solution. Sometimes it is possible to find a closed-form *exact discrete solution* that fulfills the discrete finite difference equations. The implementation can then be verified against the exact discrete solution. This is usually the best technique for verification.

Define

$$A = \frac{1 - (1-\theta)a\Delta t}{1 + \theta a\Delta t}\,.$$

Manual computations with the θ-rule results in

$$u^0 = I,$$
$$u^1 = Au^0 = AI,$$
$$u^2 = Au^1 = A^2 I,$$
$$\vdots$$
$$u^n = A^n u^{n-1} = A^n I\,.$$

We have then established the exact discrete solution as

$$u^n = IA^n.$$ (34)

One should be conscious about the different meanings of the notation on the left- and right-hand side of this equation: on the left, n is a superscript reflecting a counter of mesh points, while on the right, n is the power in an exponentiation.

Comparison of the exact discrete solution and the computed solution is done in the following function:

```
def verify_exact_discrete_solution():

    def exact_discrete_solution(n, I, a, theta, dt):
        factor = (1 - (1-theta)*a*dt)/(1 + theta*dt*a)
        return I*factor**n

    theta = 0.8; a = 2; I = 0.1; dt = 0.8
    N = int(8/dt)   # no of steps
    u, t = solver(I=I, a=a, T=N*dt, dt=dt, theta=theta)
    u_de = array([exact_discrete_solution(n, I, a, theta, dt)
                  for n in range(N+1)])
    difference = abs(u_de - u).max()   # max deviation
    tol = 1E-15   # tolerance for comparing floats
    success = difference <= tol
    return success
```

Note that one can define a function inside another function (but such a function is invisible outside the function in which it is defined). The complete program is found in the file `dc_verf2.py`.

2.3   Computing the numerical error

Now that we have evidence for a correct implementation, we are in a position to compare the computed u^n values in the `u` array with the exact u values at the mesh points, in order to study the error in the numerical solution.

Let us first make a function for the analytical solution $u_e(t) = Ie^{-at}$ of the model problem:

```
def exact_solution(t, I, a):
    return I*exp(-a*t)
```

A natural way to compare the exact and discrete solutions is to calculate their difference at the mesh points:

$$e_n = u_e(t_n) - u^n, \quad n = 0, 1, \ldots, N.$$ (35)

These numbers are conveniently computed by

```
u, t = solver(I, a, T, dt, theta)   # Numerical solution
u_e = exact_solution(t, I, a)
e = u_e - u
```

The last two statements make use of array arithmetics: `t` is an array of mesh points that we pass to `exact_solution`. This function evaluates `-a*t`, which is a scalar times an array, meaning that the scalar is multiplied with each array element. The result is an array, let us call it `tmp1`. Then `exp(tmp1)` means applying the exponential function to each element in `tmp`, resulting an array, say `tmp2`. Finally, `I*tmp2` is computed (scalar times array) and `u_e` refers to this array returned from `exact_solution`. The expression `u_e - u` is the difference between two arrays, resulting in a new array referred to by `e`.

The array `e` is the current problem's discrete *error function*. Very often we want to work with just one number reflecting the size of the error. A common choice is to integrate e_n^2 over the mesh and take the square root. Assuming the exact and discrete solution to vary linearly between the mesh points, the integral is given exactly by the Trapezoidal rule:

$$\hat{E}^2 = \Delta t \left(\frac{1}{2} e_0^2 + \frac{1}{2} e_N^2 + \sum_{n=1}^{N-1} e_n^2 \right)$$

A common approximation of this expression, for convenience, is

$$\hat{E}^2 \approx E^2 = \Delta t \sum_{n=0}^{N} e_n^2$$

The error in this approximation is not much of a concern: it means that the error measure is not exactly the Trapezoidal rule of an integral, but a slightly different measure. We could equally well have chosen other error messages, but the choice is not important as long as we use the same error measure consistently in all experiments when investigating the error.

The error measure $\hat{E}$ or E is referred to as the L_2 norm of the discrete error function. The formula for E will be frequently used:

$$E = \sqrt{\Delta t \sum_{n=0}^{N} e_n^2} \tag{36}$$

The corresponding Python code, using array arithmetics, reads

```
E = sqrt(dt*sum(e**2))
```

The `sum` function comes from `numpy` and computes the sum of the elements of an array. Also the `sqrt` function is from `numpy` and computes the square root of each element in the array argument.

Instead of doing array computing we can compute with one element at a time:

```
m = len(u)      # length of u array (alt: u.size)
u_e = zeros(m)
t = 0
```

```
for i in range(m):
    u_e[i] = exact_solution(t, a, I)
    t = t + dt
e = zeros(m)
for i in range(m):
    e[i] = u_e[i] - u[i]
s = 0  # summation variable
for i in range(m):
    s = s + e[i]**2
error = sqrt(dt*s)
```

Such element-wise computing, often called *scalar* computing, takes more code, is less readable, and runs much slower than array computing.

2.4 Plotting solutions

Having the `t` and `u` arrays, the approximate solution `u` is visualized by `plot(t, u)`:

```
from matplotlib.pyplot import *
plot(t, u)
show()
```

It will be illustrative to also plot $u_e(t)$ for comparison. Doing a `plot(t, u_e)` is not exactly what we want: the `plot` function draws straight lines between the discrete points `(t[n], u_e[n])` while $u_e(t)$ varies as an exponential function between the mesh points. The technique for showing the "exact" variation of $u_e(t)$ between the mesh points is to introduce a very fine mesh for $u_e(t)$:

```
t_e = linspace(0, T, 1001)      # fine mesh
u_e = exact_solution(t_e, I, a)
plot(t,   u,   'r-')            # red  line for u
plot(t_e, u_e, 'b-')           # blue line for u_e
```

With more than one curve in the plot we need to associate each curve with a legend. We also want appropriate names on the axis, a title, and a file containing the plot as an image for inclusion in reports. The Matplotlib package (`matplotlib.pyplot`) contains functions for this purpose. The names of the functions are similar to the plotting functions known from MATLAB. A complete plot session then becomes

```
from matplotlib.pyplot import *

figure()                        # create new plot
t_e = linspace(0, T, 1001)      # fine mesh for u_e
u_e = exact_solution(t_e, I, a)
plot(t,   u,   'r--o')          # red dashes w/circles
plot(t_e, u_e, 'b-')           # blue line for exact sol.
legend(['numerical', 'exact'])
xlabel('t')
ylabel('u')
title('theta=%g, dt=%g' % (theta, dt))
savefig('%s_%g.png' % (theta, dt))
show()
```

23

Note that `savefig` here creates a PNG file whose name reflects the values of θ and Δt so that we can easily distinguish files from different runs with θ and Δt.

A bit more sophisticated and easy-to-read filename can be generated by mapping the θ value to acronyms for the three common schemes: FE (Forward Euler, $\theta = 0$), BE (Backward Euler, $\theta = 1$), CN (Crank-Nicolson, $\theta = 0.5$). A Python dictionary is ideal for such a mapping from numbers to strings:

```
theta2name = {0: 'FE', 1: 'BE', 0.5: 'CN'}
savefig('%s_%g.png' % (theta2name[theta], dt))
```

Let us wrap up the computation of the error measure and all the plotting statements in a function `explore`. This function can be called for various θ and Δt values to see how the error varies with the method and the mesh resolution:

```
def explore(I, a, T, dt, theta=0.5, makeplot=True):
    """
    Run a case with the solver, compute error measure,
    and plot the numerical and exact solutions (if makeplot=True).
    """
    u, t = solver(I, a, T, dt, theta)  # Numerical solution
    u_e = exact_solution(t, I, a)
    e = u_e - u
    E = sqrt(dt*sum(e**2))
    if makeplot:
        figure()                        # create new plot
        t_e = linspace(0, T, 1001)      # fine mesh for u_e
        u_e = exact_solution(t_e, I, a)
        plot(t,   u,   'r--o')          # red dashes w/circles
        plot(t_e, u_e, 'b-')            # blue line for exact sol.
        legend(['numerical', 'exact'])
        xlabel('t')
        ylabel('u')
        title('theta=%g, dt=%g' % (theta, dt))
        theta2name = {0: 'FE', 1: 'BE', 0.5: 'CN'}
        savefig('%s_%g.png' % (theta2name[theta], dt))
        show()
    return E
```

The `figure()` call is key here: without it, a new `plot` command will draw the new pair of curves in the same plot window, while we want the different pairs to appear in separate windows and files. Calling `figure()` ensures this.

The complete code resides in the file `dc_plot_mpl.py`. Running this program results in

```
Terminal> python dc_plot_mpl.py
0.0   0.40:    2.105E-01
0.0   0.04:    1.449E-02
0.5   0.40:    3.362E-02
0.5   0.04:    1.887E-04
1.0   0.40:    1.030E-01
1.0   0.04:    1.382E-02
```

We observe that reducing Δt by a factor of 10 increases the accuracy for all three methods (θ values). We also see that the combination of $\theta = 0.5$ and a small time step $\Delta t = 0.04$ gives a much more accurate solution, and that $\theta = 0$ and $\theta = 0$ with $\Delta t = 0.4$ result in the least accurate solutions.

Figure 6 demonstrates that the numerical solution for $\Delta t = 0.4$ clearly lies below the exact curve, but that the accuracy improves considerably by using 1/10 of this time step.
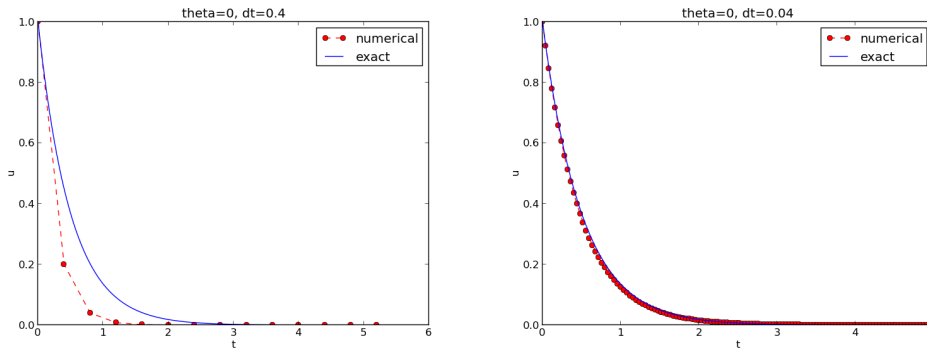
Figure 6: The Forward Euler scheme for two values of the time step.

Mounting two PNG files, as done in the figure, is easily done by the `montage` program from the ImageMagick suite:

```
                                    Terminal
Terminal> montage -background white -geometry 100% -tile 2x1 \
          FE_0.4.png FE_0.04.png FE1.png
```

The `-geometry` argument is used to specify the size of the image, and here we preserve the individual sizes of the images. The `-tile HxV` option specifies `H` images in the horizontal direction and `V` images in the vertical direction. A series of image files to be combined are then listed, with the name of the resulting combined image, here `FE1.png` at the end.

The behavior of the two other schemes are shown in Figures 7 and 8. Crank-Nicolson is obviously the most accurate scheme from a visual point of view.

2.5   Plotting with SciTools

The SciTools package provides a unified plotting interface, called Easyviz, to many different plotting packages, including Matplotlib. The syntax is very similar to that of Matplotlib and MATLAB. In fact, the plotting commands shown above look the same in SciTool's Easyviz interface, apart from the import statement, which reads

Figure 7: The Backward Euler scheme for two values of the time step.
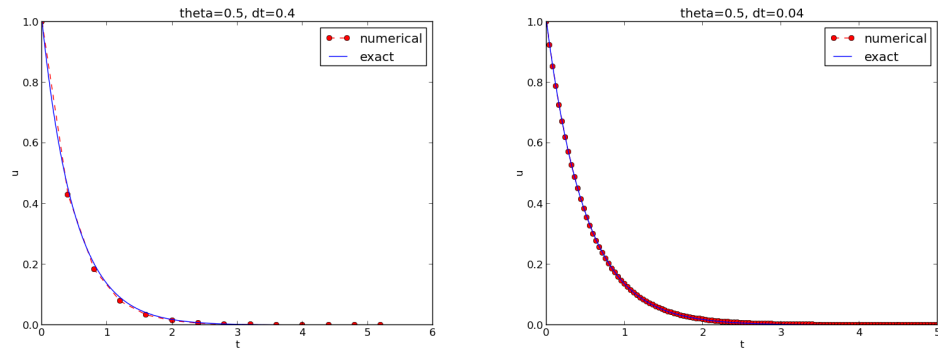
Figure 8: The Crank-Nicolson scheme for two values of the time step.

```
from scitools.std import *
```

This statement performs a `from numpy import *` as well as an import of the most common pieces of the Easyviz (`scitools.easyviz`) package, along with some additional numerical functionality.

With Easyviz one can, using an extended `plot` command, merge several plotting commands into one, using keyword arguments:

```
plot(t,   u,   'r--o',              # red dashes w/circles
     t_e, u_e, 'b-',                # blue line for exact sol.
     legend=['numerical', 'exact'],
     xlabel='t',
     ylabel='u',
     title='theta=%g, dt=%g' % (theta, dt),
     savefig='%s_%g.png' % (theta2name[theta], dt),
     show=True)
```

The `dc_plot_st.py` file contains such a demo.

By default, Easyviz employs Matplotlib for plotting, but Gnuplot and Grace are viable alternatives:

```
                                  ┌──────────┐
──────────────────────────────────│ Terminal │──────────────────────────────
                                  └──────────┘
Terminal> python dc_plot_st.py --SCITOOLS_easyviz_backend gnuplot
Terminal> python dc_plot_st.py --SCITOOLS_easyviz_backend grace
```

The backend used for creating plots (and numerous other options) can be permanently set in SciTool's configuration file.

All the Gnuplot windows are launched without any need to kill one before the next one pops up (as is the case with Matplotlib) and one can press the key 'q' anywhere in a plot window to kill it. Another advantage of Gnuplot is the automatic choice of sensible and distinguishable line types in black-and-white PostScript files (produced by `savefig('myplot.eps')`).

Regarding functionality for annotating plots with title, labels on the axis, legends, etc., we refer to the documentation of Matplotlib and SciTools for more detailed information on the syntax. The hope is that the programming syntax explained so far suffices for understanding the code and learning more from a combination of the forthcoming examples and other resources such as books and web pages.

2.6   Creating user interfaces

It is good programming practice to let programs read input from the user rather than require the user to edit the source code when trying out new values of input parameters. Reading input from the command line is a simple and flexible way of interacting with the user. Python stores all the command-line arguments in the list `sys.argv`, and there are, in principle, two ways of programming with command-line arguments in Python:

- Decide upon a sequence of parameters on the command line and read their values directly from the `sys.argv[1:]` list (`sys.argv[0]` is the just program name).

- Use option-value pairs (`--option value`) on the command line to override default values of input parameters, and use the `argparse.ArgumentParser` tool to interact with the command line.

Both strategies will be illustrated next.

Reading a sequence of command-line arguments.   The `dc_plot_mpl.py` program needs the following input data: I, a, T, an option to turn the plot on or off (`makeplot`), and a list of Δt values.

The simplest way of reading this input from the command line is to say that the first four command-line arguments correspond to the first four points in the list above, in that order, and that the rest of the command-line arguments are

the Δt values. The input given for `makeplot` can be a string among `'on'`, `'off'`, `'True'`, and `'False'`. The code for reading this input is most conveniently put in a function:

```
import sys

def read_command_line():
    if len(sys.argv) < 6:
        print 'Usage: %s I a T on/off dt1 dt2 dt3 ...' % \
              sys.argv[0]; sys.exit(1)  # abort

    I = float(sys.argv[1])
    a = float(sys.argv[2])
    T = float(sys.argv[3])
    makeplot = sys.argv[4] in ('on', 'True')
    dt_values = [float(arg) for arg in sys.argv[5:]]

    return I, a, T, makeplot, dt_values
```

One should note the following about the constructions in the program above:

- Everything on the command line ends up in a *string* in the list `sys.argv`. Explicit conversion to, e.g., a `float` object is required if the string as a number we want to compute with.

- The value of `makeplot` is determined from a boolean expression, which becomes `True` if the command-line argument is either `'on'` or `'True'`, and `False` otherwise.

- It is easy to build the list of Δt values: we simply run through the rest of the list, `sys.argv[5:]`, convert each command-line argument to `float`, and collect these `float` objects in a list, using the compact and convenient *list comprehension* syntax in Python.

The loops over θ and Δt values can be coded in a `main` function:

```
def main():
    I, a, T, makeplot, dt_values = read_command_line()
    for theta in 0, 0.5, 1:
        for dt in dt_values:
            E = explore(I, a, T, dt, theta, makeplot)
            print '%3.1f %6.2f: %12.3E' % (theta, dt, E)
```

The complete program can be found in `dc_cml.py`.

Working with an argument parser. Python's `ArgumentParser` tool in the `argparse` module makes it easy to create a professional command-line interface to any program. The documentation of 'ArgumentParser' demonstrates its versatile applications, so we shall here just list an example containing the most used features. On the command line we want to specify option value pairs for I, a, and T, e.g., `--a 3.5 --I 2 --T 2`. Including `--makeplot` turns the plot on and excluding this option turns the plot off. The Δt values can be given

as `--dt 1 0.5 0.25 0.1 0.01`. Each parameter must have a sensible default value so that we specify the option on the command line only when the default value is not suitable.

We introduce a function for defining the mentioned command-line options:

```python
def define_command_line_options():
    import argparse
    parser = argparse.ArgumentParser()
    parser.add_argument('--I', '--initial_condition', type=float,
                        default=1.0, help='initial condition, u(0)',
                        metavar='I')
    parser.add_argument('--a', type=float,
                        default=1.0, help='coefficient in ODE',
                        metavar='a')
    parser.add_argument('--T', '--stop_time', type=float,
                        default=1.0, help='end time of simulation',
                        metavar='T')
    parser.add_argument('--makeplot', action='store_true',
                        help='display plot or not')
    parser.add_argument('--dt', '--time_step_values', type=float,
                        default=[1.0], help='time step values',
                        metavar='dt', nargs='+', dest='dt_values')
    return parser
```

Each command-line option is defined through the `parser.add_argument` method. Alternative options, like the short `--I` and the more explaining `--initial_condition` can be defined. Other arguments are `type` for the Python object type, a default value, and a help string, which gets printed if the command-line argument `-h` or `--help` is included. The `metavar` argument specifies the value associated with the option when the help string is printed. For example, the option for I has this help output:

-- Terminal --
```
Terminal> python dc_argparse.py -h
  ...
  --I I, --initial_condition I
                        initial condition, u(0)
  ...
```

The structure of this output is

```
    --I metavar, --initial_condition metavar
                      help-string
```

The `--makeplot` option is a pure flag without any value, implying a true value if the flag is present and otherwise a false value. The `action='store_true'` makes an option for such a flag.

Finally, the `--dt` option demonstrates how to allow for more than one value (separated by blanks) through the `nargs='+'` keyword argument. After the command line is parsed, we get an object where the values of the options are stored as attributes. The attribute name is specified by the `dist` keyword argument, which for the `--dt` option reads `dt_values`. Without the `dest` argument, the value of option `--opt` is stored as the attribute `opt`.

The code below demonstrates how to read the command line and extract the values for each option:

```
def read_command_line():
    parser = define_command_line_options()
    args = parser.parse_args()
    print 'I={}, a={}, T={}, makeplot={}, dt_values={}'.format(
        args.I, args.a, args.T, args.makeplot, args.dt_values)
    return args.I, args.a, args.T, args.makeplot, args.dt_values
```

The `main` function remains the same as in the `dc_cml.py` code based on reading from `sys.argv` directly. A complete program using the demo above of `ArgumentParser` appears in the file `dc_argparse.py`.

2.7 Computing convergence rates

We normally expect that the error E in the numerical solution is reduced if the mesh size Δt is decreased. More specifically, many numerical methods obey a power-law relation between E and Δt:

$$E = C\Delta t^r, \tag{37}$$

where C and r are (usually unknown) constants independent of Δt. The formula (37) is viewed as an asymptotic model valid for sufficiently small Δt. How small is normally hard to estimate without doing numerical estimations of r.

The parameter r is known as the *convergence rate*. For example, if the convergence rate is 2, halving Δt reduces the error by a factor of 4. Diminishing Δt then has a greater impact on the error compared with methods that have $r = 1$. For a given value of r, we refer to the method as of r-th order. First- and second-order methods are most common in scientific computing.

Estimating r. There are two ways of estimating C and r based on a set of m simulations with corresponding pairs $(\Delta t_i, E_i)$, $i = 0, \ldots, m - 1$, and $\Delta t_i < \Delta t_{i-1}$ (i.e., decreasing cell size).

1. Take the logarithm of (37), $\ln E = r \ln \Delta t + \ln C$, and fit a straight line to the data points $(\Delta t_i, E_i)$, $i = 0, \ldots, m - 1$.

2. Consider two consecutive experiments, $(\Delta t_i, E_i)$ and $(\Delta t_{i-1}, E_{i-1})$. Dividing the equation $E_{i-1} = C\Delta t_{i-1}^r$ by $E_i = C\Delta t_i^r$ and solving for r yields

$$r_{i-1} = \frac{\ln(E_{i-1}/E_i)}{\ln(\Delta t_{i-1}/\Delta t_i)} \tag{38}$$

for $i = 1, = \ldots, m - 1$.

The disadvantage of method 1 is that (37) might not be valid for the coarsest meshes (largest Δt values), and fitting a line to all the data points is then misleading. Method 2 computes convergence rates for pairs of experiments and

allows us to see if the sequence r_i converges to some value as $i \to m - 2$. The final r_{m-2} can then be taken as the convergence rate. If the coarsest meshes have a differing rate, the corresponding time steps are probably too large for (37) to be valid. That is, those time steps lie outside the asymptotic range of Δt values where the error behave like (37).

Implementation. It is straightforward to extend the `main` function in the program `dc_argparse.py` with statements for computing $r_0, r_1, \ldots, r_{m-2}$ from (37):

```
from math import log

def main():
    I, a, T, makeplot, dt_values = read_command_line()
    r = {}  # estimated convergence rates
    for theta in 0, 0.5, 1:
        E_values = []
        for dt in dt_values:
            E = explore(I, a, T, dt, theta, makeplot=False)
            E_values.append(E)

        # Compute convergence rates
        m = len(dt_values)
        r[theta] = [log(E_values[i-1]/E_values[i])/
                        log(dt_values[i-1]/dt_values[i])
                        for i in range(1, m, 1)]

    for theta in r:
        print '\nPairwise convergence rates for theta=%g:' % theta
        print ' '.join(['%.2f' % r_ for r_ in r[theta]])
    return r
```

The program is called `dc_convrate.py`.

The `r` object is a *dictionary of lists*. The keys in this dictionary are the θ values. For example, `r[1]` holds the list of the r_i values corresponding to $\theta = 1$. In the loop `for theta in r`, the loop variable `theta` takes on the values of the keys in the dictionary `r` (in an undetermined ordering). We could simply do a `print r[theta]` inside the loop, but this would typically yield output of the convergence rates with 16 decimals:

```
[1.331919482274763, 1.1488178494691532, ...]
```

Instead, we format each number with 2 decimals, using a list comprehension to turn the list of numbers, `r[theta]`, into a list of formatted strings. Then we join these strings with a space in between to get a sequence of rates on one line in the terminal window. More generally, `d.join(list)` joins the strings in the list `list` to one string, with `d` as delimiter between `list[0]`, `list[1]`, etc.

Here is an example on the outcome of the convergence rate computations:

—————————————————— Terminal ——————————————————
```
Terminal> python dc_convrate.py --dt 0.5 0.25 0.1 0.05 0.025 0.01
...
```

31

```
Pairwise convergence rates for theta=0:
1.33 1.15 1.07 1.03 1.02

Pairwise convergence rates for theta=0.5:
2.14 2.07 2.03 2.01 2.01

Pairwise convergence rates for theta=1:
0.98 0.99 0.99 1.00 1.00
```

The Forward and Backward Euler methods seem to have an r value which stabilizes at 1, while the Crank-Nicolson seems to be a second-order method with $r = 2$.

Very often, we have some theory that predicts what r is for a numerical method. Various theoretical error measures for the θ-rule point to $r = 2$ for $\theta = 0.5$ and $r = 1$ otherwise. The computed estimates of r are in very good agreement with these theoretical values.

The strong practical application of computing convergence rates is for verification: wrong convergence rates point to errors in the code, and correct convergence rates brings evidence that the implementation is correct. Experience shows that bugs in the code easily destroys the expected convergence rate.

Debugging via convergence rates. Let us experiment with bugs and see the implication on the convergence rate. We may, for instance, forget to multiply by `a` in the denominator in the updating formula for `u[n+1]`:

```
u[n+1] = (1 - (1-theta)*a*dt)/(1 + theta*dt)*u[n]
```

Running the same `dc_convrate.py` command as above gives the expected convergence rates (!). Why? The reason is that we just specified the Δt values are relied on default values for other parameters. The default value of a is 1. Forgetting the factor `a` has then no effect. This example shows how importance it is to avoid parameters that are 1 or 0 when verifying implementations. Running the code `dc_v0.py` with $a = 2.1$ and $I = 0.1$ yields

```
                                    Terminal
Terminal> python dc_convrate.py --a 2.1 --I 0.1  \
          --dt 0.5 0.25 0.1 0.05 0.025 0.01
...
Pairwise convergence rates for theta=0:
1.49 1.18 1.07 1.04 1.02

Pairwise convergence rates for theta=0.5:
-1.42 -0.22 -0.07 -0.03 -0.01

Pairwise convergence rates for theta=1:
0.21 0.12 0.06 0.03 0.01
```

This time we see that the expected convergence rates for the Crank-Nicolson and Backward Euler methods are not obtained, while $r = 1$ for the Forward

Euler method. The reason for correct rate in the latter case is that $\theta = 0$ and the wrong `theta*dt` term in the denominator vanishes anyway.

The error

```
u[n+1] = ((1-theta)*a*dt)/(1 + theta*dt*a)*u[n]
```

manifests itself through wrong rates $r \approx 0$ for all three methods. About the same results arise from an erroneous initial condition, `u[0] = 1`, or wrong loop limits, `range(1,N)`. It seems that in this simple problem, most bugs we can think of are detected by the convergence rate test, provided the values of the input data do not hide the bug.

A `verify_convergence_rate` function could compute the dictionary of list via `main` and check if the final rate estimates (r_{m-2}) are sufficiently close to the expected ones. A tolerance of 0.1 seems appropriate, given the uncertainty in estimating r:

```
def verify_convergence_rate():
    r = main()
    tol = 0.1
    expected_rates = {0: 1, 1: 1, 0.5: 2}
    for theta in r:
        r_final = r[theta][-1]
        diff = abs(expected_rates[theta] - r_final)
        if diff > tol:
            return False
    return True  # all tests passed
```

We remark that `r[theta]` is a list and the last element in any list can be extracted by the index `-1`.

2.8   Memory-saving implementation

The memory storage requirements of our implementations so far consists mainly of the `u` and `t` arrays, both of length $N + 1$, plus some other temporary arrays that Python needs for intermediate results if we do array arithmetics in our program (e.g., `I*exp(-a*t)` needs to store `a*t` before `-` can be applied to it and then `exp`). The extremely modest storage requirements of simple ODE problems put no restrictions on the formulations of the algorithm and implementation. Nevertheless, when the methods for ODEs used here are applied to three-dimensional partial differential equation (PDE) problems, memory storage requirements suddenly become an issue.

The PDE counterpart to our model problem $u' = -a$ is a diffusion equation $u_t = a\nabla^2 u$ posed on a space-time domain. The discrete representation of this domain may in 3D be a spatial mesh of M^3 points and a time mesh of N points. A typical desired value for M is 100 in many applications, or even 1000. Storing all the computed u values, like we have done in the programs so far, demands storage of some arrays of size $M^3 N$, giving a factor of M^3 larger storage demands compared to our ODE programs. Each real number in the array for u requires 8 bytes of storage, resulting in a demand for 8 Gb of memory for only one array.

In such cases one needs good ideas on how to lower the storage requirements. Fortunately, we can usually get rid of the M^3 factor. Below we explain how this is done, and the technique is almost always applied in implementations of PDE problems.

Let us critically evaluate how much we really need to store in the computer's memory in our implementation of the θ method. To compute a new u^{n+1}, all we need is u^n. This implies that the previous $u^{n-1}, u^{n-2}, \ldots, u^0$ values do not need to be stored in an array, although this is convenient for plotting and data analysis in the program. Instead of the u array we can work with two variables for real numbers, u and u_1, representing u^{n+1} and u^n in the algorithm, respectively. At each time level, we update u from u_1 and then set u_1 = u so that the computed u^{n+1} value becomes the "previous" value u^n at the next time level. The downside is that we cannot plot the solution after the simulation is done since only the last two numbers are available. The remedy is to store computed values in a file and use the file for visualizing the solution later.

We have implemented this memory saving idea in the file `dc_memsave.py`, which is a merge of the `dc_plot_mpl.py` and `dc_argparse.py` programs, using module prefixes np for numpy and plt for matplotlib.pyplot.

The following function implements the ideas above regarding minimizing memory usage and storing the solution on file:

```python
def solver_memsave(I, a, T, dt, theta, filename='sol.dat'):
    """
    Solve u'=-a*u, u(0)=I, for t in (0,T] with steps of dt.
    Minimum use of memory. The solution is store on file
    (with name filename) for later plotting.
    """
    dt = float(dt)          # avoid integer division
    N = int(round(T/dt))    # no of intervals

    outfile = open(filename, 'w')
    # u: time level n+1, u_1: time level n
    t = 0
    u_1 = I
    outfile.write('%.16E  %.16E\n' % (t, u_1))
    for n in range(1, N+1):
        u = (1 - (1-theta)*a*dt)/(1 + theta*dt*a)*u_1
        u_1 = u
        t += dt
        outfile.write('%.16E  %.16E\n' % (t, u))
    outfile.close()
    return u, t
```

This code snippet serves as a quick introduction to file writing in Python. Reading the data in the file into arrays t and u are done by the function

```python
def read_file(filename='sol.dat'):
    infile = open(filename, 'r')
    u = [];   t = []
    for line in infile:
        words = line.split()
        if len(words) != 2:
            print 'Found more than two numbers on a line!', words
```

```
            sys.exit(1)  # abort
        t.append(float(words[0]))
        u.append(float(words[1]))
    return np.array(t), np.array(u)
```

This type of file with numbers in rows and columns is very common, and `numpy` has a function `loadtxt` which loads such tabular data into a two-dimensional array, say `data`. The number in row `i` and column `j` is then `data[i,j]`. The whole column number `j` can be extracted by `data[:,j]`. A version of `read_file` using `np.loadtxt` reads

```
def read_file_numpy(filename='sol.dat'):
    data = np.loadtxt(filename)
    t = data[:,0]
    u = data[:,1]
    return t, u
```

The present counterpart to the `explore` function from `dc_plot_mpl.py` must run `solver_memsave` and then load data from file before we can compute the error measure and make the plot:

```
def explore(I, a, T, dt, theta=0.5, makeplot=True):
    filename = 'u.dat'
    u, t = solver_memsave(I, a, T, dt, theta, filename)

    t, u = read_file(filename)
    u_e = exact_solution(t, I, a)
    e = u_e - u
    E = np.sqrt(dt*np.sum(e**2))
    if makeplot:
        plt.figure()
        ...
```

The `dc_memsave.py` file also includes command-line options `--I`, `--a`, `--T`, `--dt`, `--theta`, and `--makeplot` for controlling input parameters and making a single run. For example,

---
Terminal

```
Terminal> python dc_memsave.py --T 10 --theta 1 --dt 2
I=1.0, a=1.0, T=10.0, makeplot=True, theta=1.0, dt=2.0
theta=1.0 dt=2 Error=3.136E-01
```
---

3 Software engineering

Efficient use of differential equation models requires software that is easy to test and flexible for setting up extensive numerical experiments. This section introduces three important concepts and their applications to the exponential decay model:

- Modules

- Testing frameworks

- Implementation with classes

3.1 Making a module

The previous sections has outlined numerous different programs, all of them having their own copy of the `solver` function. Such copies of the same piece of code is against the important *Don't Repeat Yourself* (DRY) principle in programming. If we want to change the `solver` function there should be one and only one place where the change needs to be performed.

To clean up the repetitive code snippets scattered among the `dc_*.py` files, we start by collecting the various functions we want to keep for the future in one file, now called `dc_mod.py`. The following functions are copied to this file:

- `solver`

- `verify_three_steps`

- `verify_discrete_solution`

- `explore`

- `define_command_line_options`

- `read_command_line` extended to work both with `sys.argv` directly and with an `ArgumentParser` object

- `main` (with convergence rates estimation as in `dc_convrate.py`)

- `verify_convergence_rate`

We use Matplotlib for plotting. A sketch of the `dc_mod.py` file, with complete versions of the modified functions, looks as follows:

```
from numpy import *
from matplotlib.pyplot import *
import sys

def solver(I, a, T, dt, theta):
    ...

def verify_three_steps():
    ...

def verify_exact_discrete_solution():
    ...

def exact_solution(t, I, a):
    ...

def explore(I, a, T, dt, theta=0.5, makeplot=True):
    ...
```

36

```
def define_command_line_options():
    ...

def read_command_line(use_argparse=True):
    if use_argparse:
        parser = define_command_line_options()
        args = parser.parse_args()
        print 'I={}, a={}, makeplot={}, dt_values={}'.format(
            args.I, args.a, args.makeplot, args.dt_values)
        return args.I, args.a, args.makeplot, args.dt_values
    else:
        if len(sys.argv) < 6:
            print 'Usage: %s I a on/off dt1 dt2 dt3 ...' % \
                    sys.argv[0]; sys.exit(1)

        I = float(sys.argv[1])
        a = float(sys.argv[2])
        T = float(sys.argv[3])
        makeplot = sys.argv[4] in ('on', 'True')
        dt_values = [float(arg) for arg in sys.argv[5:]]

        return I, a, makeplot, dt_values

def main():
    ...
```

This `dc_mod.py` file is already a module such that we can import desired in functions in other programs. For example,

```
from decay_theta import solver
u, t = solver(I=1.0, a=3.0, T=3, dt=0.01, theta=0.5)
```

However, it should also be possible to both use `dc_mod.py` as a module *and* execute the file as a program that runs `main()`. This is accomplished by ending the file with a *test block*:

```
if __name__ == '__main__':
    main()
```

When `dc_mod.py` is used as a module, `__name__` equals the module name `decay_theta`, while `__name__` equals `'__main__'` when the file is run as a program. Optionally, we could run the verification tests if the word `verify` is present on the command line and `verify_convergence_rate` could be tested if `verify_rates` is found on the command line. The `verify_rates` argument must be removed before we read parameter values from the command line, other wise `read_command_line` (called by `main`) will not work properly.

```
if __name__ == '__main__':
    if 'verify' in sys.argv:
        if verify_three_steps() and verify_discrete_solution():
            pass # ok
        else:
            print 'Bug in the implementation!'
    elif 'verify_rates' in sys.argv:
        sys.argv.remove('verify_rates')
```

```
        if not '--dt' in sys.argv:
            print 'Must assign several dt values'
            sys.exit(1)   # abort
        if verify_convergence_rate():
            pass
        else:
            print 'Bug in the implementation!'
    else:
        # Perform simulations
        main()
```

3.2   Prefixing imported functions by the module name

Import statements of the form `from module import *` imports functions and
variables in `module.py` into the current file. For example, when doing

```
from numpy import *
from matplotlib.pyplot import *
```

we get mathematical functions like `sin` and `exp` as well as MATLAB-style func-
tions like `linspace` and `plot`, which can be called by these well-known names.
However, it sometimes becomes confusing to know where a particular function
comes from. Is it from `numpy`? Or `matplotlib.pyplot`? Or is it our own
function?

An alternative import is

```
import numpy
import matplotlib.pyplot
```

and such imports require functions to be prefixed by the module name, e.g.,

```
t = numpy.linspace(0, T, N+1)
u_e = I*numpy.exp(-a*t)
matplotlib.pyplot.plot(t, u_e)
```

This is normally regarded as a better habit because it is explicitly stated from
which module a function comes from.

The modules `numpy` and `matplotlib.pyplot` are so frequently used, and
their full names quite tedious to write, so two standard abbreviations have
evolved in the Python scientific computing community:

```
import numpy as np
import matplotlib.pyplot as plt

t = np.linspace(0, T, N+1)
u_e = I*np.exp(-a*t)
plt.plot(t, u_e)
```

A version of the `decay_theta` module where we use the `np` and `plt` prefixes is
found in the file `decay_theta_v2.py`.

The downside of prefixing functions by the module name is that mathemat-
ical expressions like $e^{-at}\sin(2\pi t)$ gets cluttered with module names,

```
numpy.exp(-a*t)*numpy.sin(2(numpy.pi*t)
# or
np.exp(-a*t)*np.sin(2*np.pi*t)
```

Such an expression looks like `exp(-a*t)*sin(2*pi*t)` in most other programming languages. Similarly, `np.linspace` and `plt.plot` look less familiar to people who are used to MATLAB and who have not adopted Python's prefix style. Whether to do `from module import *` or `import module` depends on personal taste and the problem at hand. In these writings we use `from module import` in shorter programs where similarity with MATLAB could be an advantage, and where a one-to-one correspondence between mathematical formulas and Python expressions is important. The style `import module` is preferred inside Python modules (see Exercise 9 for a demonstration).

3.3 Doctests

We have emphasized how important it is to be able to run tests in the program at any time. This was solved by calling various `verify*` functions in the previous examples. However, there exists well-established procedures and corresponding tools for automating checking of tests. We shall briefly demonstrate two important techniques: *doctest* and *unit testing*. The corresponding files are the modules `dc_mod_doctest.py` and `dc_mod_unittest.py`.

Doc strings (the first string after the function header) are used to document the purpose of functions and their arguments. Very often it is instructive to include an example on how to use the function. Interactive examples in the Python shell are most illustrative as we can see the output resulting from function calls. For example, we can in the `solver` function include an example on calling this function and printing the computed `u` and `t` arrays:

```
def solver(I, a, T, dt, theta):
    """
    Solve u'=-a*u, u(0)=I, for t in (0,T] with steps of dt.

    >>> u, t = solver(I=0.8, a=1.2, T=4, dt=0.5, theta=0.5)
    >>> for t_n, u_n in zip(t, u):
    ...     print 't=%.1f, u=%.14f' % (t_n, u_n)
    t=0.0, u=0.80000000000000
    t=0.5, u=0.43076923076923
    t=1.0, u=0.23195266272189
    t=1.5, u=0.12489758761948
    t=2.0, u=0.06725254717972
    t=2.5, u=0.03621291001985
    t=3.0, u=0.01949925924146
    t=3.5, u=0.01049960113002
    t=4.0, u=0.00565363137770
    """
    ...
```

When such interactive demonstrations are inserted in doc strings, Python's `doctest` module can be used to automate running all commands in interactive

sessions and compare new output with the output appearing in the doc string. All we have to do in the current example is to write

```
Terminal> python -m doctest dc_mod_doctest.py
```

This command imports the `doctest` module, which runs all tests. No additional command-line argument is allowed when running doctest. If any test fails, the problem is reported, e.g.,

Terminal

```
Terminal> python -m doctest dc_mod_doctest.py
**********************************************************
File "dc_mod_doctest.py", line 12, in dc_mod_doctest....
Failed example:
    for t_n, u_n in zip(t, u):
        print 't=%.1f, u=%.14f' % (t_n, u_n)
Expected:
    t=0.0, u=0.80000000000000
    t=0.5, u=0.43076923076923
    t=1.0, u=0.23195266272189
    t=1.5, u=0.12489758761948
    t=2.0, u=0.06725254717972
    t=2.5, u=0.03621291001985
    t=3.0, u=0.01949925924146
    t=3.5, u=0.01049960113002
    t=4.0, u=0.00565363137770
Got:
    t=0.0, u=0.80000000000000
    t=0.5, u=0.43076923076923
    t=1.0, u=0.23195266272189
    t=1.5, u=0.12489758761948
    t=2.0, u=0.06725254717972
**********************************************************
1 items had failures:
   1 of   2 in dc_mod_doctest.solver
***Test Failed*** 1 failures.
```

Note that in the output of `t` and `u` we write `u` with 14 digits. Writing all 16 digits is not a good idea: if the tests are run on different hardware, round-off errors might be different, and the `doctest` module detects numbers are are not presicely the same and reports failures. In the present application, where $0 < u(t) \leq 0.8$, we expect round-off errors to be of size 10^{-16}, so comparing 15 digits would probably be reliable, but we compare 14 to be on the safe side.

Doctests are highly encouraged as they do two things: 1) demonstrate how a function is used and 2) test that the function works.

Here is an example on a doctest in the `explore` function:

```
def explore(I, a, T, dt, theta=0.5, makeplot=True):
    """
    Run a case with the solver, compute error measure,
    and plot the numerical and exact solutions (if makeplot=True).

    >>> for theta in 0, 0.5, 1:
```

40

```
...      E = explore(I=1.9, a=2.1, T=5, dt=0.1, theta=theta,
...                  makeplot=False)
...      print '%.10E' % E
...
7.3565079236E-02
2.4183893110E-03
6.5013039886E-02
"""
...
```

This time we limit the output to 10 digits.

We remark that doctests are not straightforward to construct for functions that rely on command-line input and that print results to the terminal window.

3.4   Unit testing with nose

The unit testing technique consists of identifying small units of code, usually functions (or classes), and write one or more tests for each unit. One test should, ideally, not depend on the outcome of other tests. For example, the doctest in function `solver` is a unit test, and the doctest in function `explore` as well, but the latter depends on a working `solver`. Putting the error computation and plotting in `explore` in two separate functions would allow independent unit tests. In this way, the design of unit tests impacts the design of functions. The recommended practice is actually to design and write the unit tests first and *then* implement the functions!

In scientific computing it is still not obvious how to best perform unit testing. The units is naturally larger than in non-scientific software. Very often the solution procedure of a mathematical problem identifies a unit.

Basic use of nose.   The `nose` package is a versatile tool for implementing unit tests in Python. Here is a recommended way of structuring tests:

1. Collect tests in separate files with names starting with `test_`.

2. Implement tests in functions with names starting with `test_`.

3. These functions perform assertions on computing results using `assert` functions from the `nose.tools` module.

Here comes a very simple illustration of the three points. Assume that we have this function in a module `mymod`:

```
def double(n):
    return 2*n
```

In a file `test_mymod.py` we implement a test function whose purpose is to test that the function `double` works as intended:

```
import mymod
import nose.tools as nt

def test_double():
    result = mymod.double(4)
    nt.assert_equal(result, 8)
```

Running

```
Terminal> nosetests
```

makes the `nose` tool look for `test_*.py` files and run the `test*()` functions found in these files. If the `nt.assert_equal` function finds that the two arguments are equal, the test is a success, otherwise it is a failure and an exception of type `AssertionError` is raised. Instead of calling the convenience function `nt.assert_equal`, we can use Python's `assert` statement, which tests if a boolean expression is true and raises an `AssertionError` if the expression is false. Here, the statement is `assert result == 8`. A completely manual alternative is to write

```
if result != 8:
    raise AssertionError()
```

The number of failed tests and their details are reported, or an `OK` is printed if all tests passed. Imports like `import mymod` in `test_*.py` files works even if `mymod.py` is not located in the same directory as the test file (`nose` will find it without any need for manipulating `PYTHONPATH` or `sys.path`).

The advantage with the `nose` is two-fold: 1) tests are written and collected in a structured way, and 2) large collections of tests, scattered throughout a tree of folders, can be executed with one command (`nosetests`).

Demonstrating nose. Let us illustrate how to use the `nose` tool for testing key functions in the `decay_theta` module. Or more precisely, the module is called `decay_theta_unittest` with all the `verify*` functions removed as these now are outdated by the unit tests.

We design three unit tests:

1. A comparison between the computed u^n values and the exact discrete solution.

2. A comparison between the computed u^n values and precomputed, verified reference values.

3. A comparison between observed and expected convergence rates.

These tests follow very closely the code in the previously shown `verify*` functions. We start with comparing u^n, as computed by the function `solver`, to the formula for the exact discrete solution:

```
import nose.tools as nt
import sys, os
sys.path.insert(0, os.pardir)
import dc_mod_unittest as dc_mod
import numpy as np

def exact_discrete_solution(n, I, a, theta, dt):
    """Return exact discrete solution of the theta scheme."""
    dt = float(dt)  # avoid integer division
    factor = (1 - (1-theta)*a*dt)/(1 + theta*dt*a)
    return I*factor**n

def test_against_discrete_solution():
    """
    Compare result from solver against
    formula for the discrete solution.
    """
    theta = 0.8; a = 2; I = 0.1; dt = 0.8
    N = int(8/dt)   # no of steps
    u, t = dc_mod.solver(I=I, a=a, T=N*dt, dt=dt, theta=theta)
    u_de = np.array([exact_discrete_solution(n, I, a, theta, dt)
                     for n in range(N+1)])
    diff = np.abs(u_de - u).max()
    nt.assert_almost_equal(diff, 0, delta=1E-14)
```

The `nt.assert_almost_equal` is the relevant function for comparing two
real numbers. The `delta` argument specifies a tolerance for the comparison.
Alternatively, one can specify a `places` argument for the number of decimal
places to enter the comparison.

When we at some point have carefully verified the implementation, we may
store correctly computed numbers in the test program or in files for use in future
tests. Here is an example on how the outcome from the `solver` function can be
compared to what is considered to be correct results:

```
def test_solver():
    """
    Compare result from solver against
    precomputed arrays for theta=0, 0.5, 1.
    """
    I=0.8; a=1.2; T=4; dt=0.5  # fixed parameters
    precomputed = {
        't': np.array([ 0. ,   0.5,   1. ,   1.5,   2. ,   2.5,
                        3. ,   3.5,   4. ]),
        0.5: np.array(
            [ 0.8       ,  0.43076923,  0.23195266,  0.12489759,
              0.06725255,  0.03621291,  0.01949926,  0.0104996 ,
              0.00565363]),
        0: np.array(
            [  8.00000000e-01,   3.20000000e-01,
               1.28000000e-01,   5.12000000e-02,
               2.04800000e-02,   8.19200000e-03,
               3.27680000e-03,   1.31072000e-03,
               5.24288000e-04]),
        1: np.array(
            [ 0.8       ,  0.5       ,  0.3125    ,  0.1953125 ,
              0.12207031,  0.07629395,  0.04768372,  0.02980232,
              0.01862645]),
```

```
        }
    for theta in 0, 0.5, 1:
        u, t = dc_mod.solver(I, a, T, dt, theta=theta)
        diff = np.abs(u - precomputed[theta]).max()
        # Compare to 8 decimal places
        nt.assert_almost_equal(diff, 0, places=8,
                               msg='theta=%s' % theta)
```

The `precomputed` object is a dictionary with four keys: `'t'` for the time mesh, and three θ values for u^n solutions corresponding to the $\theta = 0, 0.5, 1$.

Testing for special type of input data that may cause trouble constitutes a common way of constructing unit tests. The updating formula for u^{n+1} may be incorrectly evaluated because of unintended integer divisions. For example, with

```
theta = 1; a = 1; I = 1; dt = 2
```

the nominator and denominator in the updating expression,

```
(1 - (1-theta)*a*dt)
(1 + theta*dt*a)
```

evaluate to 1 and 3, respectively, and the fraction `1/3` will call up integer division and consequently lead to `u[n+1]=0`. We construct a unit test to make sure `solver` is smart enough to avoid this problem:

```
def test_potential_integer_division():
    """Choose variables that can trigger integer division."""
    theta = 1; a = 1; I = 1; dt = 2
    N = 4
    u, t = dc_mod.solver(I=I, a=a, T=N*dt, dt=dt, theta=theta)
    u_de = np.array([exact_discrete_solution(n, I, a, theta, dt)
                     for n in range(N+1)])
    diff = np.abs(u_de - u).max()
    nt.assert_almost_equal(diff, 0, delta=1E-14)
```

The final test is to see if $\theta = 0, 1$ has convergence rate 1 and $\theta = 0.5$ has convergence rate 2:

```
def test_convergence_rates():
    """Compare empirical convergence rates to exact ones."""
    # Set command-line arguments directly in sys.argv
    sys.argv[1:] = '--I 0.8 --a 2.1 --T 5 '\
                   '--dt 0.4 0.2 0.1 0.05 0.025'.split()
    # Suppress output from dc_mod.main()
    stdout = sys.stdout  # save standard output for later use
    scratchfile = open('.tmp', 'w')  # fake standard output
    sys.stdout = scratchfile

    r = dc_mod.main()
    for theta in r:
        nt.assert_true(r[theta])  # check for non-empty list

    scratchfile.close()
```

```
    sys.stdout = stdout   # restore standard output

    expected_rates = {0: 1, 1: 1, 0.5: 2}
    for theta in r:
        r_final = r[theta][-1]
        # Compare to 1 decimal place
        nt.assert_almost_equal(expected_rates[theta], r_final,
                               places=1, msg='theta=%s' % theta)

# no need for any main
```

Nothing more is needed in the `test_dc_nose.py`. file. Running `nosetests` will report `Ran 3 tests` and an `OK` for success. Everytime we modify the `decay_theta_unittest` module we can run `nosetests` to quickly see if the edits have any impact on the test collection.

Installation of nose.   The `nose` does not come with a standard Python distribution and must therefore be installed separately. The procedure is standard and described on Nose's web pages. On Debian-based Linux systems the command is `sudo apt-get install python-nose`, and with MacPorts you run `sudo port install py27-nose`.

3.5   Classical unit testing with unittest

The classical way of implementing unit tests derives from the JUnit tool in Java where all tests are methods in a class for testing. Python comes with a module `unittest` for doing this type of unit tests. While `nose` allows simple functions for unit tests, `unittest` requires deriving a class `Test*` from `unittest.TestCase` and implementing each test as methods with names `test_*` in that class.

Basic use of unittest.   Using the `double` function in the `mymod` module introduced in the previous section, unit testing with the aid of the `unittest` module consists of writing a file `test_mymod.py` with the content

```
import unittest
import mymod

class TestMyCode(unittest.TestCase):
    def test_double(self):
        result = mymod.double(4)
        self.assertEqual(result, 8)

if __name__ == '__main__':
    unittest.main()
```

The test is run by executing the test file `test_mymod.py` as a standard Python program. There is no support in `unittest` for automatically locating and running all tests in all test files in a folder tree.

45

Those who have experience with object-oriented programming will see that the difference between using `unittest` and `nose` is minor. Programmers with no or little knowledge of classes will certainly prefer `nose`.

Demonstration of unittest.   The same tests as shown for the nose framework are reimplemented with the `TestCase` classes in the file `test_dc_unittest.py`. The tests are identical, the only difference being that with `unittest` we must write the tests as methods in a class and the assert functions have slightly different syntax.

```
import unittest
import dc_mod_unittest as decay
import numpy as np

def exact_discrete_solution(n, I, a, theta, dt):
    factor = (1 - (1-theta)*a*dt)/(1 + theta*dt*a)
    return I*factor**n

class TestDecay(unittest.TestCase):

    def test_against_discrete_solution(self):
        ...
        diff = np.abs(u_de - u).max()
        self.assertAlmostEqual(diff, 0, delta=1E-14)

    def test_solver(self):
        ...
        for theta in 0, 0.5, 1:
            ...
            self.assertAlmostEqual(diff, 0, places=8,
                                   msg='theta=%s' % theta)

    def test_potential_integer_division():
        ...
        self.assertAlmostEqual(diff, 0, delta=1E-14)

    def test_convergence_rates(self):
        ...
        for theta in r:
            ...
            self.assertAlmostEqual(...)

if __name__ == '__main__':
    unittest.main()
```

3.6   Implementing simple problem and solver classes

The θ-rule was compactly and conveniently implemented in a function `solver`. In more complicated problems it might be beneficial to use classes and introduce a class `Problem` to hold the definition of the physical problem, a class `Solver` to hold the data and methods needed to numerically solve the problem, and a class `Visualizer` to make plots. This idea will now be illustrated, resulting in code that represents an alternative to the `solver` and `explore` functions found in the `dc_mod` module.

Explaining the details of class programming in Python is considered beyond the scope of this text. Readers who are unfamiliar with Python class programming should first consult one of the many electronic Python tutorials or textbooks to come up to speed with concepts and syntax of Python classes before reading on.

The problem class. The purpose of the problem class is to store all information about the mathematical model. This usually means all the physical parameters in the problem. In the current example with exponential decay we may also add the exact solution of the ODE to the problem class. The simplest form of a problem class is therefore

```
from numpy import exp

class Problem:
    def __init__(self, I=1, a=1, T=10):
        self.T, self.I, self.a = I, float(a), T

    def exact_solution(self, t):
        I, a = self.I, self.a
        return I*exp(-a*t)
```

We could in the `exact_solution` method have written `self.I*exp(-self.a*t)`, but using local variables `I` and `a` allows a formula `I*exp(-a*t)` which looks closer to the mathematical expression Ie^{-at}. This is not an important issue with the current compact formula, but is beneficial in more complicated problems with longer formulas. We will therefore often "strip off" `self` in variables in forthcoming examples.

The class data can be set either as arguments in the constructor or at any time later, e.g.,

```
problem = Problem(T=5)
problem.T = 8
problem.dt = 1.5
```

However, it would be convenient if class `Problem` could also initialize the data from the command line. To this end, we add a method for defining a set of command-line options and a method that sets the local attributes equal to what was found on the command line. The default values associated with the command-line options are taken as the values provided to the constructor. Class `Problem` now becomes

```
class Problem:
    def __init__(self, I=1, a=1, T=10):
        self.T, self.I, self.a = I, float(a), T

    def define_command_line_options(self, parser=None):
        if parser is None:
            import argparse
            parser = argparse.ArgumentParser()
```

```
        parser.add_argument(
            '--I', '--initial_condition', type=float,
            default=self.I, help='initial condition, u(0)',
            metavar='I')
        parser.add_argument(
            '--a', type=float, default=self.a,
            help='coefficient in ODE', metavar='a')
        parser.add_argument(
            '--T', '--stop_time', type=float, default=self.T,
            help='end time of simulation', metavar='T')
        return parser

    def init_from_command_line(self, args):
        self.I, self.a, self.T = args.I, args.a, args.T

    def exact_solution(self, t):
        I, a = self.I, self.a
        return I*exp(-a*t)
```

Observe that if the user already has an `ArgumentParser` object it can be supplied, but if we do not have, class `Problem` makes one for us. Python's `None` object is used to indicate that a variable is not initialized with a value.

The solver class. The solver class stores data related to the numerical solution method and provides a function `solve` for solving the problem. A problem object must be given to the constructor so that the solver can easily look up physical data. In the present example, the data related to the numerical solution method consist of Δt and θ. We add, as in the problem class, functionality for reading Δt and θ from the command line:

```
class Solver:
    def __init__(self, problem, dt=0.1, theta=0.5):
        self.problem = problem
        self.dt, self.theta = float(dt), theta

    def define_command_line_options(self, parser):
        parser.add_argument(
            '--dt', '--time_step_value', type=float,
            default=0.5, help='time step value', metavar='dt')
        parser.add_argument(
            '--theta', type=float, default=0.5,
            help='time discretization parameter', metavar='dt')
        return parser

    def init_from_command_line(self, args):
        self.dt, self.theta = args.dt, args.theta

    def solve(self):
        from dc_mod import solver
        self.u, self.t = solver(
            self.problem.I, self.problem.a, self.problem.T,
            self.dt, self.theta)

    def error(self):
        u_e = self.problem.exact_solution(self.t)
        e = u_e - self.u
```

```
        E = sqrt(self.dt*sum(e**2))
        return E
```

Note that we here simply reuse the implementation of the numerical method from the `dc_mod` module.

The visualizer class. The purpose of the visualizer class is to plot the numerical solution stored in class `Solver`. We also add the possibility to plot the exact solution. Access to the problem and solver objects is required when making plots so the constructor must store these objects:

```
class Visualizer:
    def __init__(self, problem, solver):
        self.problem, self.solver = problem, solver

    def plot(self, include_exact=True, plt=None):
        """
        Add solver.u curve to scitools plotting object plt,
        and include the exact solution if include_exact is True.
        This plot function can be called several times (if
        the solver object has computed new solutions).
        """
        if plt is None:
            import scitools.std
            plt = scitools.std

        plt.plot(self.solver.t, self.solver.u, '--o')
        plt.hold('on')
        theta2name = {0: 'FE', 1: 'BE', 0.5: 'CN'}
        name = self.solver.theta2name.get(self.solver.theta, '')
        plt.legend('numerical %s' % name)
        if include_exact:
            t_e = linspace(0, self.problem.T, 1001)
            u_e = self.problem.exact_solution(t_e)
            plt.plot(t_e, u_e, 'b-')
            plt.legend('exact')
        plt.xlabel('t')
        plt.ylabel('u')
        plt.title('theta=%g, dt=%g' %
                  (self.solver.theta, self.solver.dt))
        plt.savefig('%s_%g.png' % (name, self.solver.dt))
        return plt
```

The `plt` object in the `plot` method is worth a comment. The idea is that `plot` can add a numerical solution curve to an existing plot. Calling `plot` with a `plt` object, which has to be a `scitools.std` object in this implementation, will just add the curve `self.solver.u` as a dashed line with circles at the mesh points (leaving the color of the curve up to the plotting tool). This functionality allows plots with several solutions: just make a loop where new data is set in the problem and/or solver classes, the solver's `solve()` method is called, the most recent numerical solution is plotted by the `plot(plt)` method in the visualizer object (Exercise 10 describes a problem setting where this functionality is explored).

Combing the objects. Eventually we need to show how the classes `Problem`, `Solver`, and `Visualizer` play together:

```
def main():
    problem = Problem()
    solver = Solver(problem)
    viz = Visualizer(problem, solver)

    # Read input from the command line
    parser = problem.define_command_line_options()
    parser = solver. define_command_line_options(parser)
    args = parser.parse_args()
    problem.init_from_command_line(args)
    solver. init_from_command_line(args)

    # Solve and plot
    solver.solve()
    plt = viz.plot()
    E = solver.error()
    if E is not None:
        print 'Error: %.4E' % E
    plt.show()
```

The file `dc_class.py` constitutes a module with the three classes and the `main` function.

3.7 Implementing more advanced problem and solver classes

The previous `Problem` and `Solver` classes containing parameters soon get much repetitive code when the number of parameters increases. Much of this code can be parameterized and be made more compact. For this purpose, we decide to collect all parameters in a dictionary, `self.prms`, with two associated dictionaries `self.types` and `self.help` for holding associated object types and help strings. Provided a problem, solver, or visualizer class defines these three dictionaries in the constructor, using default or user-supplied values of the parameters, we can create a super class `Parameters` with general code for defining command-line options and reading them as well as methods for setting and getting a parameter. A `Problem` or `Solver` class will then inherit command-line functionality and the set/get methods from the `Parameters` class.

A generic class for parameters. A simplified version of the parameter class looks as follows:

```
class Parameters:
    def set(self, **parameters):
        for name in parameters:
            self.prms[name] = parameters[name]

    def get(self, name):
        return self.prms[name]

    def define_command_line_options(self, parser=None):
        if parser is None:
```

```
        import argparse
        parser = argparse.ArgumentParser()

    for name in self.prms:
        tp = self.types[name] if name in self.types else str
        help = self.help[name] if name in self.help else None
        parser.add_argument(
            '--' + name, default=self.get(name), metavar=name,
            type=tp, help=help)

    return parser

def init_from_command_line(self, args):
    for name in self.prms:
        self.prms[name] = getattr(args, name)
```

The file `class_dc_verf1.py` contains a slightly more advanced version of class `Parameters` where we in the `set` and `get` functions test for valid parameter names and raise exceptions with informative messages if any name is not registered.

The problem class. A class `Problem` for the exponential decay ODE with parameters a, I, and T can now be coded as

```
class Problem(Parameters):
    """
    Physical parameters for the problem u'=-a*u, u(0)=I,
    with t in [0,T].
    """
    def __init__(self):
        self.prms = dict(I=1, a=1, T=10)
        self.types = dict(I=float, a=float, T=float)
        self.help = dict(I='initial condition, u(0)',
                         a='coefficient in ODE',
                         T='end time of simulation')

    def exact_solution(self, t):
        I, a = self.get('I'), self.get('a')
        return I*np.exp(-a*t)
```

The solver class. Also the solver class is derived from class `Parameters` and works with the `prms`, `types`, and `help` dictionaries in the same way as class `Problem`. Otherwise, the code is very similar to class `Solver` in the `decay_class.py` file:

```
class Solver(Parameters):
    def __init__(self, problem):
        self.problem = problem
        self.prms = dict(dt=0.5, theta=0.5)
        self.types = dict(dt=float, theta=float)
        self.help = dict(dt='time step value',
                         theta='time discretization parameter')

    def solve(self):
```

51

```
        #from dc_mod import solver
        from decay_theta import solver
        self.u, self.t = solver(
            self.problem.get('I'),
            self.problem.get('a'),
            self.problem.get('T'),
            self.get('dt'),
            self.get('theta'))

    def error(self):
        try:
            u_e = self.problem.exact_solution(self.t)
            e = u_e - self.u
            E = np.sqrt(self.get('dt')*np.sum(e**2))
        except AttributeError:
            E = None
        return E
```

The visualizer class. Class `Visualizer` can be identical to the one in the `decay_class.py` file since the class does not need any parameters. However, a few adjustments in the `plot` method is necessary since parameters are accessed as, e.g., `problem.get('T')` rather than `problem.T`. The details are found in the file `class_dc_verf1.py`.

Finally, we need a function that solves a real problem using the classes `Problem`, `Solver`, and `Visualizer`. This function can be just like `main` in the `class_dc_v1.py` file.

The advantage with the `Parameters` class is that it scales to problems with a large number of physical and numerical parameters: as long as the parameters are defined once via a dictionary, the compact code in class `Parameters` can handle any collection of parameters of any size.

4 Performing scientific experiments

The goal of this section is to explore the behavior of a numerical method for a differential equation and show how scientific experiments can be set up and reported. We address the ODE problem

$$u'(t) = -au(t), \quad u(0) = I, \quad 0 < t \leq T, \tag{39}$$

numerically discretized by the θ-rule:

$$u^{n+1} = \frac{1 - (1 - \theta)a\Delta t}{1 + \theta a\Delta t}u^n, \quad u^0 = I.$$

Our aim is to plot u^n against the exact solution $u_e = Ie^{-at}$ for various choices of the parameters in this numerical problem: I, a, Δt, and θ. In particular, we are interested in how the discrete solution compares with the exact solution when the Δt parameter is varied and θ takes on the three values corresponding to the Forward Euler, Backward Euler, and Crank-Nicolson schemes ($\theta = 0, 1, 0.5$, respectively).

A verified implementation for computing u^n and plotting u^n together with u_{e} is found in the file `dc_mod.py`. This program admits command-line arguments to specify a series of Δt values and will run a loop over these values and $\theta = 0, 0.5, 1$. We make a slight edit of how the plots are designed: the numerical solution is specified with line type `'r--o'` (dashed red lines with dots at the mesh points), and the `show()` command is removed to avoid a lot of plot windows popping up on the computer screen (but hardcopies of the plot are still stored in files via `savefig`). The slightly modified program has the name `experiments/dc_mod.py`. All files associated with the scientific investigation are collected in a subfolder `experiments`.

Running the experiments is easy since the `dc_mod.py` program already has the loops over θ and Δt implemented:

```
                               Terminal
Terminal> python dc_mod.py --I 1 --a 2 --makeplot \
          --T 5 --dt 0.5 0.25 0.1 0.05
```

A lot of image files `FE_*.png`, `BE_*.png`, and `CN_*.png` are generated. We want to combine all the `FE_*.png` files in a table fashion in one file, with two images in each row, starting with the largest Δt in the upper left corner and decreasing the value as we go to the right and down. This can be done using the `montage` program:

```
                               Terminal
Terminal> montage -background white -geometry 100% -tile 2x \
          file1.png file2.png ... result.png
```

Running manual commands is boring, and errors may easily sneak in. Both for automating manual work and documenting the operating system commands we actually issued in the experiment, we should write a *script* (little program). The script takes a list of Δt values on the command line as input and makes three combined images, one for each θ value, displaying the quality of the numerical solution as Δt varies. For example,

```
                               Terminal
Terminal> python dc_exper0.py 0.5 0.25 0.1 0.05
```

results in three images `FE.png`, `CN.png`, and `BE.png`, each with four plots corresponding to the four Δt values. Each plot compares the numerical solution with the exact one. The latter image is shown in Figure 9.

Ideally, the script should be scalable in the sense that it works for any number of Δt values, which is the case for this particular implementation:

Figure 9: Illustration of the Backward Euler method for four time step values.

```
import os, sys, glob

# The command line must contain dt values
if len(sys.argv) > 1:
    dt_values = [float(arg) for arg in sys.argv[1:]]
else:
    print 'Usage: %s dt1 dt2 dt3 ...';  sys.exit(1)  # abort

# Fixed physical parameters
I = 1
a = 2
T = 5

# Run module file as a stand-alone application
cmd = 'python dc_mod.py --I %g --a %g --makeplot --T %g' % \
      (I, a, T)
dt_values_str = ' '.join([str(v) for v in dt_values])
cmd += ' --dt %s' % dt_values_str
print cmd
failure = os.system(cmd)
if failure:
    print 'Command failed:', cmd; sys.exit(1)
```

```
# Combine images into rows with 2 plots in each row
montage_commands = []
for method in 'BE', 'CN', 'FE':
    imagefiles = ['%s_%s.png' % (method, dt) for dt in dt_values]
    montage_commands.append(
        'montage -background white -geometry 100%' + \
        ' -tile 2x %s %s.png' % (' '.join(imagefiles), method))

for cmd in montage_commands:
    print cmd
    failure = os.system(cmd)
    if failure:
        print 'Command failed:', cmd; sys.exit(1)

# Remove the files generated by dc_mod.py
filenames = glob.glob('*_*.png')
for filename in filenames:
    os.remove(filename)
```

This file is available as `experiments/dc_exper0.py`.

We may comment upon many useful constructs in this script:

- `[float(arg) for arg in sys.argv[1:]]` builds a list of real numbers from all the command-line arguments.

- `failure = os.system(cmd)` runs an operating system command, e.g., another program. The execution successful only if `failure` is zero.

- Unsuccessful execution usually makes it meaningless to continue the program, and therefore we abort the program with `sys.exit(1)`. Any argument different from 0 signifies that our program stopped with a failure.

- `['%s_%s.png' % (method, dt) for dt in dt_values]` builds a list of filenames from a list of numbers (`dt_values`).

- All `montage` commands for creating composite figures are stored in a list and thereafter executed in a loop.

- `glob.glob('*_*.png')` returns a list of the names of all files in the current folder where the filename matches the *Unix wildcard notation* `*_*.png` (meaning "any text, underscore, any text, and then '.png'").

- `os.remove(filename)` removes the file with name `filename`.

4.1   Interpreting output from other programs

Programs that run other programs, like `dc_exper0.py` does, will often need to interpret output from the other programs. Let us demonstrate how this is done in Python by extracting the relations between θ, Δt, and the error E as written to the terminal window by the `dc_mod.py` program, which is being executed by `dc_exper0.py`. We will

55

- read the output from the `dc_mod.py` program

- interpret this output and store the E values in arrays for each θ value

- plot E versus Δt, for each θ, in a log-log plot

The simple `os.system(cmd)` call does not allow us to read the output from running `cmd`. Instead we need to invoke a bit more involved procedure:

```
from subprocess import Popen, PIPE, STDOUT
p = Popen(cmd, shell=True, stdout=PIPE, stderr=STDOUT)
output, dummy = p.communicate()
failure = p.returncode
if failure:
    print 'Command failed:', cmd; sys.exit(1)
```

The command stored in `cmd` is run and all text that is written to the standard output *and* the standard error is available in the string `output`. The text in `output` is what appeared in the terminal window while running `cmd`.

Our next task is to run through the `output` string, line by line, and if the current line prints θ, Δt, and E, we split the line into these three pieces and store the data. The chosen storage structure is a dictionary `errors` with keys `dt` to hold the Δt values, and three θ keys to hold the corresponding E values. The relevant code lines are

```
errors = {'dt': dt_values, 1: [], 0: [], 0.5: []}
for line in output.splitlines():
    words = line.split()
    if words[0] in ('0.0', '0.5', '1.0'):  # line with E?
        # typical line: 0.0   1.25:     7.463E+00
        theta = float(words[0])
        E = float(words[2])
        errors[theta].append(E)
```

Note that we do not bother to store the Δt values as we read them from `output`, because we already have these values in the `dt_values` list.

We are now ready to plot E versus Δt for $\theta = 0, 0.5, 1$:

```
import matplotlib.pyplot as plt
#import scitools.std as plt
plt.loglog(errors['dt'], errors[0], 'ro-')
plt.hold('on')  # MATLAB style...
plt.loglog(errors['dt'], errors[0.5], 'b+-')
plt.loglog(errors['dt'], errors[1], 'gx-')
plt.legend(['FE', 'CN', 'BE'], loc='upper left')
plt.xlabel('log(time step)')
plt.ylabel('log(error)')
plt.title('Error vs time step')
plt.savefig('error_BE_CN_FE.png')
```

Plots occasionally need some manual adjustments. Here, the axis of the log-log plot look nicer if we adapt them strictly to the data, see Figure 10. To this end, we need to compute $\min E$ and $\max E$, and later specify the extent of the axes:

```
errors = {'dt': dt_values, 1: [], 0: [], 0.5: []}
min_E = 1E+20; max_E = -min_E  # keep track of min/max E for axis
for line in output.splitlines():
    ...
        ...
        min_E = min(min_E, E);  max_E = max(max_E, E)

import matplotlib.pyplot as plt
plt.loglog(errors['dt'], errors[0], 'ro-')
...
plt.axis([min(dt_values), max(dt_values), min_E, max_E])
...
```


Figure 10: Default plot (left) and manually adjusted plot (right).

The complete program, incorporating the code snippets above, is found in `experiments/dc_exper1.py`. This program can hopefully be reused in a number of other occasions where one needs to run experiments, extract data from the output of programs, make plots, and combine several plots in a figure file.

4.2   Making a report

The results of running computer experiments are best documented in a little report containing the problem to be solved, key code segments, and the plots from a series of experiments. At least the part of the report containing the plots should be automatically generated by the script that performs the set of experiments, because in that script we know exactly which input data that were used to generate a specific plot, thereby ensuring that each figure is connected to the right data. Take a look at an example to see what we have in mind:

> http://hplgit.github.com/INF5620/doc/writing_reports/sphinx-cloud/

Plain HTML.   Scientific reports can be written in a variety of formats. Here we begin with the HTML format which allows efficient browsing of all the experiments in any web browser. An extended version of the `dc_exper1.py` from the

57

last section, called `dc_exper1_html.py`, adds code at the end for creating an HTML file with a summary, a section on the mathematical problem, a section on the numerical method, a section on the `solver` function implementing the method, and a section with subsections containing figures that shows the results of experiments where Δt is varied for $\theta = 0, 0.5, 1$. The mentioned Python file contains all the details for writing this HTML report.

HTML with MathJax. Scientific reports usually need mathematical formulas and hence mathematical typesetting. In plain HTML, as used in the `dc_exper1_html.py` file, we have to use just the keyboard characters to write mathematics. However, there is an extension to HTML, called MathJax, that allows formulas and equations to be typeset with LaTeX syntax and nicely rendered in web browsers. A relatively small subset of LaTeX environments is supported, but the syntax for formulas is quite rich. Inline formulas are look like `\( u'=-au \)` while equations are surrounded by `$$` signs. Inside such signs, one can use `\[ u'=-au \]` for unnumbered equations, or `\begin{equation}` and `\end{equation}` surrounding `u'=-au` for numbered equations, or `\begin{align}` and `\end{align}` for multiple aligned equations, with (`align`) or without (`align*`) numbers and labels.

The file `dc_exper1_mathjax.py` contains all the details for turning the previous plain HTML report into web pages with nicely typeset mathematics.

LaTeX. The *de facto* language for mathematical typesetting and scientific report writing is LaTeX. A number of very sophisticated packages have been added to the language over a period of three decades, allowing very fine-tuned layout and typesetting. For output in the PDF format, LaTeX is the definite choice when it comes to quality. The LaTeX language used to write the reports has typically a lot of commands involving backslashes and braces. For output on the web, using HTML (and not the PDF directly in the browser window), LaTeX struggles with delivering high quality typesetting. Other tools, especially Sphinx, gives better results and can also produce nice-looking PDFs.

Sphinx. Sphinx is a typesetting language with similarities to HTML and LaTeX, but with much less tagging. It has recently become very popular for software documentation and mathematical reports. Sphinx can utilize MathJax or LaTeX for mathematical formulas and equations, but has limitations compared to both tools. The Sphinx syntax is an extension of the reStructuredText language, and comes with rich support for fancy layout of web pages. In particular, Sphinx can easily be combined with various layout *themes* that give a certain look and feel to the web site and that offers table of contents, navigation, search facilities, etc.

Markdown. A recently popular format for easy writing of web pages is Markdown. Text is written very much like one would do in email, using spacing and special characters to naturally format the code instead of heavily tagging the

text as in LaTeX and HTML. With the tool Pandoc one can go from Markdown to a variety of formats. HTML is a common output format, but LaTeX, epub, XML, OpenOffice, MediaWiki, and MS Word are some other possibilities.

Wiki formats. A range of wiki formats are popular for creating notes on the web, especially documents which allow groups of people to edit and evolve the content. Apart from MediaWiki (the wiki format used for Wikipedia), wiki formats have no support for mathematical typesetting and also limited tools for displaying computer code in nice ways. Wiki formats are therefore less suitable for scientific reports compared to the other formats mentioned here.

Doconce. Since it is difficult to choose the right tool or format for writing a scientific report, it is advantageous to write the content in a format that easily can be translated to LaTeX, HTML, Sphinx, Markdown, and wikis. Doconce is such a tool. It is similar to Pandoc, but offers some special convenient features for writing about mathematics and programming. The tagging is modest, somewhere between LaTeX and Markdown.

The HTML, LaTeX PDF, Sphinx, and Doconce formats for the scientific report whose content is outlined above, are exemplified with source codes and results at the web pages associated with this teaching material.

4.3 Publishing a complete project

A report documenting scientific investigations should be accompanied by all the software and data used for the investigations so that others have a possibility to redo the work and assess the qualify of the results. This possibility is important for *reproducible research* and hence reaching reliable scientific conclusions.

One way of documenting a complete project is to make a folder tree with all relevant files. Preferably, the tree is published at some project hosting site like Bitbucket, GitHub, or Googlecode so that others can download it as a tarfile, zipfile, or clone the files directly using a version control system like Mercurial or Git. For the investigations outlined in Section 4.2, we can create a folder tree with files

```
src/dc_exper1_mathjax.py
doc/report.html
doc/run.sh
```

The `run.sh` file is a simple Bash script listing the `python` command we used to generate the experiments that are documented in `report.html`.

5 Exercises

Exercise 1: Experiment with integer division

Explain what happens in the following computations, where some are (mathematically) unexpected:

```
>>> dt = 1
>>> T = 5
>>> N = T/dt
>>> N
5
>>> type(N)
<type 'int'>
>>> from numpy import linspace
>>> linspace(0, 1, 2)
array([ 0.,   1.])
>>> linspace(0, 1, 2.4)
array([ 0.,   1.])
>>> linspace(0, 1, 2.9)
array([ 0.,   1.])
>>> theta = 1
>>> a = 1
>>> (1 - (1-theta)*a*dt)/(1 + theta*dt*a)
0
>>> (1 - (1-theta)*a*dt)
1
>>> (1 + theta*dt*a)
2
```

Filename: `pyproblems.txt`.

Exercise 2: Experiment with wrong computations

Consider the `solver` function in the `dc_v1.py` and the following call:

```
u, t = solver(I=1, a=1, T=7, dt=2, theta=1)
```

The output becomes

```
    t= 0.000 u=1
    t= 2.000 u=0
    t= 4.000 u=0
    t= 6.000 u=0
```

Print out the result of all intermediate computations and use `type(v)` to see the object type of the result stored in `v`. Examine the intermediate calculations and explain why `u` is wrong and why we compute up to $t = 6$ only even though we specified $T = 7$. Filename: `dc_v1_err.py`.

Exercise 3: Implement specialized functions

Implement a specialized Python function `ForwardEuler` that solves the problem $u' = -au$, $u(0) = I$, using the Forward Euler method. Do not reimplement the solution algorithm, but let the `ForwardEuler` function call the `solver` function. Import this latter function from the module `dc_mod`. Also make similar functions `BackwardEuler` and `CrankNicolson`. Filename: `decay_FE_BE_CN.py`.

Exercise 4: Plot the error function

Solve the problem $u' = -au$, $u(0) = I$, using the Forward Euler, Backward Euler, and Crank-Nicolson schemes. For each scheme, plot the error function $e_n = u_e(t_n) - u^n$ for Δt, $\frac{1}{4}\Delta t$, and $\frac{1}{8}\Delta t$, where u_e is the exact solution of the ODE and u^n is the numerical solution at mesh point t_n. Filename: `decay_plot_error.py`.

Exercise 5: Compare methods for a give time mesh

Make a program that imports the `solver` function from the `dc_mod` module and offers a function `compare(dt, I, a)` for comparing, in a plot, the methods corresponding to $\theta = 0, 0.5, 1$ and the exact solution. This plot shows the accuracy of the metods for a given time mesh. Read input data for the problem from the command line using appropriate functions in the `dc_mod` module (the `--dt` option for giving several time step values can be reused: just use the first time step value for the computations). Filename: `decay_compare_theta.py`.

Exercise 6: Change formatting of numbers and debug

The `dc_memsave.py` program writes the time values and solution values to a file which looks like

```
0.0000000000000000E+00   1.0000000000000000E+00
2.0000000000000001E-01   8.3333333333333337E-01
4.0000000000000002E-01   6.9444444444444453E-01
6.0000000000000009E-01   5.7870370370370383E-01
8.0000000000000004E-01   4.8225308641975323E-01
1.0000000000000000E+00   4.0187757201646102E-01
1.2000000000000000E+00   3.3489797668038418E-01
1.3999999999999999E+00   2.7908164723365347E-01
```

Modify the file output such that it looks like

```
0.000   1.00000
0.200   0.83333
0.400   0.69444
0.600   0.57870
0.800   0.48225
1.000   0.40188
1.200   0.33490
1.400   0.27908
```

Run

Terminal

```
Terminal> python decay8_v2.py --T 10 --theta 1 --dt 0.2 --makeplot
```

The program just prints `Bug in the implementation!` and does not show the plot. What went wrong? Filename: `dc_memsave_v2.py`.

Answer. With only 5 decimals in the file, the `verify` function compares truncated elements `u`, accurate only to 10^{-5} with the exact discrete solution and applies a far too small `tol` value. `tol` must be `1E-4`.

Exercise 7: Write a doctest

Type in the following program and equip the `roots` function with a doctest:

```
import sys
# This sqrt(x) returns real if x>0 and complex if x<0
from numpy.lib.scimath import sqrt

def roots(a, b, c):
    """
    Return the roots of the quadratic polynomial
    p(x) = a*x**2 + b*x + c.

    The roots are real or complex objects.
    """
    q = b**2 - 4*a*c
    r1 = (-b + sqrt(q))/(2*a)
    r2 = (-b - sqrt(q))/(2*a)
    return r1, r2

a, b, c = [float(arg) for arg in sys.argv[1:]]
print roots(a, b, c)
```

Make sure to test both real and complex roots. Write out numbers with 14 digits or less. Filename: `doctest_roots.py`.

Exercise 8: Write a nose test

Make a nose test for the `roots` function in Exercise 7. Filename: `test_roots.py`.

Exercise 9: Make a module

Let

$$q(t) = \frac{RAe^{at}}{R + A(e^{at} - 1)} .$$

Make a Python module `q_module` containing two functions `q(t)` and `dqdt(t)` for computing $q(t)$ and $q'(t)$, respectively. Perform a `from numpy import *` in this module. Import `q` and `dqdt` in another file using a construction `from q_module import *`. All objects available in this file is given by `dir()`. Print `dir()` and `len(dir())`. Then change the import of `numpy` in `q_module.py` to `import numpy as np`. What is the effect of this import on the number of objects in `dir()` in a file that does `from q_module import *`?

Filename: `q_module.py`.

Exercise 10: Make use of a class implementation

We want to solve the exponential decay problem $u' = -au$, $u(0) = I$, for several Δt values and $\theta = 0, 0.5, 1$. For each Δt value, we want to make a plot where the three solutions corresponding to $\theta = 0, 0.5, 1$ appear along with the exact solution. Write a function `experiment` to accomplish this. The function should import the classes `Problem`, `Solver`, and `Visualizer` from the `decay_class1`

module and make use of these. A new command-line option `--dt_values` must be added to allow the user to specify the Δt values on the command line (the options `--dt` and `--theta` have then no effect when running the `experiment` function). Note that the classes in the `decay_class1` module should *not* be modified. Filename: `dc_class_exper.py`.

6 Analysis of the θ-rule for a decay ODE

We address the ODE for exponential decay,

$$u'(t) = -au(t), \quad u(0) = I, \tag{40}$$

where a and I are given constant. This problem is solved by the θ-rule finite difference scheme, resulting in the recursive equations

$$u^{n+1} = \frac{1 - (1 - \theta)a\Delta t}{1 + \theta a\Delta t}u^n \tag{41}$$

for the numerical solution u^{n+1}, which approximates the exact solution u_{e} at time point t_{n+1}. For constant mesh spacing, which we assume here, $t_{n+1} = (n+1)\Delta t$.

6.1 Discouraging numerical solutions

Choosing $I = 1$, $a = 2$, and running experiments with $\theta = 1, 0.5, 0$ for $\Delta t = 1.25, 0.75, 0.5, 0.1$, gives the results in Figures 11, 12, and 13.

The characteristics of the displayed curves can be summarized as follows:

- The Backward Euler scheme always give a monotone solution, lying above the exact curve.

- The Crank-Nicolson scheme gives the most accurate results, but for $\Delta t = 1.25$ the solution oscillates.

- The Forward Euler scheme gives a growing, oscillating solution for $\Delta t = 1.25$; a decaying, oscillating solution for $\Delta t = 0.75$; a strange solution $u^n = 0$ for $n \geq 1$ when $\Delta t = 0.5$; and a solution seemingly as accurate as the one by the Backward Euler scheme for $\Delta t = 0.1$, but the curve lies below the exact solution.

Since the exact solution of our model problem is a monotone function, $u(t) = Ie^{-at}$, some of these results are indeed alarming!

6.2 Experimental investigation of oscillatory solutions

We may ask the question: Under what circumstances, i.e., values of the input data I, a, and Δt will the Forward Euler and Crank-Nicolson schemes result in oscillatory solutions?
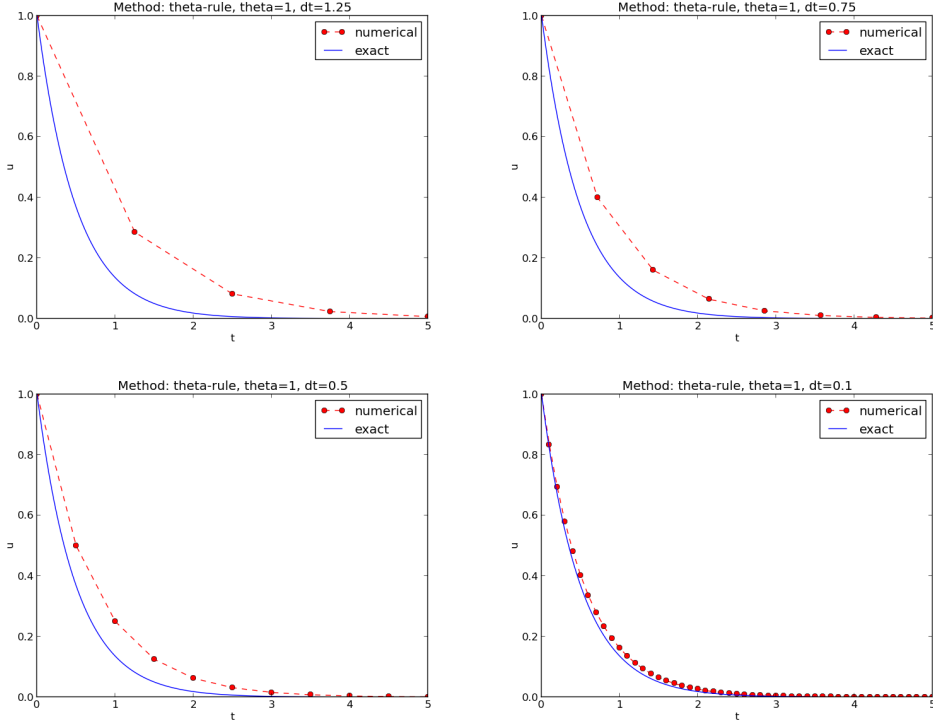
Figure 11: Backward Euler.

We may set up an experiment where we loop over values of I, a, and Δt. For each experiment, we flat the solution as oscillatory if

$$u^n > u^{n-1},$$

for some value of n, since we expect u^n to decay with n, but oscillations make u increase over a time step. We will quickly see that oscillations are independent of I, but do depend on a and Δt. Therefore, we introduce a two-dimensional function $B(a, \Delta t)$ which is 1 if oscillations occur and 0 otherwise. We can visualize B as a contour plot (lines for which $B = \text{const}$). The contour $B = 0.5$ will correspond to the borderline between oscillator regions $B = 1$ and monotone regions in the $a - \Delta t$ plane.

The B function is defined at discrete a and Δt values. Say we have given P a values, $a_0, \ldots, a_{P-1}$, and Q Δt values, $\Delta t_0, \ldots, \Delta t_{Q-1}$. These a_i and Δt_j values, $i = 0, \ldots, P-1$, $j = 0, \ldots, Q-1$, form a rectangular mesh of $P \times Q$ points in the plane. At each point $(a_i, \Delta t_j)$, we associate the corresponding value of $B(a_i, \Delta t_j)$, denoted B_{ij}. The B_{ij} values are naturally stored in a two-dimensional array. Both Matplotlib and SciTools can create a plot of the contour line $B_{ij} = 0.5$ dividing the oscillatory and monotone regions. The file
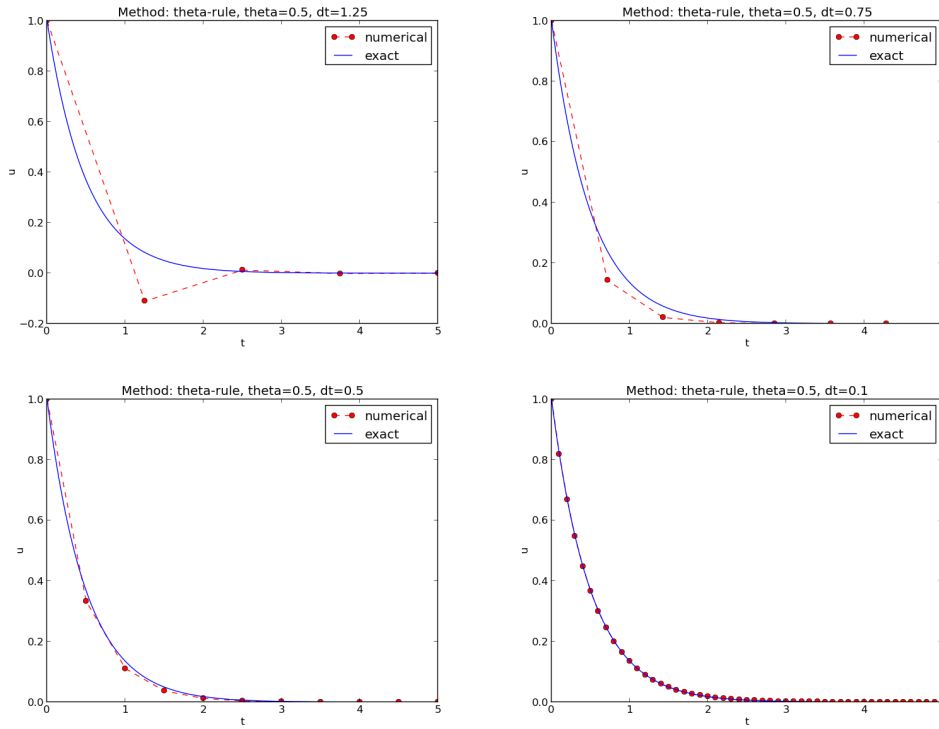
Figure 12: Crank-Nicolson.

`dc_osc_regions.py` contains all nuts and bolts to produce the $B = 0.5$ line in Figures 14 and decay:analysis:B:CN (the oscillatory region is above the line):

```python
from dc_mod import solver
import numpy as np
import scitools.std as st

def non_physical_behavior(I, a, T, dt, theta):
    """
    Given lists/arrays a and dt, and numbers I, dt, and theta,
    make a two-dimensional contour line B=0.5, where B=1>0.5
    means oscillatory (unstable) solution, and B=0<0.5 means
    monotone solution of u'=-au.
    """
    a = np.asarray(a); dt = np.asarray(dt)  # must be arrays
    B = np.zeros((len(a), len(dt)))         # results
    for i in range(len(a)):
        for j in range(len(dt)):
            u, t = solver(I, a[i], T, dt[j], theta)
            # Does u have the right monotone decay properties?
            correct_qualitative_behavior = True
            for n in range(1, len(u)):
                if u[n] > u[n-1]:  # Not decaying?
                    correct_qualitative_behavior = False
```
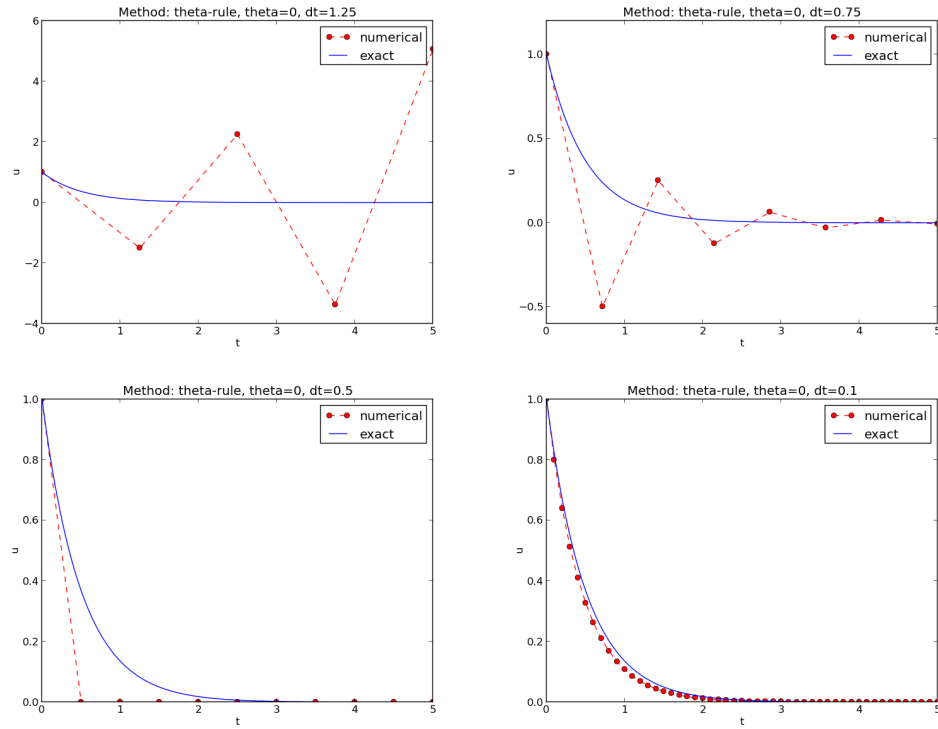
Figure 13: Forward Euler.

```
                      break  # Jump out of loop
              B[i,j] = float(correct_qualitative_behavior)
    a_, dt_ = st.ndgrid(a, dt)  # make mesh of a and dt values
    st.contour(a_, dt_, B, 1)
    st.grid('on')
    st.title('theta=%g' % theta)
    st.xlabel('a'); st.ylabel('dt')
    st.savefig('osc_region_theta_%s.png' % theta)
    st.savefig('osc_region_theta_%s.eps' % theta)

non_physical_behavior(
    I=1,
    a=np.linspace(0.01, 4, 22),
    dt=np.linspace(0.01, 4, 22),
    T=6,
    theta=0.5)
```

By looking at the curves in the figures one may guess that $a\Delta t$ must be less than a critical limit to avoid the undesired oscillations. This limit seems to be about 2 for Crank-Nicolson and 1 for Forward Euler. We shall now establish a mathematical analysis of the discrete model that can explain the observations in our numerical experiments.

Figure 14: Forward Euler scheme.

Figure 15: Crank-Nicolson scheme.

67

6.3 Exact numerical solution

Starting with $u^0 = I$, the simple recursion (41) can be applied repeatedly n times, with the result that

$$u^{n+1} = \frac{1 - (1-\theta)a\Delta t}{1 + \theta a\Delta t} u^n \,. \tag{42}$$

Difference equations where all terms are linear in u^{n+1}, u^n, and maybe u^{n-1}, u^{n-2}, etc., are called *homogeneous, linear* difference equations, and their solutions are generally of the form $u^n = A^n$. Inserting this expression and dividing by A^{n+1} gives a polynomial equation in A. In the present case we get

$$A = \frac{1 - (1-\theta)a\Delta t}{1 + \theta a\Delta t} \,.$$

This is a solution technique of wider applicability than repeated use of the recursion (41).

Regardless of the solution approach, we have obtained a formula for u^n. This formula can explain everything what we see in the figures above, but it also gives us a more general insight into accuracy and stability properties of the three schemes.

6.4 Stability

Since u^n is a factor

$$A = \frac{1 - (1-\theta)a\Delta t}{1 + \theta a\Delta t} \tag{43}$$

raised to an integer power (n), we realize that $A < 0$ will for odd powers imply $u^n < 0$ and for even power result in $u^n > 0$, i.e., a solution that oscillates between the mesh points. We have that $A < 0$ when

$$(1-\theta)a\Delta t > 1 \,. \tag{44}$$

Since $A > 0$ is a requirement for having a numerical solution with the same basic property (monotonicity) as the exact solution, we may say that $A > 0$ is a *stability criterion*. Expressed in terms of Δt the stability criterion reads

$$\Delta t \le \frac{1}{(1-\theta)a} \,. \tag{45}$$

The Backward Euler scheme is always stable (since $A < 0$ is impossible, while non-oscillating solutions for Forward Euler and Crank-Nicolson demand $\Delta t \le 1/a$ and $\Delta t \le 2/a$, respectively. The relation between Δt and a look reasonable: a smaller a means faster decay and hence a need for smaller time steps.

Looking at Figure 13, we see that with $a\Delta t = 2 \cdot 1.25 = 2.5$, $A = -1.5$, and the ssolution $u^n = (-1.5)^n$ oscillates *and* grows. With $a\Delta t = 2 \cdot 0.75 = 1.5$,

$A = -0.5$, $u^n = (-0.5)^n$ decays but oscillates. The peculiar case $\Delta t = 0.5$, where the Forward Euler scheme produces a solution that is stuck on the t axis, corresponds to $A = 0$ and therefore $u^0 = I = 1$ and $u^n = 0$ for $n \geq 1$. The decaying oscillations in the Crank-Nicolson scheme for $\Delta t = 1.25$ is easily explaned by $A = -0.25$.

The factor A is called *amplification factor* since the solution at a time level times A gives the solution at the next time level. For a decay process, we must obviously have $|A| \leq 1$ for all Δt, which is fulfilled for $\theta \geq 1/2$. Arbitrarily large values of u can be generated when $|A| > 1$ and n is large enough. The numerical solution is then totally irrelvant to the ODE modeling decay processes.

6.5 Comparing Amplification Factors

After establishing how A impacts the qualitative features of the solution, we mau now look more into how well the numerical amplification factor approximates the exact one. The exact solution reads $u(t) = Ie^{-at}$, which can be rewritten as

$$u_e(t_n) = Ie^{-an\Delta t} = I(e^{-a\Delta t})^n .\tag{46}$$

From this formula we see that the exact amplification factor is

$$A_e = e^{-a\Delta t} .\tag{47}$$

We realize that the amplification factors depend on a and Δt through the product $a\Delta t$. Therefore, it is convenient to introduce a symbol for this product, $p = a\Delta t$, and view A and A_e as functions of p. Figure 16 shows these functions. Crank-Nicolson is clearly closest to the exact amplification factor, but that method has the unfortunate oscillatory behavior when $p > 2$.

6.6 Series Expansion of Amplification Factors

As an alternative to the visual understanding inherent in Figure 16, there is a strong tradition in numerical analysis to investigate approximation errors when the discretization parameter, here Δt, becomes small. In the present case we let p be our small discretization parameter, and it makes sense to simplify the expressions for A and A_e by using Taylor polynomials around $p = 0$. The Taylor polynomials are accurate for small p and greatly simplifies the comparison of the analytical expressions since we then can compare polynomials, term by term.

Calculating the Taylor series for A_e is easily done by hand, but the three versions of A for $theta = 0, 1, \frac{1}{2}$ lead to more cumbersome calculations. Nowadays, analytical computations can benefit greatly by symbolic computer algebra system. The Python package `sympy` represents a powerful computer algebra system, not as sophisticated as the famous Maple and Mathematica systems, but free and very easy to integrate with our numerical computations in Python.

When using `sympy`, it is convenient to enter the interactive Python mode where we can write expressions and statements and immediately see the results.

Figure 16: Comparison of amplification factors.

Here is a simple example. We strongly recommend to use IPython for interactive sessions, although our typesetting will apply the prompt `>>>` (associated with the primitive Python shell that results from writing `python` in a terminal window).

Let us enter `sympy` in a Python shell and show how we can find the Taylor series for e^{-p}:

```
>>> from sympy import *
>>> p = Symbol('p')  # all variables must be declared as Symbols
>>> A_e = exp(-p)
>>>
>>> # First 6 terms of the Taylor series of A_e
>>> A_e.series(p, 6)
1 + (1/2)*p**2 - p - 1/6*p**3 - 1/120*p**5 + (1/24)*p**4 + O(p**6)
```

Lines with `>>>` represent input lines and lines without this prompt represents the result of computations. Apart from the order of the powers, the computed formula is easily recognized as the beginning of the Taylor series for e^{-p}.

Let us define the numerical amplification factor where p and θ enter the formula as symbols:

```
>>> theta = Symbol('theta')
>>> A = (1-(1-theta)*p)/(1+theta*p)
```

To work with the factor for the Backward Euler scheme we can substitute a value for `theta`, here 1:

70

```
>>> A.subs(theta, 1)
1/(1 + p)
```

Similarly, we can substitute `theta` by $1/2$ for Crank-Nicolson, preferably using an exact rational representation of $1/2$ in `sympy`:

```
>>> half = Rational(1,2)
>>> A.subs(theta, half)
1/(1 + (1/2)*p)*(1 - 1/2*p)
```

The Taylor series of the amplification factor for the Crank-Nicolson scheme can be computed as

```
>>> half = Rational(1,2)
>>> A.subs(theta, half).series(p, 4)
1 + (1/2)*p**2 - p - 1/4*p**3 + O(p**4)
```

We are now in a position to compare Taylor series:

```
>>> FE = A_e.series(p, 4) - A.subs(theta, 0).series(p, 4)
>>> BE = A_e.series(p, 4) - A.subs(theta, 1).series(p, 4)
>>> CN = A_e.series(p, 4) - A.subs(theta, half).series(p, 4)
>>> FE
(1/2)*p**2 - 1/6*p**3 + O(p**4)
>>> BE
-1/2*p**2 + (5/6)*p**3 + O(p**4)
>>> CN
(1/12)*p**3 + O(p**4)
```

From these expressions we see that the error $A - A_e \sim \mathcal{O}(p^2)$ for the Forward and Backward Euler schemes, while $A - A_e \sim \mathcal{O}(p^3)$ for the Crank-Nicolson scheme. It is the *leading order term*, i.e., the term of the lowest order (degree), that is of interest, because as $p \to 0$, this term is (much) bigger than the higher-order terms.

Now, a is a given parameter in the problem, while Δt is what we can vary. One therefore usually writes the error expressions in terms Δt. When then have

$$A - A_e = \begin{cases} \mathcal{O}(\Delta t^2), & \text{Forward and Backward Euler,} \\ \mathcal{O}(\Delta t^3), & \text{Crank-Nicolson} \end{cases} \tag{48}$$

What is the significance of this result? If we halve Δt, the error in amplification factor at a time level will be reduced by a factor of 4 in the Forward and Backward Euler schemes, and by a factor of 8 in the Crank-Nicolson scheme. That is, as we reduce Δt to obtain more accurate results, the Crank-Nicolson scheme can be efficiently reduce the error than the other schemes.

6.7   Local error

Definition of local error or (local truncation error, no don't use truncation here, that is with Taylor series and residual): take one time step (from some arbitrary "initial" time) and measure the error over this single step.

71

Global error: introduce some norm of the difference of the solutions over $[0, T]$ (max/infinity norm, L^p).

Both errors can be investigated analytically.

6.8   Analytical comparison of schemes

An alternative comparison of the schemes is to look at the ratio A/A_e, or the error $1 - A/A_e$ in this ratio:

```
>>> FE = 1 - (A.subs(theta, 0)/A_e).series(p, 4)
>>> BE = 1 - (A.subs(theta, 1)/A_e).series(p, 4)
>>> CN = 1 - (A.subs(theta, half)/A_e).series(p, 4)
>>> FE
(1/2)*p**2 + (1/3)*p**3 + O(p**4)
>>> BE
-1/2*p**2 + (1/3)*p**3 + O(p**4)
>>> CN
(1/12)*p**3 + O(p**4)
```

The leading-order terms have the same powers as in the same as in the analysis of $A - A_e$.

6.9   The real (global) error at a point

The error in the amplification factor reflects the error when progressing from time level t_n to t_{n-1}. To investigate the real error at a point, known as the *global error*, we look at $u^n - u_e(t_n)$ for some n:

```
>>> n = Symbol('n')
>>> u_e = exp(-p*n)
>>> u_n = A**n
>>> FE = u_e.series(p, 4) - u_n.subs(theta, 0).series(p, 4)
>>> BE = u_e.series(p, 4) - u_n.subs(theta, 1).series(p, 4)
>>> CN = u_e.series(p, 4) - u_n.subs(theta, half).series(p, 4)
>>> FE
(1/2)*n*p**2 - 1/2*n**2*p**3 + (1/3)*n*p**3 + O(p**4)
>>> BE
(1/2)*n**2*p**3 - 1/2*n*p**2 + (1/3)*n*p**3 + O(p**4)
>>> CN
(1/12)*n*p**3 + O(p**4)
```

For a fixed time t, the parameter n in these expressions increases as $p \to 0$ since $t = n\Delta t$. That is, $n = t/\Delta t = at/\Delta t$, and the leading-order error terms therefore become $\frac{1}{2}at\Delta t$ for the Forward and Backward Euler scheme, and $\frac{1}{12}at\Delta t^2$ for the Crank-Nicolson scheme. The global error is therefore of second order in Δt for the latter scheme and first order for the former schemes.

6.10   Integrated errors

The formulas for various error measures have so far measured the error at one time point. Many prefer to use the error integrated over the whole time interval

of interest: $[0, T]$. An immediate practical problem arises, however, since the numerical solution is only known at the mesh points, while an integration will need this solution also at the points between the mesh points. Let $\tilde{u}$ be a continuous representation of the numerical solution, usually obtained by drawing straight lines between the values at the mesh points. Then a common measure of the global error is the so-called L^2 error:

$$E_2 = \sqrt{\int_0^T (u_{\mathrm{e}}(t) - \tilde{u}(t))^2 dt}\,. \tag{49}$$

A family of such measures is the L^p errors, defined as

$$E_p = \left(\int_0^T (u_{\mathrm{e}}(t) - \tilde{u}(t))^p dt \right)^{1/p}\,. \tag{50}$$

For $p = 1$ we just take the absolute value of the integrand.

Strictly speaking, it is questionable in a finite difference method to introduce an additional approximation in the error meausure, namely how $\tilde{u}$ varies between the mesh points. Some may argue and say that the numerical solution is defined at the mesh points only and that we should approximate the integrals above by numerical methods involving the integrand at just the mesh points. The numerical integration method also represents an approximation, but a discrete integration procedure is consistent with having only discrete values of the integrand.

For uniformly distributed mesh points we have the well-known Trapezoidal rule,

$$E_2 \approx \left(\Delta t \left(\frac{1}{2}(u_{\mathrm{e}}(0) - u^0)^2 + \frac{1}{2}(u_{\mathrm{e}}(T) - u^N)^2 + \sum_{k=1}^{N-1} (u_{\mathrm{e}}(k\Delta t) - u^k)^2 \right) \right)^{1/2}\,. \tag{51}$$

In case the mesh points are arbitrarily spaced, we have an immediate generalization in terms of the sum of the various trapezoids:

$$E_2 \approx \left(\frac{1}{2} \sum_{k=0}^{N-1} (t_{k+1} - t_k) \left((u_{\mathrm{e}}(k\Delta t) - u^k)^2 + (u_{\mathrm{e}}((k+1)\Delta t) - u^{k+1})^2 \right) \right)^{1/2}\,. \tag{52}$$

A simpler approximation is to use rectangles whose heights are determined by the left (or right) value in each interval:

$$E_2 \approx \left(\sum_{k=0}^{N-1} (t_{k+1} - t_k)(u_{\mathrm{e}}(k\Delta t) - u^k)^2 \right)^{1/2}\,. \tag{53}$$

With uniformly distributed mesh points we get the simplification

$$E_2 \approx \left(\Delta t \sum_{k=0}^{N-1} (u_{\mathrm{e}}(k\Delta t) - u^k)^2 \right)^{1/2}\,. \tag{54}$$

Suppose that in a program the u^k values are available as elements in the array u, while the $u_e(k\Delta t)$ values are available as elements in the array u_e. The formula (54) can then be calculated as follows by array arithmetics in Python:

```
E2 = sqrt(dt*sum((u_e - u)**2))
```

This is exactly the "array formula" that popped up in Section 2.3.

Integrated error measures sum up the contributions from each mesh point, so we must expect the global error to be larger than the local error. Roughly speaking, if $|u_e(t_n) - u^n| \sim p^r$, we have

$$E \approx \sqrt{\Delta t \sum_{i=0}^{N} p^{2r}} \approx a\Delta t^{r-1/2},$$

because $\sum_{i=0}^{N} p^{2r} = p^{2r} \sum_{i=0}^{N} 1 \sim p^{2r}\frac{1}{2}N^2$, $N = T/Deltat$, and $p = a\Delta t$.

7 Exercises

Exercise 11: Explore the θ-rule for exponential growth

Solve the ODE $u' = -au$ with $a < 0$ such that the ODE models exponential growth. Run experiments with θ and Δt using the dc_exper1.py code modified to your needs. Are there any numerical artifacts? Filename: growth_exper1.py.

Exercise 12: Summarize investigations in a report

Write a scientific report about the findings in Exercise 11. You can use examples from Section 4.2 to see how various formats can be used for scientific reports. Filename: growth_analysis.pdf.

Exercise 13: Plot amplification factors for exponential growth

Modify the dc_ampf_plot.py code to visualize the amplification factors for $\theta = 0, 0.5, 1$ and the exact amplification factor in case of exponential growth as in Exercise 11. Explain the artifacts seen in Exercise 11. Filename: growth_ampf_plot.py.

8 Model extensions

It is time to consider generalizations of the simple decay model $u = -au$, where a is constant, and also to look at other numerical solution methods.

8.1 Extension to a variable coefficient

In the ODE for decay, $u' = -au$, we now consider the case where a depends on time:

$$u'(t) = -a(t)u(t), \quad t \in (0, T], \quad u(0) = I. \tag{55}$$

A Forward Euler scheme consist of evaluating (55) at $t = t_n$ and approximating the derivative with a forward difference $[D_t^+ u]^n$:

$$\frac{u^{n+1} - u^n}{\Delta t} = -a(t_n)u^n. \tag{56}$$

The Backward Euler scheme becomes

$$\frac{u^n - u^{n-1}}{\Delta t} = -a(t_n)u^n. \tag{57}$$

The Crank-Nicolson method builds on sampling the ODE at $t_{n+\frac{1}{2}}$. We can evaluate a at $t_{n+\frac{1}{2}}$ and use an average for u at times t_n and t_{n+1}:

$$\frac{u^{n+1} - u^n}{\Delta t} = -a(t_{n+\frac{1}{2}})\frac{1}{2}(u^n + u^{n+1}). \tag{58}$$

Alternatively, we can use an average for the product au:

$$\frac{u^{n+1} - u^n}{\Delta t} = -\frac{1}{2}(a(t_n)u^n + a(t_{n+1})u^{n+1}). \tag{59}$$

The θ-rule unifies the three mentioned schemes,

$$\frac{u^{n+1} - u^n}{\Delta t} = -a((1-\theta)t_n + \theta t_{n+1})((1-\theta)u^n + \theta u^{n+1}). \tag{60}$$

or,

$$\frac{u^{n+1} - u^n}{\Delta t} = -(1-\theta)a(t_n)u^n - \theta a(t_{n+1})u^{n+1}. \tag{61}$$

With the finite difference operator notation the Forward Euler and Backward Euler schemes can be summarized as

$$[D_t^+ u = -au]^n, \tag{62}$$

$$[D_t^- u = -au]^n. \tag{63}$$

The Crank-Nicolson and θ schemes depend on whether we evaluate a at the sample point for the ODE or if we use an average. The various versions are written as

$$[D_t u = -a\overline{u}^t]^{n+\frac{1}{2}}, \tag{64}$$

$$[D_t u = -\overline{au}^t]^{n+\frac{1}{2}}, \tag{65}$$

$$[D_t u = -a\overline{u}^{t,\theta}]^{n+\frac{1}{2}}, \tag{66}$$

$$[D_t u = -\overline{au}^{t,\theta}]^{n+\frac{1}{2}}. \tag{67}$$

8.2 Extension to a source term

A further extension of the model ODE is to include a source term $b(t)$:

$$u'(t) = -a(t)u(t) + b(t), \quad t \in (0, T], \quad u(0) = I. \tag{68}$$

Schemes. The time point where we sample the ODE determines where $b(t)$ is evaluated. For the Crank-Nicolson scheme and the θ-rule we have a choice of whether to evaluate $a(t)$ and $b(t)$ at the correct point or use an average. The chosen strategy becomes particularly clear if we write up the schemes in the operator notation:

$$[D_t^+ u = -au + b]^n, \tag{69}$$

$$[D_t^- u = -au + b]^n, \tag{70}$$

$$[D_t u = -a\overline{u}^t + b]^{n+\frac{1}{2}}, \tag{71}$$

$$[D_t u = \overline{-au + b}^t]^{n+\frac{1}{2}}, \tag{72}$$

$$[D_t u = -a\overline{u}^{t,\theta} + b]^{n+\theta}, \tag{73}$$

$$[D_t u = \overline{-au + b}^{t,\theta}]^{n+\theta}. \tag{74}$$

Implementation. Writing out the latter θ-rule, using (32) and (33), we get

$$\frac{u^{n+1} - u^n}{\Delta t} = \theta(-a^{n+1}u^{n+1} + b^{n+1})) + (1 - \theta)(-a^n u^n + b^n)),$$

where a^n means evaluating a at $t = t_n$ and similar for a^{n+1}, b^n, and b^{n+1}. We solve for u^{n+1}:

$$u^{n+1} = ((1 - \Delta t(1 - \theta)a^n)u^n + \Delta t(\theta b^{n+1} + (1 - \theta)b^n))(1 + \Delta t\theta a^{n+1})^{-1}. \tag{75}$$

A suitable implementation where $a(t)$ and $b(t)$ are given as Python function is given next.

```python
def solver(I, a, b, T, dt, theta):
    """
    Solve u'=-a(t)*u + b(t), u(0)=I,
    for t in (0,T] with steps of dt.
    a and b are Python functions of t.
    """
    dt = float(dt)            # avoid integer division
    N = int(round(T/dt))      # no of time intervals
    T = N*dt                  # adjust T to fit time step dt
    u = zeros(N+1)            # array of u[n] values
    t = linspace(0, T, N+1)   # time mesh

    u[0] = I                  # assign initial condition
    for n in range(0, N):     # n=0,1,...,N-1
        u[n+1] = ((1 - dt*(1-theta)*a(t[n]))*u[n] + \
```

```
                 dt*(theta*b(t[n+1]) + (1-theta)*b(t[n])))/\
                 (1 + dt*theta*a(t[n+1]))
    return u, t
```

This function is found in the file `dc_vc.py`.

Verification.   One common method for verifying implementations of differential equation solvers is to choose $u_e(t)$ as a linear function, adjust variable coefficients to allow this solution, and then see if the numerical solution coincides with the linear solution at the mesh points. The rationale behind this test is that most numerical methods are capable of reproducing a linear solution (to machine precision).

Let $u_e(t) = ct + I$, which fulfills $u_e(0) = I$. We can choose any $a(t)$ and insert this $u_e(t)$ in the differential equation to get

$$c = -a(t)u_e(t) + b(t) \, .$$

Any function $u_e(t)$ is the correct solution if we choose

$$b(t) = c + a(t)u_e(t) \, .$$

A relevant implementation of the test is

```
def verify_linear_solution():
    def exact_solution(t):
        return c*t + I

    def a(t):
        return t**0.5  # can be arbitrary

    def b(t):
        return c + a(t)*exact_solution(t)

    theta = 0; I = 0.1; dt = 0.1; c = -0.5
    T = 4
    N = int(T/dt)  # no of steps
    u, t = solver(I=I, a=a, b=b, T=N*dt, dt=dt, theta=theta)
    u_e = array([exact_solution(tn) for tn in t])
    difference = abs(u_e - u).max()  # max deviation
    tol = 1E-15  # tolerance for comparing floats
    success = difference <= tol
    return success
```

8.3   Extension to systems of ODEs

Many ODE models involves more than one unknown function and more than one equation. Here is an example of two unknown functions $u(t)$ and $v(t)$ (modeling, e.g., the radioactive decay of two substances):

$$u' = -a_u u + a_v v, \tag{76}$$
$$v' = -a_v v + a_u u, \tag{77}$$

for constants $a_u, a_v > 0$. Applying the Forward Euler method to each equation results in simple updating formula

$$u^{n+1} = u^n + \Delta t(-a_u u^n + a_v v^n), \tag{78}$$

$$v^{n+1} = u^n + \Delta t(-a_v v^n + a_u u^n). \tag{79}$$

However, the Crank-Nicolson or Backward Euler schemes results in a 2×2 linear system for the new unknowns. The latter schemes gives

$$u^{n+1} = u^n + \Delta t(-a_u u^{n+1} + a_v v^{n+1}), \tag{80}$$

$$v^{n+1} = v^n + \Delta t(-a_v v^{n+1} + a_u u^{n+1}), . \tag{81}$$

and bringing u^{n+1} as well as v^{n+1} on the left-hand side results in

$$(1 + \Delta t a_u)u^{n+1} + a_v v^{n+1}) = u^n, \tag{82}$$

$$a_u u^{n+1} + (1 + \Delta t a_v)v^{n+1} = v^n, . \tag{83}$$

which is a system of two coupled, linear, algebraic equations in two unknowns.

9 General first-order ODEs

9.1 Generic form

ODEs are commonly written in a generic form

$$u' = f(u, t), \quad u(0) = I, \tag{84}$$

where $f(u, t)$ is a prescribed function. As an example, our most general exponential decay model (68) has $f(u, t) = -a(t)u(t) + b(t)$.

The unknown u in (84) may either be a scalar function of time t, or a vector valued function of t in case of a *system of ODEs*:

$$u(t) = (u^{(0)}(t), u^{(1)}(t), \ldots, u^{(m-1)}(t)).$$

In that case, the right-hand side is vector-valued function with m components,

$$
\begin{aligned}
f(u, t) = (&f^{(0)}(u^{(0)}(t), \ldots, u^{(m-1)}(t)), \\
&f^{(1)}(u^{(0)}(t), \ldots, u^{(m-1)}(t)), \\
&\vdots, \\
&f^{(m-1)}(u^{(0)}(t), \ldots, u^{(m-1)}(t))).
\end{aligned}
$$

Actually, any system of ODEs can be written in the form (84), but higher-order ODEs then need auxiliary unknown functions.

9.2 The Odespy software

A wide range of methods and software exist for solving (84). Many of methods are accessible through a unified Python interface offered by the Odespy package. Odespy features simple Python implementations of the most fundamental schemes as well as Python interfaces to several famous packages for solving ODEs: ODEPACK, Vode, rkc.f, "rkf45.f": http://www.netlib.org/ode/rkf45.f", "Radau5": "http://www.unige.ch/ hairer/software.html, as well as the ODE solvers in SciPy, SymPy, and odelab.

The usage of Odespy follows this setup for the ODE $u' = -au$ solved by the famous 4th-order Runge-Kutta method:

```
def f(u, t):
    return -a*u

import odespy
import numpy as np

I = 1; a = 2; T = 6; dt = 1
solver = odespy.RK(f)
solver.set_initial_condition(I)
t_mesh = np.linspace(0, T, N+1)
u, t = solver.solve(t_mesh)
```

9.3 Example: Runge-Kutta methods

Since all solvers have the same interface, modulo different set of parameters to the solvers' constructors, one can easily make a list of solver objects and run a loop for comparing (a lot of) solvers. The code below, found in complete form in `dc_odespy.py` compares the famous Runge-Kutta methods of orders 2, 3, and 4 with the Backward Euler scheme and the exact solution of the decay equation $u' = -au$. Figure 17 shows the results.

```
import numpy as np
import scitools.std as plt
import sys

def f(u, t):
    return -a*u

I = 1; a = 2; T = 6
dt = float(sys.argv[1]) if len(sys.argv) >= 2 else 0.75
N = int(round(T/dt))
t = np.linspace(0, N*dt, N+1)

solvers = [odespy.RK2(f),
           odespy.RK3(f),
           odespy.RK4(f),
           odespy.BackwardEuler(f, nonlinear_solver='Newton')]

legends = []
for solver in solvers:
    solver.set_initial_condition(I)
    u, t = solver.solve(t)
```

```
    plt.plot(t, u)
    plt.hold('on')
    legends.append(solver.__class__.__name__)

# Compare with exact solution plotted on a very fine mesh
t_fine = np.linspace(0, T, 10001)
u_e = I*np.exp(-a*t_fine)
plt.plot(t_fine, u_e, '-') # avoid markers by specifying line type
legends.append('exact')

plt.legend(legends)
plt.title('Time step: %g' % dt)
plt.show()
```

This code might deserve a couple of comments. We use SciTools for plotting, because with the Matplotlib and Gnuplot backends, curves are automatically given colors *and* markers, the latter being important when PNG plots are printed in reports in black and white. (The default Matplotlib and Gnuplot behavior gives colored lines, which are difficult to distinguish. However, Gnuplot automatically introduces different line styles if output in the Encapsulated PostScript format is specified via `savefig`). The automatic adding of markers is not suitable for a very finely resolved line, like the one for `u_e` in this case, and then we specify the line type as a solid line (`-`), leaving the color up to the backend used for plotting. The legends are based on the class names of the solvers, and in Python the name of a the class type (as a string) of an object `obj` is obtained by `obj.__class__.__name__`.
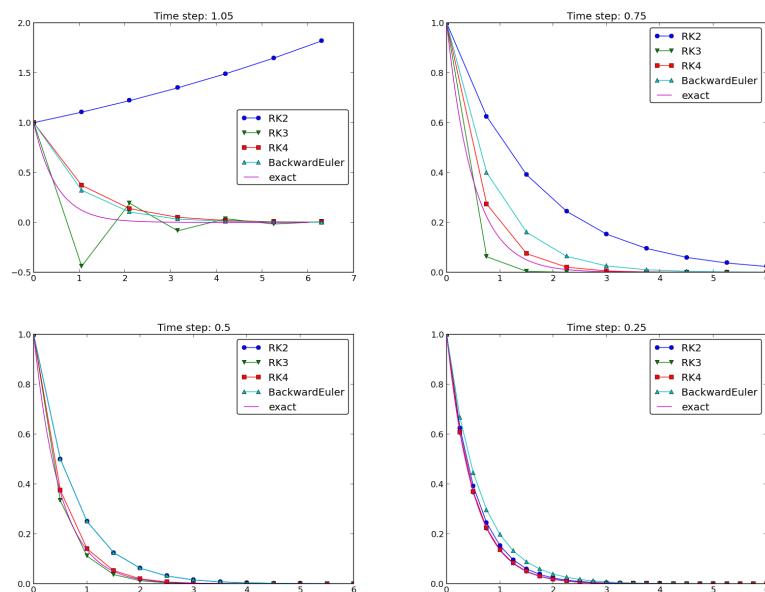

Figure 17:   Behavior of different schemes for the decay equation.

80

Since we have quite long time steps, we have included the only relevant θ-rule for large time steps, the BackwardEuler scheme ($\theta = 1$), as well in Figure 17. These and other experiments reveal that

- The 2-nd order Runge-Kutta method (`RK2`) is unstable for $\Delta t > 1$ and decays slower than the Backward Euler scheme, which is known to decay slower than the Crank-Nicolson and Forward Euler schemes. However, for fine $\Delta t = 0.25$ the 2-nd order Runge-Kutta method decays faster than the Backward Euler scheme and lies closer to the exact curve. This means that the 2-nd order Runge-Kutta method approaches the exact solution faster than the Backward Euler scheme, although it is less accurate for large Δt values.

- The 3-rd order Runge-Kutta method (`RK3`) is oscillating for $\Delta t = 1.05$ and lies partly above and below the exact curve for $\Delta t = 0.75$. This behavior is much like what one observes for the Crank-Nicolson scheme. For finer Δt, the 3-rd order Runge-Kutta method converges quickly to the exact solution.

- The 4-th order Runge-Kutta method (`RK4`) is slightly inferior to the Backward Euler scheme on the coarsest mesh, but is then clearly superior to all the other schemes. It is definitely the method of choise for all the tested schemes.

Remark about using the θ-rule in Odespy. The Odespy package assumes that the ODE is written as $u' = f(u, t)$ with an f that is possibly nonlinear in u. The θ-rule for $u' = f(u, t)$ leads to

$$u^{n+1} = u^n + \Delta t \left(\theta f(u^{n+1}, t_{n+1}) + (1 - \theta) f(u^n, t_n) \right),$$

which is a *nonlinear equation* in u^{n+1}. Odespy's implementation of the θ-rule (`ThetaRule`) and the specialized Backward Euler and Crank-Nicolson schemes (called `BackwardEuler` and `MidpointImplicit`, respectively) must invoke iterative methods for solving the nonlinear equation in u^{n+1}. This is done even when f is linear in u, as in the model problem $u' = -au$, where we can easily solve for u^{n+1}. In the present example, we need to use Newton's method to ensure that Odespy is capable of solving the equation for u^{n+1} for large Δt. (The reason is that the Forward Euler method is used to compute the initial guess for the nonlinear iteration method, and this Forward Euler may give very wrong values for large Δt. The Newton method is not sensitive to a bad initial guess in linear problems.)

9.4 Example: Adaptive Runge-Kutta methods

Odespy offers solution methods that can adapt the size of Δt with time to match a desired accuracy in solution. Intuitively, small time steps will be chosen in

areas where the solution is changing rapidly, while larger time steps can be used where the solution is slowly varying. Some kind of *error estimater* is used to adjust the next time step at each time level.

A very popular adaptive method for solving ODEs is the Dormand-Prince Runge-Kutta method of order 4 and 5. The 5th-order method is used as a reference solution and the difference between the 4th- and 5th-order methods is used as an indicator of the error in the numerical solution. The Dormand-Prince method is the default choice in MATLAB's famous `ode45` routine.

We can easily set up Odespy to use the Dormand-Prince method and see how it selects the optimal time steps. To this end, we request only one time step from $t = 0$ to $t = T$ and ask the method to compute the necessary non-uniform time mesh to meet a certain error tolerance. The code goes like

```
import odespy
import numpy as np
import dc_mod
import sys
#import matplotlib.pyplot as plt
import scitools.std as plt

def f(u, t):
    return -a*u

def exact_solution(t):
    return I*np.exp(-a*t)

I = 1; a = 2; T = 5
tol = float(sys.argv[1])
solver = odespy.DormandPrince(f, atol=tol, rtol=0.1*tol)

N = 1  # just one step - let the scheme find its intermediate points
t_mesh = np.linspace(0, T, N+1)
t_fine = np.linspace(0, T, 10001)

solver.set_initial_condition(I)
u, t = solver.solve(t_mesh)

# u and t will only consist of [I, u^N] and [0,T]
# solver.u_all and solver.t_all contains all computed points
plt.plot(solver.t_all, solver.u_all, 'ko')
plt.hold('on')
plt.plot(t_fine, exact_solution(t_fine), 'b-')
plt.legend(['tol=%.0E' % tol, 'exact'])
plt.savefig('tmp_odespy_adaptive.png')
plt.show()
```

Running four cases with tolerances 10^{-1}, 10^{-3}, 10^{-5}, and 10^{-7}, gives the results in Figure 18. Intuitively, one would expect denser points in the beginning of the decay and larger time steps when the solution flattens out.

9.5   Other schemes

Next we list some well-known methods for $u' = f(u, t)$, valid both for a single ODE (scalar u) and systems of ODEs (vector u).
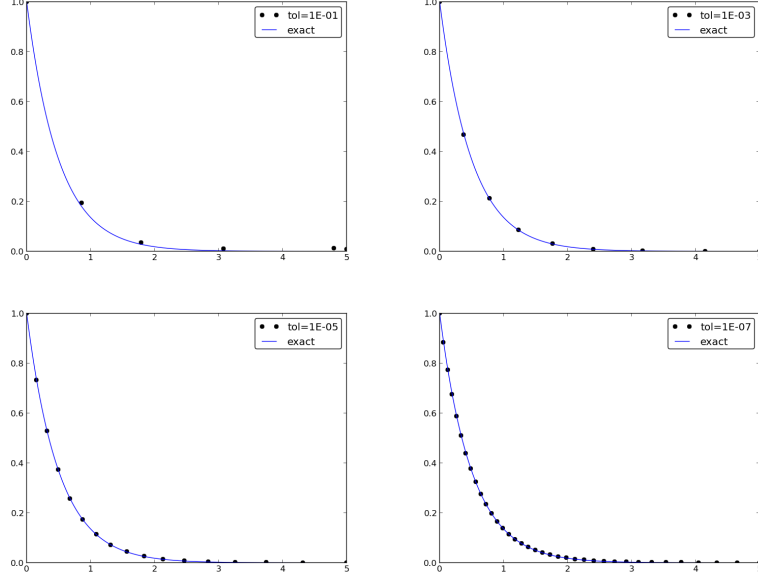
Figure 18: Choice of adaptive time mesh by the Dormand-Prince method for difference tolerances.

Implicit 2-step backward scheme. The implicit backward method with 2 steps applies a three-level backward difference as approximation to $u'(t)$,

$$u'(t_{n+1}) \approx \frac{3u^{n+1} - 4u^n + u^{n-1}}{2\Delta t},$$

which is an approximation of order Δt^2 to the first derivative. The resulting scheme for $u' = f(u,t)$ reads

$$u^{n+1} = \frac{4}{3}u^n - \frac{1}{3}u^{n-1} + \frac{2}{3}\Delta t f(u^{n+1}, t_{n+1}).$$

The scheme is implicit and requires solution of nonlinear equations when f is nonlinear in u. The standard 1st-order Backward Euler method or the Crank-Nicolson scheme can be used for the first step.

The Leapfrog scheme. The derivative of u at some point t_n can be approximated by a centeral difference over two time steps,

$$u'(t_n) \approx \frac{u^{n+1} - u^{n-1}}{2\Delta t}, \tag{85}$$

which is an approximation of second order in Δt. This approximation gives the scheme

$$u^{n+1} = u^{n-1} + \Delta t f(u^n, t_n). \tag{86}$$

Some other scheme must be used as starter to compute u^1. Observe that (86) is an explicit scheme, and that a nonlinear f (in u) is trivial to handle.

2nd-order Runge-Kutta scheme. The two-step scheme

$$u^* = u^n + \Delta t f(u^n, t_n), \tag{87}$$

$$u^{n+1} = u^n + \Delta t \frac{1}{2} \left(f(u^n, t_n) + f(u^*, t_{n+1}) \right), \tag{88}$$

essentially applies a Crank-Nicolson method to the ODE, but replaces the term $f(u^{n+1}, t_{n+1})$ by a prediction $f(u^*, t_{n+1})$ based on a Forward Euler step. The scheme (87)-(88) is known as Huen's method, but is also a 2nd-order Runge-Kutta method. The scheme is explicit, and the error is expected to behave as Δt^2.

2nd-order Adams-Bashforth scheme. The following method is known as the 2nd-order Adams-Bashforth scheme:

$$u^{n+1} = u^n + \frac{1}{2} \Delta t \left(3f(u^n, t_n) - f(u^{n-1}, t_{n-1}) \right). \tag{89}$$

The scheme is explicit and requires another one-step scheme to compute u^1 (the Forward Euler scheme or Heun's method, for instance). As the name implies, the scheme is of order Δt^2.

3rd-order Adams-Bashforth scheme. Another explicit scheme, involving four time levels, is the 3rd-order Adams-Bashforth scheme

$$u^{n+1} = u^n + \frac{1}{12} \left(23f(u^n, t_n) - 16f(u^{n-1}, t_{n-1}) + 5f(u^{n-2}, t_{n-2}) \right). \tag{90}$$

The numerical error is of order Δt^3, and the scheme needs some method for computing u^1 and u^2.

10 Exercises

Exercise 14: Implement the 2-step backward scheme

Implement the 2-step backward method (9.5) for the model $u'(t) = -a(t)u(t) + b(t)$, $u(0) = I$. Allow the first step to be computed by either the Backward Euler scheme or the Crank-Nicolson scheme.

Run a test case where a constant and $b = 0$ and determine if the choice of a first-order scheme (Backward Euler) has any impact on the overall accuracy of this scheme, which is expected to be of second order in Δt. Filename: `dc_backward2step.py`.

Exercise 15: Implement the Leapfrog scheme

Implement the Leapfrog scheme (86) for the model $u'(t) = -a(t)u(t) + b(t)$, $u(0) = I$. Since the Leapfrog scheme is explicit in time, it is most convenient to use the explicit Forward Euler scheme for computing u^1.

Determine if the choice of a first-order scheme (Forward Euler) for the first time step has any impact on the overall accuracy of the Leapfrog scheme, which is expected to be of second order in Δt (see Exercise 15). Filename: `dc_leapfrog.py`.

Exercise 16: Experiment with the Leapfrog scheme

Set up a set of experiments to demonstrate that the Leapfrog scheme (86) as implemented in Exercise 15 is associated with numerical artifacts (instabilities). Filename: `dc_leapfrog_exper.py`.

Exercise 17: Analyze the Leapfrog scheme

The purpose of this exercise is to analyze and explain instabilities of the Leapfrog scheme (86). Consider the case where a is constant and $b = 0$. Assume that an exact solution of the discrete equations has the form $u^n = A^n$, where A is an amplification factor to be determined. Use `sympy` to compute A. Since the governing polynomial for A has two roots, A_1 and A_2, u^n is a linear combination $u^n = C_1 A_1^n + C_2 A_2^n$. Filename: `dc_leapfrog_anaysis.py`.

Exercise 18: Implement the 2nd-order Adams-Bashforth scheme

Implement the 2nd-order Adams-Bashforth method (89) for the decay problem $u' = -a(t)u + b(t)$, $u(0) = I$, $t \in (0, T]$. Use the Forward Euler method for the first step such that the overall scheme is explicit. Filename: `dc_AdamBashforth2.py`.

Exercise 19: Implement the 3rd-order Adams-Bashforth scheme

Implement the 3rd-order Adams-Bashforth method (90) for the decay problem $u' = -a(t)u + b(t)$, $u(0) = I$, $t \in (0, T]$. Since the scheme is explicit, allow it to be started by two steps with the Forward Euler method. Investigate experimentally the case where $b = 0$ and a is a constant: Can we have oscillatory solutions for large Δt? Filename: `dc_AdamBashforth3.py`.

Exercise 20: Generalize a class implementation

Consider the file `dc_class.py` where the exponential decay problem $u' = -au$, $u(0) = I$, is implemented via three classes `Problem`, `Solver`, and `Visualizer`.

Extend the classes to handle the more general problem

$$u'(t) = -a(t)u(t) + b(t), \quad u(0) = I, \ t \in (0, T],$$

using the θ-rule for discretization.

In the case with arbitrary functions $a(t)$ and $b(t)$ the problem class is no longer guaranteed to provide an exact solution. Let the `exact_solution` in class `Problem` return `None` if the exact solution for the particular problem is not available. Modify classes `Solver` and `Visualizer` accordingly.

Add test functions `test_*()` for the nose testing tool in the module. Also add a demo example for a rough model of a parachute jumper where $b(t) = g$ (acceleration of gravity) and

$$a(t) = \left\{ \begin{array}{ll} a_0, & 0 \le t \le t_p, \\ ka_0, & t > t_p, \end{array} \right.$$

where t_p is the point of time the parachute is released and k is the factor of increased air resistance. Scale the model using the terminal velocity of the free fall, $v = \sqrt{g/a_0}$, as velocity scale. Make a demo of the scaled model with $k = 50$. Filename: `dc_class2.py`.

Exercise 21: Generalize a advanced class implementation

Solve Exercise 20 by utilizing the class implementations in `dc_class_oo.py`. Filename: `dc_class3.py`.

Exercise 22: Make a unified implementation of many schemes

Consider the linear ODE problem $u'(t) = -a(t)u(t) + b(t)$, $u(0) = I$. Many solution schemes for this problem can be written in the (explicit) form

$$u^{n+1} = \sum_{j=0}^{m} c_j u^{n-j}, \tag{91}$$

for some choice of $c_0, \ldots, c_m$. Find thhe c_j coefficients for the θ-rule, the three-level backward scheme, the Leapfrog scheme, the 2nd-order Runge-Kutta method, and the 3rd-order Adams-Bashforth scheme.

Make a class `ExpDecay` that implements (91), with subclasses specifying lists $c_0, \ldots, c_m$ for the various methods. The subclasses also need extra lists for methods needed to start schemes with $m > 0$. Verify the implementation by testing with a linear solution $(u_e(t) = ct + d)$.

test

Filename: `decay_schemes_oo.py`.

11 Applications of exponential decay models

This section presents derives a set of mathematical models that all end up with ODEs of the type $u' = -au + b$, where the solution exhibits exponential growth or decay. The applications are taken from biology, finance, and physics, and cover population growth or decay, compound interest and inflation, radioactive decay, cooling of bodies, pressure variations in the atmosphere, and air resistance on falling or rising bodies.

11.1 Evolution of a population

Let N be the number of individuals in a population occupying some spatial domain. Although N is then an integer we compute with N as a real number. Assuming that in a time interval Δt the number of newcomers to the populations (newborns) is proportional to N, with proportionality constant $\bar{b}$, the population is increased to

$$N(t + \Delta t) = N(t) + \bar{b}N(t).$$

It is obvious that a long time interval Δt will result in more newcomers and hence a larger $\bar{b}$. Therefore, we introduce $b = \bar{b}/\Delta t$: the number of newcomers per unit time and per individual. We must then multiply b by the length of the time interval considered and by the population size to get the total number of new individuals, $b\Delta tN$.

If the number of removals from the population (deaths) is also proportional to N, with proportionalty constant $d\Delta t$, the population evolves according to

$$N(t + \Delta t) = N(t) + b\Delta tN(t) - d\Delta tN(t).$$

Dividing by Δt and letting $\Delta t \to 0$, we get the ODE

$$N' = (b - d)N. \tag{92}$$

This ODE is of the form $u' = -au$ with $a = d - b$ and $u = N$. In a population where the death rate (d) is larger than then newborn rate (b), $a > 0$, and the population experiences exponential decay.

In some populations there is an immigration of individuals into the spatial domain. With I individuals coming in per time unit, the equation for the population change becomes

$$N(t + \Delta t) = N(t) + b\Delta tN(t) - d\Delta tN(t) + \Delta tI.$$

The corresponding ODE reads

$$N' = (b - d)N + I. \tag{93}$$

This ODE can be written in the general form $u' = -au + b$ if we define a as $d - b$ and b as I. Note that the symbol b represents two different quantities (the source term in the general model $u' = -au + b$ and the birth rate in the population model).

This simple exponential model for populations has many applications, including animal populations and cell cultures. The parameter a may depend on time if the environment changes (independent of u). For example, if the food or nutrition supply suddenly drops an amount A at $t = s$, we can introduce

$$a(t) = \begin{cases} a_0, & t > s, \\ a_0 - A, & t \geq s, \end{cases}$$

Similarly, b in the model $u' = -au + b$ may depend on time if the influx b of new individuals changes with time.

The growth rate a of a population decreases if the environment has limited. Suppose the environment can sustain M individuals. Taking $a(t)$ as a linear function, $a(t) = a_0(1 - u/M)$, where a_0 is the initial growth rate, results in the equation

$$u' = a_0(1 - u/M)u, \tag{94}$$

known as the *logistic equation*. We realize that the growth rate reaches zero when $u = M$, which means that u cannot increase further and hence $u(0) \leq u \leq M$.

11.2 Compound interest and inflation

Say the annual interest rate is r percent and that the bank adds the interest once a year to your investment. If u^n is the investment in year n, the investment in year u^{n+1} grows to

$$u^{n+1} = u^n + \frac{r}{100}u^n.$$

In reality, the interest rate is added every day. We therefore introduce a parameter m for the number of periods per year where the interest is added to the investment. If n counts the periods, we have

$$u^{n+1} = u^n + \frac{r}{100m}u^n.$$

Starting with an investment u^0, and assuming that r is constant in time, we get

$$\begin{aligned} u^{n+1} &= \left(1\frac{r}{100m}\right)u^n \\ &= \left(1\frac{r}{100m}\right)^2 u^{n-1} \\ &\;\;\vdots \\ &= \left(1\frac{r}{100m}\right)^{n+1} u^0 \end{aligned}$$

Introducing a real counter for years, t, we have that $n = mt$, so we can write

$$u^{mt} = \left(1\frac{r}{100m}\right)^{mt} u^0.$$

Continuous compounding means that the interest is added continuously, or $m \to \infty$. The result is

$$u(t) = u(0)e^{rt/100} \, . \tag{95}$$

This means that u is governed by a differential equation

$$u' = \frac{r}{100}u \, . \tag{96}$$

For $r > 0$ this equation models continuous compound interest, while for $r < 0$ the equation models continuous inflation, with r as the annual inflation rate measured in percent.

Introducing $a = \frac{r}{100}$, continuous inflation of an initial fortune I is then a process exhibiting exponential decay according to

$$u' = -au, \quad u(0) = I \, .$$

11.3   Radioactive Decay

An atomic nucleus of an unstable atom may lose energy by emitting ionizing particles and thereby be transformed to a nucleus with a different number of protons and neutrons. This process is called radioactive decay. Actually, the process is stochastic when viewed for a single atom, because it is impossible to predict exactly when a particular atom emit a particle. Nevertheless, with a large number of atoms, N, one may view the process as deterministic and compute the mean behavior of the decay.

Suppose at time t, the number of the original atom type is $N(t)$. A basic model assumption is that the transformation of the atoms of the original type in a small time interval Δt is proportional to N, so that

$$N(t + \Delta t) = N(t) - a\Delta t N(t),$$

where a is a constant. Introducing $u = N(t)/N(0)$, dividing by Δt and letting $\Delta t \to 0$ gives the following ODE:

$$u' = -au, \quad u(0) = 1 \, . \tag{97}$$

The parameter a can for a given nucleus be expressed through the *half-life* $t_{1/2}$, which is the time taken for the decay to reduce the initial amount by one half, i.e., $u(t_{1/2}) = 0.5$. With $u(t) = e^{-at}$, we get $t_{1/2} = a^{-1}\ln 2$ or $a = \ln 2/t_{1/2}$.

11.4   Newton's law of cooling

When a body at some temperature is placed in a cooling environment, experience shows that the temperature falls rapidly in the beginning, and then the changes in temperature levels off until the body's temperature equals that of the surroundings. Newton carried out some experiments on cooling hot iron and found that the temperature evolved as a "geometric progression at times in arithmetic progression", meaning that the temperature decayed exponentially.

Later, this result was formulated as a differential equation: the rate of change of temperature in a body is proportional to to the temperature difference between the body and its surroundings. This statement is known as *Newton's law of cooling*, which can be mathematically expressed as

$$\frac{dT}{dt} = -k(T - T_s),$$ (98)

where T is the temperature of the body, T_s is the temperature of the surroundings, t is time, and k is a positive constant. Equation (98) is primarily viewed as an empirical law, valid when heat is efficiently convected away from the surface of the body by a flowing fluid such as air at constant temperature T_s. The constant k reflects the transfer of heat from the body to the surroundings and must be determined from physical experiments.

   To bring (98) on the form of a standard decay model $u' = -au$, we introduce $u = T - T_s$ and assume that T_s is constant in time. Inserting $T = u + T_s$ in (98) reveals that $u' = -au$ with $a = k$.

11.5   Decay of atmospheric pressure with altitude

Vertical equilibrium of air in the atmosphere is governed by the equation

$$\frac{dp}{dz} = -\varrho g$$ (99)

. Here, $p(z)$ is the air pressure, ϱ is the density of air, and $g = 9.807$ m/s^2 is a standard value of the acceleration of gravity. (Equation (11.5) follows directly from the general Navier-Stokes equations under the assumption that the air does not move.)

   The pressure is related to density and temperature through the ideal gas law

$$\varrho = \frac{Mp}{R^*T},$$ (100)

where M is the molar mass of Earth's air (0.029 kg/mol), R^* is the universal gas constant (8.314 Nm/(mol K)), and T is the temperature. All variables p, ϱ, and T vary with the height z. Inserting (100) in (11.5) results in an ODE for exponential decay with a variable coefficient:

$$\frac{dp}{dz} = -\frac{Mg}{R^*T(z)}p$$ (101)

Multiple atmospheric layers.   The atmosphere can be modeled by seven layers where (101) T can be taken as linear in each layer:

$$T(z) = \bar{T}_i + L(z)(z - h_i),$$

where $z = h_i$ denotes the bottom of the layer number i, having temperature $\bar{T}_i$, and L_i is a constant in layer number i. The table below lists h_i (m), $\bar{T}_i$ (K), and L_i (K/m) for the layers $i = 0, \dots, 6$.

i	h_i	T_i	L_i
0	0	288	-0.0065
1	11,000	216	0.0
2	20,000	216	0.001
3	32,000	228	0.0028
4	47,000	270	0.0
5	51,000	270	-0.0028
6	71,000	214	-0.002

For implementation it might be convenient to write (101) on the form

$$\frac{dp}{dz} = -\frac{Mg}{R^*(\bar{T}(z) + L(z)(z - h(z)))} p, \tag{102}$$

where $\bar{T}(z)$, $L(z)$, and $h(z)$ are piecewise constant functions with values given in the table. The value of the pressure at the ground $z = 0$ is 101325 Pa.

Simplification: $L = 0$. One commonly used simplification is to assume that the temperature is constant within eacy layer. This means that $L = 0$.

Simplification: one-layer model. Another commonly used approximation is to work with one layer instead of seven. This one-layer model is based on $T(z) = T_0 - Lz$, with sea level standard temperature $T_0 = 288$ K and temperature lapse rate $L = 0.0065$ K/m. Furthermore, p_0 is the sea level atmospheric pressure $1 \cdot 10^5$ Pa.

11.6   Vertical motion of a body in a viscous fluid

A body moving vertically through a fluid (liquid or gas) is subject to three different types of forces: the gravity force, the drag force, and the buoyancy force.

Overview of forces. The gravity force is $F_g = -mg$, where m is the mass of the body and g is the acceleration of gravity. The uplift or buoyancy force ("Archimedes force") $F_b = \varrho g V$, where ϱ is the density of the fluid and V is the volume of the body.

The drag force is of two types, depending on the Reynolds number Re $= dv\varrho/\mu$, where d is the diameter of the body in the direction perpendicular to the flow, v is the velocity of the body, and μ is the dynamic viscosity of the fluid. When Re $\ll 1$, the drag force is the so-called Stokes' drag, which for a spherical body of diameter d reads

$$F_d^{(S)} = 3\pi d\mu v. \tag{103}$$

For large Re,

$$F_d^{(q)} = -\frac{1}{2} C_D \varrho A |v| v, \tag{104}$$

where C_D is a dimensionless drag coefficient depending on the body's shape, A is the cross-sectional area (produced by a cutting plane $y = \text{const}$ through the thickest part of the body). The superscripts q and S in $F_d^{(S)}$ and $F_d^{(q)}$ indicates Stokes drag and quadratic drag, respectively.

Equation of motion. All the mentioned forces act in the vertical direction. Newton's second law of motion applied to the body says that the sum of these forces must equal the mass of the body times its acceleration a in the vertical direction.

$$ma = F_g + F_d^{(S)} + F_b,$$

if we choose to work with the Stokes drag. Inserting the expressions for the forces yields

$$ma = -mg - 3\pi d\mu v + \varrho g V\,.$$

The unknowns here are v and a, i.e., we have two unknowns but only one equation. From kinematics in physics we know that the acceleration is the time derivative of the velocity: $a = dv/dt$. This is our second equation. We can easily eliminate a and get a single differential equation for v:

$$m\frac{dv}{dt} = -mg - 3\pi d\mu v + \varrho g V\,.$$

A small rewrite of this equation is handy: We express m as $\varrho_b V$, where ϱ_b is the density of the body, and we divide by the mass to get

$$v'(t) = -\frac{3\pi d\mu}{\varrho_b V}v + g\left(\frac{\varrho}{\varrho_b} - 1\right)\,. \tag{105}$$

We may introduce the constants

$$a = \frac{3\pi d\mu}{\varrho_b V}, \quad b = g\left(\frac{\varrho}{\varrho_b} - 1\right), \tag{106}$$

so that the structure of the differential equation becomes evident:

$$v'(t) = -av(t) + b\,. \tag{107}$$

The corresponding initial condition is $v(0) = v_0$ for some prescribed initial velocity v_0.

This derivation can be repeated with the quadratic drag force $F_d^{(q)}$, yielding

$$v'(t) = \frac{1}{2}C_D\frac{\varrho A}{\varrho_b V}| - v|v + g\left(\frac{\varrho}{\varrho_b} - 1\right)\,. \tag{108}$$

Defining

$$a = \frac{1}{2}C_D\frac{\varrho A}{\varrho_b V}, \tag{109}$$

and b as above, we can write (108)

$$v'(t) = -a|v|v + b. \tag{110}$$

Terminal velocity. An interesting aspect of (107) and (110) is whether we can approach a constant, so-called *terminal velocity* v_T, as $t \to \infty$. This means that $v'(t) \to 0$ as $t \to \infty$ and

$$0 = -av_T + b$$

and

$$0 = -a|v_T|v_T + b.$$

In the former case, $v_T = b/a$, while in the latter case, assuming a falling body with $v_T < 0$, $v_T = -\sqrt{|b|/a}$.

A Crank-Nicolson scheme. Both governing equations, the Stokes' drag model (107) and the quadratic drag model (110), can be readily solved by the Forward Euler scheme. The Crank-Nicolson method gives a nonlinear equation in v when applied to (110):

$$\frac{v^{n+1} - v^n}{\Delta t} = -a^{n+1/2}\frac{1}{2}(|v^{n+1}|v^{n+1} + a^n|v^n|v^n) + b^{n+1/2}. \tag{111}$$

However, instead of approximating the term $-a|v|v$ by an arithmetic average, we can use a *geometric average*:

$$|v|v^{n+1/2} \approx |v|^n v^{n+1}. \tag{112}$$

The error is of second order in Δt, just as for the arithmetic average. Now,

$$\frac{v^{n+1} - v^n}{\Delta t} = -a^{n+1/2}|v^n|v^{n+1} + b^{n+1/2}.$$

resulting in the scheme

$$v^{n+1} = v_n + \Delta t(1 + \Delta t a^{n+1/2}|v^n|)^{-1}b^{n+1/2} \tag{113}$$

Physical data. Suitable values of μ are $1.8{\cdot}10^{-5}$ Pa s for air and $8.9{\cdot}10^{-4}$ Pa s for water. Densities can be taken as 1.2 kg/m^3 for air and as $1.0 \cdot 10^3$ kg/m^3 for water. For considerable vertical displacement in the atmosphere one should take into account that the density of air varies with height, see Section 11.5. One possible density variation arises from the one-layer model in Section 11.5.

The drag coefficient C_D depends heavily on the shape of the body. Some values are 0.45 for a sphere, 0.42 for a semi-sphere, 1.05 for a cube, 0.82 for a long cylinder (with the length along the vertical direction), 0.75 for a rocket, 1.0-1.3 for a man in upright position, 1.3 for a flat plat perpendicular to the flow, and 0.04 for a streamlined (droplet-like) body.

Verification. To verify the program, one may assume a heavy body in air such that the F_b force can be neglected, and assume a small velocity such that the air resistance F_d can also be neglected. Setting $\varrho = 0$ removes both these terms from the equation. The motion then leads to the velocity $v(t) = v_0 - gt$, which is linear in t and therefore should be reproduced to machine precision (say tolerance 10^{-15}) by an implementation.

Another verification, but not as powerful as the one above, can be based on computing the terminal velocity and comparing with the exact expressions.

12   Exercises

Exercise 23: Simulate the pressure drop in the atmosphere

We consider the models for atmospheric pressure in Section 11.5. Make a program with three functions,

- one computing the pressure $p(z)$ using a seven-layer model and varying L,

- one computing $p(z)$ using a seven-layer model, but with constant temperature in each layer, and

- one computing $p(z)$ based on the one-layer model.

How can these implementations be verified? Should ease of verification impact how you code the functions? Compare the three models. Filename: `atmospheric_pressure.py`.

Exercise 24: Make a program for vertical motion in a fluid

Implement the Stokes' drag model and the quadratic drag model from Section 11.6 in a program, using the Crank-Nicolson scheme as explained. At each time level, compute the Reynolds number Re and choose the Stokes' drag model if Re $= \epsilon$, where typically $\epsilon = 0.1$. Otherwise, use the quadratic drag model. Make a nose test that runs the two suggestion verification tests in Section 11.6. Filename: `vertical_motion.py`.

Exercise 25: Plot forces acting in vertical motion in a fluid

Extend the program from Exercise 24 to compute the forces $F_d^{(q)}$, $F_d^{(S)}$, F_g, and F_b. Plot the latter two and the relevant drag force, depending on the value of the Reynolds number. Filename: `vertical_motion_force_plot.py`.

Exercise 26: Simulate a free fall of a parachute jumper

Apply the program from Exercise 24 to a parachute jumber in free fall before the parachute opens. Set the density of the human body as $\varrho_b = 1003 \text{ kg/m}^3$ and the mass as $m = 80$ kg, implying $V = m/\varrho_b = 0.08 \text{ m}^3$. One can base

the cross-sectional area A on the height 1.8 m and a width of 50 cm, giving giving $A \approx \pi R^2 = 0.9$ m^2. The density of air decreases with height, and here one may use the value 0.79 kg/m^3 which is relevant for about 5000 m height. The C_D coefficient can be set as 1.0. Start with $v_0 = 0$. Filename: `parachute_jumper1.py`.

Exercise 27: Simulate a parachute jumper

A parachute jumper moves a vertical distance of about 5000 m where A and C_D change dramatically when the parachute is released. Moreover, the density of air changes with altitude. Add a differential equation for the altitude z,

$$z'(t) = v(t), \tag{114}$$

and a differential equation for the air density, based on the information for the one-layer atmospheric model in Section 11.5:

$$p'(z) = -\frac{Mg}{R^*(T_0 + Lz)}p, \tag{115}$$

$$\varrho = p\frac{M}{RT^*}. \tag{116}$$

Extend the program from Exercise 24 to handle three coupled differential equations. With the Crank-Nicolson method, one must solve a 3×3 system of equations at each time level, since p depends on z, which depends on v, which depends on ϱ, which depends on p. Switching to the Forward Euler method avoids the need to solve coupled equations (but demands smaller Δt for accuracy). Make sure A and C_D can be changed at some time t_p, e.g., the time it takes to reach the terminal velocity in free fall. Make nose tests of this program based on the same verification cases as outlined in Exercise 24.

Use the program to simulate a jump from $z = 5000$ m to the ground $z = 0$. Filename: `parachute_jumper2.py`.

Exercise 28: Simulate a rising ball in water

A ball with the size of a soccer ball is placed in somewhat deep water, and we seek to model its motion upwards. The buoyancy force F_b is now the driving force, and the gravity force F_g is small. Set $A = \pi a^2$ with $a = 11$ cm, the mass of the ball is 0.43 kg, the density of water is 1000 kg/m^3, and C_D is 0.2. Start with $v_0 = 0$ and see how the ball rises.

Filename: `rising_ball.py`.

Exercise 29: Radioactive decay of Carbon-14

The Carbon-14 isotope, whose radioactive decay is used extensively in dating organic material that is tens of thousands of years old, has a half-life of 5, 730

years. Determine the age of an organic material that contains 8.4% of its initial amount of Carbon-14. Use a time unit of 1 year in the computations. The uncertainty in the half time of Carbon-14 is ± 40 years. What is the corresponding uncertainty in the estimate of the age? (Use simulations, not the exact solution.) Filename: `carbon14.py`.

Exercise 30: Radioactive decay of Carbon-14

Consider two radioactive substances A and B. The nuclei in substance A decay to form nuclei of type B with a half-life $A_{1/2}$, while substance B decay to form type A nuclei with a half-life $B_{1/2}$. Letting u_A and u_B be the fractions of the initial amount of material in substance A and B, respectively, the following system of ODEs governs the evolution of $u_A(t)$ and $u_B(t)$:

$$u'_A = \ln 2 u_B / B_{1/2} - \ln 2 u_A / A_{1/2}, \qquad\qquad (117)$$

$$u'_B = \ln 2 u_A / A_{1/2} - \ln 2 u_B / B_{1/2}, \qquad\qquad (118)$$

with $u_A(0) = u_B(0) = 1$.

Make a simulation program that solves for $u_A(t)$ and $u_B(t)$. Verify the implementation by computing analytically the limiting values of u_A and u_B as $t \to \infty$ (assume $u'_A, u'_B \to 0$) and comparing these with those obtained numerically.

Run the program for the case of $A_{1/2} = 10$ minutes and $B_{1/2} = 50$ minutes. Use a time unit of 1 minute. Plot u_A and u_B versus time in the same plot. Filename: `radioactive_decay_2subst.py`.

Exercise 31: Determine time of murder from body temperature

A detective measures the temperature of a dead body to be 26.7 C at 2 pm. One our later the temperature is 25.8 C. The question is when death occured.

Assume that Newton's law of cooling (98) is an appropriate mathematical model for the evolution of the temperature in the body. First, determine k in (98) by formulating a Forward Euler approximation with one time step from time 2 am to time 3 am, where knowning the two temperatures allows for finding k. Thereafter, simulate the temperature evolution from the time of murder, taken as $t = 0$, when $T = 37$ C, until the temperature reaches 26.7 C. The corresponding time allows for answering when deach occured.

Filename: `detective.py`.

Exercise 32: Compute the signum function by solving an ODE

Consider the ODE problem

$$y'(x) = \begin{cases} -1, & x < 0, \\ 1, & x \geq 0 \end{cases} \qquad x \in (-1, 1], \quad y(1-) = 1 \,.$$

The solution is $y(x) = |x|$.

Using a mesh $x_0 = -1$, $x_1 = 0$, and $x_2 = 1$, calculate by hand y_1 and y_2 from the Forward Euler, Backward Euler, Crank-Nicolson, and Leapfrog methods. Use all of the former three methods for computing the y_1 value to be used in the Leapfrog calculation of y_2. Thereafter, visualize how these schemes perform for a uniformly partitioned mesh with $N = 10$ and $N = 11$ points. Filename: `signum.py`.

Exercise 33: Simulate growth of a fortune with random interest rate

The goal of this exercise is to compute the value of a fortune subject to inflation and a random interest rate. Suppose that the inflation is constant at i percent per year and that the annual interest rate, p, changes randomly at each time step, starting at some value p_0 at $t = 0$. The random change is from a value p^n at $t = t_n$ to $p_n + 0.5$ with probability 0.25 and $p_n - 0.5$ with probability 0.25. No change occurs with probability 0.5. Use a time step of one month, $p_0 = i$, initial fortune scaled to 1, and simulate 1000 scenarios of length 20 years. Compute the mean evolution of the money and the standard deviation. Plot these two curves along with the $p^n - i$ curve. Filename: `random_interest.py`.

Hint 1. The following code snippet computes p^{n+1}:

```
import random

def new_interest_rate(p_n, increment=0.5):
    r = random.random()  # uniformly distr. random number in [0,1)
    if 0 <= r < 0.25:
        p_np1 = p_n + increment
    elif 0.25 <= r < 0.5:
        p_np1 = p_n - increment
    else:
        p_np1 = p_n
    return p_np1
```

Hint 2. If $u_i(t)$ is the value of the fortune in experiment number i, $i = 0, \ldots, N - 1$, the mean evolution of the fortune is

$$\bar{u}(t) = \frac{1}{N} \sum_{i=0}^{N-1} u_i(t),$$

and the standard deviation

$$s(t) = \sqrt{\frac{1}{N-1} \left(-(\bar{u}(t))^2 + \sum_{i=0}^{N-1} (u_i(t))^2 \right)}.$$

Suppose $u_i(t)$ is stored in an array `u`. The mean and the standard deviation of the fortune is most efficiently computed by using to accumulation arrays, `sum_u`

and `sum_u2`, and performing `sum_u += u` and `sum_u2 += u**2` after every experiment. This technique avoids storing all the $u_i(t)$ time series for computing the statistics.

Index